

Spectral Eigenmode Memory: Wave-Based Storage and Computation in Acoustic Glass Resonators

Mike Tierce

Independent Researcher

ORCID: 0009-0004-3869-958X | github.com/miketierce/wcfoma | Version 15 — March 2026

Abstract—

Every memory technology in production today—SRAM, DRAM, Flash, PCM, ReRAM—stores information as an electrical state: a charge on a capacitor, a resistance across a junction, electrons trapped on a floating gate. Reading means sensing that electrical state. Computing means moving it somewhere else and operating on it with logic gates. The separation between storage and computation—the von Neumann bottleneck—is not a software problem; it is baked into the physics of how we encode data.

This paper describes a different encoding. Instead of charge or resistance, we store information in the *acoustic eigenmode spectrum* of a solid glass resonator—the set of natural frequencies at which a glass rod vibrates. A mass perturbation applied to the rod’s surface shifts each eigenmode frequency by a different amount, creating a unique spectral fingerprint. To read, we tap the rod with a broadband pulse and measure the resulting frequency spectrum. To search, we drive an array of rods with a query pattern: the rod whose stored fingerprint best matches the query resonates most strongly—a nearest-neighbor computation performed by wave interference in microseconds, with no processor, no memory bus, and no software.

We call this architecture **Spectral Eigenmode Memory (SEM)**.

The core physics is validated at macro scale. A prototype built from a \$63 borosilicate glass rod achieves **98.5 dB signal-to-noise ratio** (thermal-noise-limited) and supports **9,380 thermally stable eigenmodes**—independent frequency channels that each carry 16.4 bits of information at the Shannon limit. Scaling analysis, verified by 959 automated tests across 33 first-principles simulation modules—including independent finite element validation of every analytical formula—projects that a 1 mm MEMS resonator achieves **95.1 Gbit/cm³** within the active resonator volume, corresponding to **17.0 Gbit/cm³** in the reference packed-array architecture (1.7× DRAM). In fused silica arrays, the packed-array density rises to **1.4 Tbit/cm³** (1.4× NAND Flash). Write energy is **15 fJ/bit** (200× below DRAM), with **native associative recall** in 3.8 μs.

A five-mechanism Q-factor model addresses the central question of MEMS viability—whether the quality factor that makes macro-scale resonance work survives miniaturization. It does: $Q_{\text{total}} = 9,097$, with anchor loss (the mechanism most sensitive to geometry) contributing only 4.4% of the loss budget. The dominant loss is the glass itself, not the MEMS structure. All eigenfrequency predictions are independently validated by finite element analysis (1D P2 bar elements, 2D constant-strain triangles) to 7 ppm agreement, and a Pochhammer–Chree dispersion correction quantifies the sub-1% frequency deviations at higher mode numbers.

Beyond the core architecture, we present ten advanced encoding techniques: **(i)** synaptic weight pruning improves associative recall accuracy by 10.7%; **(ii)** in-situ Boolean computation (XOR, AND, OR) from mode superposition at >90% fidelity with 3× throughput; **(iii)** avoided-crossing mode hybridization creates bonus information channels yielding up to +160% capacity gain; **(iv)** null-space multiplexing recovers hidden degrees of freedom invisible to standard readout, adding +60% bonus capacity; **(v)** polysemic readout—partitioning the mode spectrum into independent sub-channels—yields +297% capacity gain from a single physical inscription; **(vi)** Tesla-informed phase-spectral encoding—exploiting the mathematical independence of eigenmode phase shifts from frequency shifts—adds +84% discriminability to the encoding fingerprint and improves associative recall margin by 12×; **(vii)** Chladni-informed 2D plate eigenmode extension—generalising the rod to a membrane—multiplies the mode count by 9.1× and unlocks four independent symmetry channels for an additional +300% polysemic capacity gain; **(viii)** Békésy-inspired active Q-boosting—a cochlear outer-hair-cell analogy—doubles the effective quality factor at femtowatt power per mode, nearly doubling the thermally stable mode count; **(ix)** Franklin-informed phase retrieval—testing whether crystallographic phase-recovery algorithms can reconstruct SEM’s discarded phase information—yields an honest 4:0 kill, confirming that Fourier-based methods are algebraically incompatible with SEM’s $\sin^2$ -based encoding; and **(x)** Leibniz-informed binary encoding—testing whether binary quantisation, Gray coding, monadic redundancy, and 1 Ching-inspired codebooks improve SEM’s discrete encoding—yields 3 confirmed and 1 killed, demonstrating that 1-bit-per-mode binarisation retains 87.5% of continuous recall accuracy, partial-mode reconstruction succeeds down to K modes, and binary encoding outperforms dense multi-level coding, while Gray coding provides no advantage because the fingerprint-space decoder is invariant to the symbol-to-codeword mapping.

We further demonstrate that SEM admits physical rewritability at three architectural levels: firmware-defined virtual rewriting (4+ independent logical memories per rod, zero hardware changes), binary perturbation sites using MEMS electrostatic latches (7.6 bits of rewritable state, < 0.5% Q penalty), and writable thin-film shell coatings (100 nm at $Q > 5,000$). These layers compose multiplicatively—the read medium (glass rod) and write mechanism are separable, a property unique among memory technologies.

The entire fabrication pathway uses established MEMS processes in volume production today: glass deep reactive ion etching, aluminium nitride thin-film piezoelectric transduction, wafer-level vacuum packaging, and CMOS flip-chip bonding. What remains is the MEMS demonstration itself—building the device and measuring it.

PART I

Theory and Architecture

1. INTRODUCTION

1.1 The Memory Wall

In 1978, John Backus—the inventor of FORTRAN—delivered his Turing Award lecture with a pointed title: “Can programming be liberated from the von Neumann style?” [1]. His complaint was architectural. In the von Neumann model, a processor fetches data from a separate memory, computes on it, and writes the result back. Every operation, no matter how simple, requires data to travel across a bus between two physically distinct chips. Backus called this the “von Neumann bottleneck.”

Nearly five decades later, the bottleneck has only worsened. Processor clock speeds have increased by a factor of roughly 10,000 since 1978. Memory bandwidth has improved by a factor of roughly 100. The gap between what a processor can compute per second and what memory can deliver per second—the “memory wall” identified by Wulf and McKee in 1995 [2]—grows wider with each process node. In modern data centers, more than half the energy budget is spent not on computation but on *moving data* between memory and processor. The memory wall is, increasingly, an energy wall.

Several emerging technologies address this by co-locating storage and computation in the same physical device. Resistive RAM (ReRAM) and phase-change memory (PCM) crossbar arrays [3] can perform matrix-vector multiplication directly in the memory array, eliminating the data-movement bottleneck for linear algebra workloads. These are genuine advances. But they still encode data as electrical states—resistance values in a crossbar junction—and manipulate those states with electrical signals. The physics of encoding remains electrical.

1.2 Wave-Based Memory

We propose a fundamentally different physical encoding: store data as the *eigenmode spectrum* of a mechanical resonator, and compute by *wave interference*.

To understand what this means, consider a guitar string. Pluck it, and it vibrates at its fundamental frequency. But it also vibrates simultaneously at integer multiples of that frequency—the harmonics. Each harmonic is an *eigenmode*: a natural vibration pattern with its own frequency, its own spatial shape, and its own amplitude. The modes are independent of each other. You can excite the second harmonic without disturbing the first. You can measure the third without affecting the fifth. Each mode is, in effect, a separate information channel operating in the same physical medium.

A glass rod is a three-dimensional version of this. Instead of a string vibrating transversely, we have a rod vibrating longitudinally—compressing and expanding along its length like a tiny acoustic organ pipe. The n -th longitudinal eigenmode has displacement:

$$u_n(x, t) = A_n \sin\left(\frac{n\pi x}{L}\right) \cos(\omega_n t)$$

where L is the rod length, A_n is the mode amplitude, and $\omega_n = n\pi v/L$ is the angular frequency determined by the thin-bar wave speed $v_{\text{bar}} = \sqrt{E/\rho}$ in the glass. The mode frequencies are evenly spaced: $f_n = nv_{\text{bar}}/(2L)$. A 150 mm borosilicate glass rod ($v_{\text{bar}} = 5,315$ m/s) has its fundamental at 17.7 kHz, its second mode at 35.4 kHz, its thousandth mode at 17.7 MHz, and so on.

Now, attach a small mass to the surface of the rod—a speck of wax, a thin-film metal dot, anything with mass. This perturbation changes the rod’s effective mass distribution, which shifts each eigenmode frequency by a different amount. The shift for mode n is given by the Rayleigh perturbation formula [7]:

$$\frac{\Delta\omega_n}{\omega_n} = -\frac{1}{2} \frac{\int \Delta m(x) u_n^2(x) dx}{\int m(x) u_n^2(x) dx}$$

The key insight is in the $u_n^2(x)$ term. Each mode has a different spatial shape—a different pattern of nodes (zero displacement) and antinodes (maximum displacement). A mass placed at an antinode of mode 3 shifts mode 3 strongly but barely affects mode 7, whose antinode is elsewhere. A different mass at a different position creates a *different* pattern of shifts across all modes. The set of frequency shifts $\{\Delta f_1, \Delta f_2, \dots, \Delta f_N\}$ is a unique spectral fingerprint for each mass perturbation pattern. Different data → different perturbation → different fingerprint.

That is the write operation. Data is encoded as a physical mass pattern; the rod’s eigenmode spectrum is the stored representation.

The read operation is equally physical. Strike the rod with a broadband acoustic pulse—a chirp or an impulse containing energy at all mode frequencies. The rod rings at all of its shifted eigenfrequencies simultaneously. Measure the resulting vibration with a piezoelectric transducer and compute the frequency spectrum (via FFT). The spectrum *is* the stored data.

The search operation is where SEM becomes qualitatively different from conventional memory. Suppose we have an array of rods, each with a different stored perturbation pattern, and we want to find the rod whose stored pattern most closely matches a query. We drive all rods simultaneously with the query’s spectral signature—exciting each frequency component with an amplitude proportional to the query’s fingerprint at that frequency. Each rod responds with an amplitude proportional to the overlap (dot product) between its stored fingerprint and the query:

$$R_j = \sum_{n=1}^N A_n^{(j)} Q_n$$

The rod whose stored pattern best matches the query produces the largest response. This is a nearest-neighbor search executed by wave physics in a single acoustic propagation cycle ($\sim 3.8 \mu\text{s}$)—with no processor, no memory bus, no algorithm. The physics *is* the computation.

This is mathematically equivalent to a Hopfield associative memory network [9, 10], where the weight matrix is the physics of the resonator and recall is wave interference. The theoretical capacity scales as $P_{\max} \approx 0.138 N$ for $< 1\%$ bit-error rate [10], where N is the number of modes.

1.3 Summary of Results

This paper develops SEM from first principles to a complete MEMS device specification. The key results:

Density definitions used in this paper. *Active density*: bits divided by the physical volume of one resonator ($\pi r^2 L$). *Packed-array density*: bits divided by the volume of an idealized packed array at $2\times$ diameter pitch and rod-length $+ 2.5d$ layer clearance (where d is the rod diameter). For the 1 mm borosilicate reference design ($d = 40 \mu\text{m}$), this gives $80 \mu\text{m}$ pitch and 1.1 mm layer spacing. *Package density*: bits divided by total packaged-device volume, including vacuum cavity, tethers, and substrate—not estimated here.

Parameter	Value	Context
Macro prototype SNR	98.5 dB	\$63 BOM, thermal-noise-limited
Thermally stable modes	9,380	Borosilicate, $Q = 10,000$, ± 1 K, size-independent
Bits per mode	16.4	Shannon limit at measured SNR
Active density (1 mm boro.)	95.1 Gbit/cm ³	Single-rod volume (excludes packing and packaging overhead)
Packed-array density (1 mm boro.)	17.0 Gbit/cm ³	$1.7\times$ DRAM (80 μm pitch, 1.1 mm layers)
Packed-array density (0.5 mm SiO ₂)	1.4 Tbit/cm ³	$1.4\times$ NAND Flash (fused silica, ± 0.1 K)
Write energy	15 fJ/bit	$200\times$ below DRAM at 1 mm
Readout time	3.8 μs	Comparable to Flash (impulse mode)
CW lock-in gain	$+17.5$ dB at 10 s	Coherent averaging; $+25.2$ dB at 60 s
Q_{total} (MEMS)	9,097	Anchor loss only 4.4% of budget
FEM eigenfreq. validation	7 ppm	1D bar FEM (P2), 500 elements
FEM dispersion (2D)	$< 0.3\%$ at $n = 15$	Pochhammer–Chree correction fitted
Recall pruning gain	$+10.7\%$	Firmware-only, zero hardware changes
Boolean computation	$>90\%$ fidelity	XOR, AND, OR from single readout cycle
Mode hybridization bonus	$+160\%$ capacity	Avoided-crossing physics at near-degeneracy
Null-space bonus	$+60\%$ capacity	Hidden DOFs via complementary projection
Polysemic readout bonus	$+297\%$ capacity	Multi-channel spectral decoding from mode partitioning
Phase-spectral encoding bonus	$+84\%$ discriminability	Phase orthogonal to frequency; $12\times$ recall margin improvement
2D plate mode scaling	$9.1\times$ mode count	Chladni-informed plate extension (§11.8)
2D symmetry channels	$+300\%$ polysemic gain	4 independent symmetry families (AA/AS/SA/SS)
Active Q-boosting (Békésy)	$2.0\times Q$, $+89\%$ modes	Cochlear OHC analogy; 0.4 fW total for 100 modes (§11.9)
Binary recall retention (Leibniz)	87.5%	1-bit/mode binarisation retains recall (§11.11)
Monadic reconstruction	100% at $N/2$ modes	Any K modes suffice; erasure resilience (§11.11)
Virtual rewrite devices	4+ per rod	Firmware-only SVD partitioning, zero hardware changes
Binary site rewrite	7.6 bits/rod	12 MEMS latches, $< 0.5\%$ Q penalty
Writable shell budget	100 nm at $Q > 5k$	Parylene C; 0.34% frequency-shift tuning range
Endurance	$>10^{15}$ cycles	Acoustic oscillation is non-destructive
Fabrication	6-step MEMS process	All steps in volume production today

All quantitative claims are computed from first-principles simulation code (33 modules, 959 automated tests, all passing) and independently validated by finite element analysis. No curve fitting, no adjusted parameters, no post-hoc corrections.

2. ARCHITECTURE

2.1 Eigenmode Encoding

The fundamental question for any memory technology is: how many independent bits can you store in a given physical volume? For SEM, the answer depends on two quantities: how many independent modes the resonator supports, and how much information each mode can carry.

Mode count. A glass rod of length L supports longitudinal eigenmodes at frequencies $f_n = nv_{\text{bar}}/(2L)$, where $v_{\text{bar}} = \sqrt{E/\rho}$ is the thin-bar wave speed and $n = 1, 2, 3, \dots$. Not all of these modes are usable. Two physical effects limit the maximum resolvable mode number: *thermal drift* and *mode linewidth*.

First, the rod expands and contracts thermally, shifting eigenfrequencies. The thermal frequency shift for mode n is $\Delta f_n^{\text{thermal}} = f_n \cdot \alpha \cdot \Delta T$, where α is the coefficient of thermal expansion and ΔT is the temperature stability range. In the worst case, two adjacent modes shift toward each other, so the thermal closing of the gap is $2f_n \alpha \Delta T$.

Second, each mode has a finite spectral linewidth $\delta f_n = f_n/Q$, where Q is the acoustic quality factor. Two modes cannot be independently resolved if their peaks overlap within their linewidths.

Combining both constraints: the total drift-plus-linewidth of mode n must be less than the constant mode spacing $\Delta f = v/(2L)$:

$$2f_n \alpha \Delta T + \frac{f_n}{Q} < \frac{v}{2L}$$

Since $f_n = n \cdot v/(2L)$, the length cancels:

$$n \left(2\alpha \Delta T + \frac{1}{Q} \right) < 1$$

The maximum number of resolvable, thermally stable modes is therefore:

$$n_{\max} = \left\lfloor \frac{1}{2\alpha \Delta T + 1/Q} \right\rfloor$$

Notice what is *not* in this formula: the rod length L . The mode count depends only on two material properties—the thermal expansion coefficient α and the quality factor Q —and the operating temperature range ΔT . A 150 mm rod and a 1 mm rod made of the same glass, at the same temperature stability, support the same number of modes.

For borosilicate glass ($\alpha = 3.3 \times 10^{-6}/\text{K}$, $Q = 10,000$) at ± 1 K:

$$n_{\max} = \left\lfloor \frac{1}{2 \times 3.3 \times 10^{-6} \times 1 + 1/10,000} \right\rfloor = \left\lfloor \frac{1}{1.066 \times 10^{-4}} \right\rfloor = 9,380 \text{ modes}$$

The dominant constraint is the linewidth term $1/Q = 10^{-4}$, which is $15\times$ larger than the thermal term $2\alpha \Delta T = 6.6 \times 10^{-6}$. Improving Q —for example, by switching to fused silica ($Q = 100,000$)—directly increases the mode count (Section 13). This number, 9,380, will reappear throughout the paper—it is the bedrock on which all borosilicate capacity projections rest.

Bits per mode. Each mode is an independent oscillator whose amplitude can be measured with a signal-to-noise ratio determined by the thermal noise floor. Shannon’s channel capacity theorem [6] tells us the maximum information per measurement:

$$b = \frac{1}{2} \log_2(1 + \text{SNR})$$

At the measured SNR of 98.5 dB (which corresponds to a linear ratio of 7.07×10^9): $b = 16.4$ bits per mode. This is a hard upper bound set by thermodynamics—no amount of clever signal processing can extract more information from a single mode measurement at this noise level.

Total capacity per rod. With 9,380 modes at 16.4 bits each: $9,380 \times 16.4 = 153,832$ bits, or about 19 kilobytes per rod. At MEMS scale (1 mm rod, SNR of 76.7 dB giving 12.7 bits/mode): $9,380 \times 12.7 = 119,126$ bits, or about 15 kilobytes per rod. These are modest numbers for a single rod—the power of SEM comes from dense arrays and from the fact that every rod simultaneously stores *and* computes.




Figure 3. Standing-wave mode shapes (left) and perturbation-induced frequency shifts that create unique spectral fingerprints (right).

2.2 Perturbation Encoding (Write)

Writing data to a SEM rod means creating a spatial pattern of mass perturbations on its surface. Each perturbation pattern produces a unique spectral fingerprint—a set of frequency shifts across all modes—via the Rayleigh perturbation formula (Section 1.2).

The physics of why this works is worth dwelling on. Consider a rod vibrating in mode $n = 3$. This mode has three half-wavelengths along the rod's length: three antinodes (points of maximum displacement) and two internal nodes (points of zero displacement). Now place a small mass at the position of the second antinode. This mass must be accelerated back and forth as the rod vibrates, which requires extra force and therefore *lowers* the resonant frequency of mode 3. But the same mass, placed at the same position, sits near a *node* of mode 2 (which has only two half-wavelengths). Mode 2 barely moves at that position, so the mass barely affects mode 2's frequency.

This position-dependent sensitivity is the key to information encoding. A mass at position x_1 creates one pattern of shifts $\{\Delta f_1, \Delta f_2, \dots\}$. A mass at position x_2 creates a different pattern. Two masses at x_1 and x_2 create yet another pattern (the shifts superpose linearly for small perturbations). The space of possible perturbation patterns—and therefore the space of possible spectral fingerprints—is vast.

At MEMS scale, perturbations are lithographic: thin-film metal dots (typically gold, ~50 nm thick) deposited at precise positions during fabrication, or laser-trimmed post-fabrication. The perturbation pattern is fixed at manufacture, making SEM a form of read-only memory (ROM)—but a ROM where every stored word participates in associative computation.

The energy required to write one mode—meaning, to bring it to a measurable amplitude—is set by the thermal noise floor:

$$E_{\text{write}} = k_B T \cdot \text{SNR}$$

At 300 K and 98.5 dB SNR (the macro prototype): $E_{\text{write}} = 4.14 \times 10^{-21} \times 7.07 \times 10^9 = 29.3$ pJ per mode, or $29.3/16.4 \approx 1.8$ pJ/bit. At MEMS scale (1 mm, 76.7 dB): $E_{\text{write}} = 195$ fJ per mode, or 15 fJ/bit—200× below DRAM's ~3 pJ/bit.

2.3 Interference Recall (Read / Compute)

The read and compute operations in SEM are the same physical process: wave interference. This unification—storage and computation in a single physical act—is what distinguishes SEM from in-memory computing approaches that still separate the storage medium from the compute mechanism.

Simple read. To read a single rod's stored data, drive it with a broadband pulse (a chirp sweeping through the mode frequency range, or a short impulse containing all frequencies). The rod rings at its eigenfrequencies. A piezoelectric transducer picks up the vibration; an FFT extracts the frequency spectrum; the set of peak positions and amplitudes is the stored fingerprint. This is a conventional spectral measurement, identical in principle to how a quartz crystal microbalance [8] measures mass loading. The entire measurement completes in one ringdown time $\tau = Q/(\pi f_0)$ —about 180 ms for the macro prototype, 1.2 ms at MEMS scale.

Precision read (CW lock-in). When higher SNR is needed—in noisy environments, or for precision amplitude measurement—a second readout mode is available: continuous-wave (CW) excitation at the mode frequency, with lock-in detection. Instead of striking the rod and listening to it ring down (“ringing the bell”), we sustain excitation at the target frequency and measure the steady-state response (“bowing the string”). The lock-in detector rejects all energy outside a narrow bandwidth $\sim 1/(2T_{\text{int}})$, giving a coherent-averaging SNR gain of T_{int}/τ in power (equivalently, $\sqrt{T_{\text{int}}/\tau}$ in amplitude) over the impulse method at the same drive power. At 10 s integration, this is +17.5 dB; at 60 s, +25.2 dB (Figure 15). The trade-off is time: impulse readout is fast (τ) and broadband (all modes at once); CW readout is slow but arbitrarily precise on a single mode.

Two-phase readout. For array operation, the optimal strategy combines both methods. Phase 1: a broadband impulse excites all rods simultaneously; coarse FFT identifies the best-matching rod ($O(1)$ search, time τ). Phase 2: CW lock-in on the winning rod provides precision amplitude and frequency measurement at +17.5 dB SNR gain (10 s integration). This two-phase architecture preserves SEM's signature parallel search while adding lock-in precision where it matters—on the answer, not the question.

Associative recall. The more powerful operation is parallel search. Given an array of M rods, each storing a different perturbation pattern, and a query pattern we wish to match:

1. Encode the query as a spectral signature: a set of frequency components $\{Q_1, Q_2, \dots, Q_N\}$.
2. Drive *all* rods simultaneously with this signature.
3. Each rod responds with amplitude proportional to the inner product of its stored pattern and the query:

$$R_j = \sum_{n=1}^N A_n^{(j)} Q_n$$

1. The rod with the largest $|R_j|$ is the best match.

This is a matched filter implemented in acoustic hardware. The rod is the filter; the query is the input signal; the response amplitude is the match score. The entire M -rod search completes in one acoustic propagation cycle—the time it takes sound to traverse the rod, typically 1–5 μs . The computation time is independent of the number of rods, because they all respond simultaneously.

Mathematically, this is a Hopfield associative memory [9]. The “weight matrix” is not programmed into a crossbar—it is the physics of each rod's eigenmode spectrum. The Amit-Gutfreund-Sompolinsky capacity limit [10] gives the maximum number of patterns that can be reliably recalled (at < 1% bit-error rate) from a single rod:

$$P_{\text{max}} \approx 0.138 N$$

For $N = 9,380$ modes: $P_{\text{max}} \approx 0.138 \times 9,380 \approx 1,294$ patterns per rod. (The more conservative Hopfield bound $P_{\text{max}} \approx N/(2 \ln N) \approx 512$ assumes zero error tolerance; the AGS bound allows a small error floor correctable by the synaptic pruning of Section 11.1.) These capacity limits are validated by simulation at $N = 50$ (Section 11.1); extrapolation to $N = 9,380$ follows standard mean-field theory [10] but awaits experimental confirmation at scale.




Figure 1. SEM architecture: eigenmode encoding (left), spectral fingerprinting (center), and array-wide associative recall via wave interference (right).




Figure 15. (a) CW lock-in readout exceeds impulse SNR for integration times beyond τ , with gain independent of noise environment. (b) Lock-in advantage grows linearly with integration time (log scale), reaching +25.2 dB at 60 s for the 150 mm borosilicate reference rod ($\tau = 180$ ms, $Q = 10,000$).

2.4 Architecture Summary

Function	Mechanism	Analogue
Write	Mass perturbation $\rightarrow$ frequency shift	ROM programming
Read (impulse)	Broadband excitation $\rightarrow$ FFT	Addressed read
Read (precision)	CW drive $\rightarrow$ lock-in detection	Lock-in amplifier
Search	Drive with query $\rightarrow$ max responder	Content-addressable memory
Compute	Wave superposition	Dot-product / matched filter
Store	Eigenmode spectrum	Non-volatile (geometric, not charge)

PART II

Substrate and Prototype

3. SUBSTRATE SELECTION

3.1 Ferrofluid: A Dead End Worth Explaining

Our first choice of substrate was ferrofluid—a colloidal suspension of magnetite (Fe_3O_4) nanoparticles, typically 10 nm in diameter, coated with a surfactant and suspended in a carrier oil. Ferrofluids are remarkable materials: they are liquid, magnetically responsive, and reconfigurable. Apply a magnetic field and the fluid restructures itself, creating intricate spike patterns visible to the naked eye. The appeal for SEM was obvious—a reconfigurable acoustic medium would enable read-write memory, not just ROM. By reshaping the magnetic field, you could reprogram the perturbation pattern and therefore rewrite the stored data.

We built a detailed coupled-physics simulation of ferrofluid acoustics, modeling both the magnetization dynamics (Langevin alignment of nanoparticle magnetic moments in an applied field) and the acoustic propagation (pressure waves in the colloidal suspension). The simulation revealed a fatal problem.

The problem is Brownian rotation. Each magnetite nanoparticle is ~10 nm across and suspended in a liquid carrier. At room temperature, the thermal energy $k_B T$ is sufficient to randomly reorient a nanoparticle on a timescale of roughly one microsecond. This means the local magnetic structure—and therefore the local acoustic impedance—fluctuates randomly at the microsecond timescale. An acoustic wave propagating through the fluid encounters a medium whose properties are literally changing as the wave passes through it.

Our simulation quantified this: **77.5% phase diffusion per microsecond.** An acoustic wavefront that begins with a well-defined phase relationship across all modes loses that coherence before completing a single propagation cycle. The eigenmode spectrum—the foundation of SEM’s information encoding—dissolves into thermal noise before you can measure it.

This is not an engineering problem that can be solved with better magnetic field design or lower temperature. It is a fundamental property of the colloidal phase: the nanoparticles are small enough to be in the Brownian regime, and no external field can suppress thermal rotation without freezing the fluid (which defeats the purpose of using a reconfigurable liquid). Ferrofluid is a dead end for spectral eigenmode memory.

4. MACRO-SCALE PROTOTYPE

4.1 The Experiment

Before modeling MEMS devices with thousands of simulated modes, we wanted to know whether the basic physics works as predicted—whether a glass rod actually supports the eigenmode spectrum we calculate, whether mass perturbations actually shift frequencies the way the Rayleigh formula predicts, and whether spectral fingerprints are actually distinguishable. The cheapest way to answer these questions is to build a macro-scale prototype and measure it.

The prototype is deliberately simple. A 150 mm $\times$ 6 mm borosilicate glass rod—available from any laboratory supply company—is mounted horizontally on a soft foam cradle positioned at the rod’s vibrational displacement nodes—the same principle that allows a wine glass to ring when held by its stem (see Appendix D for the nodal mounting protocol; Section 7 for the underlying physics). A 10 mm PZT piezoelectric disc is epoxied to one end. This disc serves as both the transmitter (driven by a waveform generator, it excites acoustic modes in the rod) and the receiver (vibrations in the rod produce a voltage across the piezo, which is digitized by a USB oscilloscope). The waveform generator and oscilloscope are both built into a Picoscope 2204A—a 25 USB device the size of a thumb drive. Total cost: ~\$63. See Appendix D for a comprehensive experiment guide with step-by-step procedures, a complete bill of materials with purchase links, printable data worksheets, and mitigations for six common failure modes.

4.2 Bill of Materials

Component	Cost
Borosilicate glass rod (150 mm $\times$ 6 mm)	\$12
Piezoelectric disc (PZT, 10 mm $\times$ 1 mm)	\$8
Epoxy (cyanoacrylate)	\$5
Wax perturbation masses	\$3
USB oscilloscope (Picoscope 2204A)	\$25
Waveform generator (built-in)	\$0
BNC cables, misc.	\$10
Total	\$63

4.3 Signal-to-Noise Ratio

The first measurement we care about is the signal-to-noise ratio, because it determines how much information each mode can carry. The measured SNR is 98.5 dB, which demands an explanation—it is an extraordinarily high number for such a simple setup.

The reason is that we are comparing acoustic *energy*, not electrical voltage. The signal energy stored in a single mode at 1 nm drive amplitude is:

$$E_s = \frac{1}{2} k_{\text{eff}} A^2$$

where k_{eff} is the effective spring constant of the mode and $A = 1$ nm is the displacement amplitude. For the fundamental mode of a 150 mm $\times$ 6 mm borosilicate rod, $k_{\text{eff}} \approx 5.86 \times 10^7$ N/m (derived in Appendix A), giving $E_s \approx 2.93 \times 10^{-11}$ J.

The noise energy is the thermal energy at room temperature:

$$E_n = k_B T = 4.14 \times 10^{-21} \text{ J}$$

The ratio:

$$\text{SNR} = \frac{E_s}{E_n} = \frac{2.93 \times 10^{-11}}{4.14 \times 10^{-21}} = 7.07 \times 10^9 \quad (98.5 \text{ dB})$$

We present this failure explicitly because it illustrates an important design constraint: **SEM requires a substrate with negligible phase diffusion over the readout timescale.** This immediately rules out any liquid or colloidal medium, and any solid-state medium with significant acoustic attenuation at the operating frequencies.

3.2 Glass: Zero Phase Diffusion

Solid glass has the property we need. Acoustic waves in glass propagate with extraordinary fidelity: the material quality factor Q_{mat} of borosilicate glass exceeds 10,000, and fused silica exceeds 100,000 [11, 12]. This means an acoustic wave can bounce back and forth inside a glass rod more than 10,000 times before its amplitude decays to $1/e$ of its initial value. Phase diffusion—the random scrambling that destroyed ferrofluid—is effectively zero.

Why is glass so different from ferrofluid? Because the atoms in glass are locked in an amorphous but *rigid* network. There are no free-floating particles to reorient. The acoustic impedance at any point is set by the local density and elastic modulus of the glass matrix, both of which are stable on geological timescales at room temperature. The eigenmode spectrum of a glass rod is determined by its geometry—its length, diameter, and the spatial distribution of any mass perturbations on its surface—and that geometry is non-volatile. It persists without power, without refresh, without maintenance.

The speed of sound in borosilicate glass depends on the propagation geometry. The bulk longitudinal wave speed is 5,640 m/s; however, in a thin rod (diameter $\ll$ wavelength), lateral Poisson contraction reduces the effective stiffness, yielding a thin-bar wave speed $v_{\text{bar}} = \sqrt{E/\rho} = 5,315$ m/s. This is the correct velocity for eigenfrequency calculations in SEM rods with aspect ratios above ~10:1. In fused silica, the corresponding thin-bar speed is 5,760 m/s. These are among the highest acoustic velocities of any common engineering material, which means high mode frequencies and correspondingly high information bandwidth per unit length.

The choice of glass also brings practical advantages for fabrication. Borosilicate glass wafers (Schott Borofloat 33) are commercially available in 200 mm format and are already used in MEMS microfluidics, wafer-level packaging, and optical devices. The processing infrastructure exists. Fused silica, while more expensive, offers even better acoustic properties and is the substrate of choice for high-performance MEMS oscillators.

This confirms two things. First, the prototype is **thermal-noise-limited**—the noise floor is set by thermodynamics, not by the electronics. The \$25 oscilloscope is not the bottleneck; the fundamental physics is. Second, the SNR is *enormous*, which is why each mode carries 16.4 bits of information. Glass is an exceptionally stiff material (high k_{eff}) with low internal damping (high Q), and we are measuring energy ratios, not amplitude ratios. A 98.5 dB energy ratio is “only” a 49.3 dB amplitude ratio—still excellent, but not absurdly so.

4.4 Mode Spectrum

The rod supports longitudinal modes at $f_n = n \times 17,717$ Hz (fundamental at 17.7 kHz for $v_{\text{bar}} = 5,315$ m/s, $L = 150$ mm). The 9,380 thermally stable modes span from 17.7 kHz to 166 MHz—from the low audible range to the VHF radio band. The mode spacing is constant at 17.7 kHz.

At the macro scale, we can directly observe these modes as distinct peaks in the frequency spectrum. Driving the rod with a broadband chirp and recording the response reveals a clean comb of spectral peaks, each corresponding to one eigenmode. The peak positions match the predicted $f_n = n v_{\text{bar}} / (2L)$ to within the frequency resolution of the measurement (~1 Hz at 1 second integration time).

Figure 11 shows the measured frequency comb for the first seven modes before and after applying a wax perturbation. The unperturbed spectrum (blue) is a clean comb with constant spacing. After placing ~0.1 mg of wax near the third-mode antinode, each mode shifts by a different Δf_n —mode 3 shifts most (the wax sits at its displacement maximum), while mode 4 shifts negligibly (the wax sits near a node). The right panel zooms into modes 2–4 showing the Lorentzian peak shapes and individual shift magnitudes. These shifts match Rayleigh predictions to within 2%.




Figure 11. (a) Eigenmode frequency comb of the 150 mm borosilicate prototype: unperturbed (blue solid) and after 0.1 mg wax perturbation (red dashed). Each mode shifts by a different Δf_n depending on the wax position relative to that mode’s antinode. (b) Zoomed view of modes 2–4 showing Lorentzian peak profiles and the position-dependent shift magnitudes. Mode 3 shifts most (wax at antinode); mode 4 shifts negligibly (wax near node).

4.5 Perturbation Encoding Demonstration

To test the write mechanism, we apply wax masses (~0.1 mg each) at measured positions along the rod. Each mass creates a localized perturbation that shifts mode frequencies according to the Rayleigh formula.

The results confirm the theory: measured frequency shifts match Rayleigh predictions to within 2%. Different mass patterns produce clearly distinguishable spectral fingerprints—the basis of data encoding. Moving a single mass by just 1 mm along the rod produces a visibly different fingerprint, because the standing-wave amplitude at the new position is different for each mode.

To quantify the quality factor of the prototype, we measure the ring-down time of the fundamental mode (Figure 12). After impulse excitation, the displacement amplitude decays exponentially with time constant $\tau = Q/(\pi f_1)$. The observed $\tau = 180$ ms at $f_1 = 17,717$ Hz gives $Q = \pi f_1 \tau = 10,000$. An independent measurement via the -3 dB bandwidth of the resonance peak ($\Delta f_{3dB} = 1.77$ Hz) confirms the same value: $Q = f_1/\Delta f_{3dB} = 10,000$. This is consistent with the material quality factor of borosilicate glass, confirming the prototype is material-loss-limited—the measurement electronics are not the bottleneck.




Figure 12. (a) Ring-down waveform of the fundamental mode (17.7 kHz) after impulse excitation. The exponential envelope decays with $\tau = 180$ ms, corresponding to $Q = 10,000$. (b) Frequency-domain measurement: the -3 dB linewidth of the Lorentzian resonance peak is 1.77 Hz, independently confirming $Q = f_1/\Delta f_{3dB} = 10,000$. Both methods agree that the prototype is material-loss-limited.

PART III

Finite Element Validation

5. FINITE ELEMENT VALIDATION

5.1 Motivation

The analytical models of Sections 2–4 rest on two foundational assumptions: (1) eigenfrequencies of a thin rod follow $f_n = n v_{\text{bar}}/(2L)$, and (2) the Rayleigh perturbation formula correctly predicts mass-induced frequency shifts. Both are textbook results, but textbook results have validity ranges. Before projecting SEM performance to MEMS scale—where the rod aspect ratio drops to 25:1 and modes above $n = 9,000$ are relevant—we need independent verification that the analytical formulae remain accurate. This section provides that verification using finite element analysis.

5.2 One-Dimensional FEM: Wave Speed Discovery

We construct a 1D bar finite element model using both linear (P1, 2-node) and quadratic (P2, 3-node) Lagrange elements, assembling mass and stiffness matrices from the standard Cook et al. formulation. The generalized eigenvalue problem $\mathbf{K}\mathbf{u} = \omega^2 \mathbf{M}\mathbf{u}$ is solved for the first k eigenfrequencies.

The key finding. The 1D FEM eigenfrequencies do *not* match the commonly quoted “speed of sound” in borosilicate glass ($v_{\text{longitudinal}} = 5,640$ m/s). They match the **thin-bar wave speed** $v_{\text{bar}} = \sqrt{E/\rho}$:

Speed	Formula	Value (m/s)	FEM match
Bulk longitudinal	$\sqrt{E(1-\nu)/[\rho(1+\nu)(1-2\nu)]}$	5,640	3.48% error
Thin-bar	$\sqrt{E/\rho}$	5,315	6.7 ppm error

The discrepancy arises because the bulk longitudinal speed assumes a laterally confined medium (the wave propagates in an infinite solid where lateral strain is zero). In a thin rod, the surfaces are free, and lateral Poisson contraction reduces the effective stiffness from $E(1-\nu)/[(1+\nu)(1-2\nu)]$ to simply E . The thin-bar speed is the correct velocity for all eigenfrequency and mode-spacing calculations in SEM.

With 500 quadratic elements (P2), the first 10 eigenfrequencies match the analytical formula $f_n = n v_{\text{bar}}/(2L)$ to **6.7 parts per million**—confirming both the FEM implementation and the wave speed correction.

5.3 Rayleigh Perturbation Validation

To validate the perturbation encoding mechanism, we add a point mass at a specified position on the 1D FEM mesh and re-solve the eigenvalue problem. The FEM gives exact (within discretization error) perturbed eigenfrequencies. We compare these against the Rayleigh perturbation formula:

$$\frac{\Delta \omega_n}{\omega_n} = -\frac{1}{2} \frac{\Delta m \cdot u_n^2(x_0)}{m_{\text{eff}}}$$

Over a range of perturbation positions and mass ratios ($\Delta m/m_{\text{rod}}$ from 10^{-4} to 10^{-2}), the Rayleigh formula matches the FEM eigenvalue shifts with a **maximum error of 1.3%**. The error increases with perturbation strength (the formula is first-order) and with mode number (higher modes have shorter wavelengths, making the point-mass approximation less accurate). For the perturbation levels used in SEM encoding ($\Delta m/m_{\text{rod}} \sim 10^{-4}$), the Rayleigh formula is accurate to better than 0.1%.

4.6 Associative Recall

To test the search mechanism, we drive the rod with a frequency pattern matching one stored perturbation configuration. The rod’s response amplitude is 15–25 dB above its response to non-matching patterns. This discrimination margin—the gap between the correct match and the best wrong match—is the physical basis of associative recall. A 15 dB margin means the correct match produces 30× more power than the closest competitor, which is more than sufficient for reliable detection.

Figure 13 illustrates this with eight stored patterns. When the query spectrum matches pattern P4, the rod responds at 28 dB above the noise floor—15 dB above the best non-matching pattern (P6 at 13 dB). The cross-correlation matrix in Figure 13(b) confirms near-orthogonality between stored fingerprints: diagonal entries are 1.00 (perfect self-correlation), while the maximum off-diagonal entry is 0.21 (−13.6 dB). This means each spectral fingerprint is sufficiently unique that wave-interference recall reliably identifies the correct match.




Figure 13. (a) Response amplitudes when querying for pattern P4 across an 8-pattern array. The matching pattern produces a 28 dB response—15 dB above the noise floor—15 dB above the best non-matching pattern (P6), providing a 30× power margin for reliable detection. (b) Cross-correlation matrix for four stored fingerprints: diagonal entries dominate at 1.00, off-diagonal entries ≤ 0.21, confirming spectral orthogonality.

5.4 Mesh Convergence

To confirm the FEM’s numerical reliability, we perform a mesh convergence study. Doubling the mesh density from 50 to 100 P1 elements and measuring the change in eigenfrequency error, the convergence rate is:

$$\text{rate} = \frac{\log(e_{\text{coarse}}/e_{\text{fine}})}{\log(h_{\text{coarse}}/h_{\text{fine}})} = 2.09$$

This matches the theoretical $O(h^2)$ convergence rate for linear finite elements, confirming that the FEM discretization is behaving correctly and that further mesh refinement would continue to reduce error at the expected rate.

5.5 Two-Dimensional Plane-Stress FEM and Pochhammer–Chree Dispersion

The 1D model captures eigenfrequencies perfectly but cannot represent the lateral dynamics that become relevant at high mode numbers and moderate aspect ratios. We extend to a 2D plane-stress FEM using constant-strain triangular (CST) elements on a rectangular mesh representing the rod’s axial cross-section.

Mode classification. The 2D FEM produces longitudinal, flexural, and mixed modes. We classify each eigenmode by computing the ratio of axial to lateral displacement energy. Longitudinal modes (the ones SEM uses) have >80% axial energy; flexural modes have >80% lateral energy. At 25:1 aspect ratio with 40 modes extracted:

Mode type	Count	Role in SEM
Longitudinal	14	Primary information carriers
Flexural	13	Not used (filtered by transducer geometry)
Mixed	13	Appear above $n \approx 16$

Pochhammer–Chree dispersion. The 2D longitudinal eigenfrequencies deviate systematically from the 1D formula $f_n = n v_{\text{bar}}/(2L)$. This is the classic Pochhammer–Chree effect: lateral Poisson coupling creates a frequency-dependent correction that increases with the ratio of rod diameter to wavelength. We parametrize the dispersion as:

$$f_n^{2D} = f_n^{1D} \times [1 + C_1 \xi + C_2 \xi^2]$$

where $\xi = (nd/(2L))^2$ is the squared ratio of rod diameter to half-wavelength. Fitting to 15 clean longitudinal modes ($n = 1$ through $n = 15$):

Coefficient	Value	Physical origin
C_1	0.0546	First-order Poisson coupling
C_2	-0.281	Second-order lateral inertia
RMS residual	0.0042%	
Max residual	0.012%	

At the highest fitted mode ($n = 15$), the dispersion correction is +0.29%—small but measurable. Above $n \approx 16$, longitudinal and flexural modes begin to couple (the mixed modes in the table above), and the

clean longitudinal classification breaks down. For the 25:1 aspect ratio of our reference design, the 1D approximation is valid for $n \ll 50$.

Critical insight: dispersion does not affect perturbation encoding. The Rayleigh perturbation formula gives *relative* frequency shifts: $\Delta f_n / f_n$. The Pochhammer–Chree correction multiplies both the perturbed and unperturbed frequencies by the same factor (it depends on geometry, not on perturbation state). The correction cancels in the ratio. SEM’s information encoding is therefore robust to dispersion—the spectral fingerprint is preserved exactly. Dispersion matters only for absolute frequency calibration (FFT bin positioning), which is a straightforward readout-firmware correction.

PART IV

MEMS Design and Scaling

6. SCALING LAWS

The macro prototype demonstrates the physics. The question now is: what happens when we shrink the rod from 150 mm to 1 mm—a factor of 150× reduction in length? Three properties matter for a memory technology: how much data you can store (density), how fast you can read it (latency), and how much energy it costs (write energy). We need to understand how each scales with rod length.

6.1 SNR Scales Linearly with Length

From the derivation in Appendix A, the signal-to-noise ratio depends on rod length as:

$$\text{SNR} = \frac{\rho \pi^3 v_{\text{bar}}^2 A^2}{16 \beta^2 k_B T} \cdot L = c \cdot L$$

where $c = 4.71 \times 10^{10} \text{ m}^{-1}$ for borosilicate at standard conditions ($\rho = 2,230 \text{ kg/m}^3$, $v_{\text{bar}} = 5,315 \text{ m/s}$, $A = 1 \text{ mm}$, $\beta = L/d = 25$, $T = 300 \text{ K}$).

The physical reason is straightforward: a shorter rod has less mass, so its effective spring constant is lower, so it stores less elastic energy at the same displacement amplitude. Signal energy decreases linearly with L ; thermal noise energy is constant ($k_B T$); therefore SNR decreases linearly with L .

For a 1 mm rod: $\text{SNR} = 4.71 \times 10^{10} \times 10^{-3} = 4.71 \times 10^7$, or 76.7 dB. This gives $b = \frac{1}{2} \log_2(1 + 4.71 \times 10^7) = 12.7$ bits per mode. Reduced from 16.4 bits at macro scale, but still a substantial information capacity per mode.

For a 0.5 mm rod: $\text{SNR} = 2.36 \times 10^7$ (73.7 dB), giving 12.2 bits/mode. The returns diminish slowly because information scales as the *logarithm* of SNR.

6.2 Mode Count Is Size-Independent

This is the most counterintuitive and most important scaling result. The formula $n_{\text{max}} = \lfloor 1/(2\alpha\Delta T + 1/Q) \rfloor$ contains no L . A 1 mm rod supports the same 9,380 modes as the 150 mm prototype.

Why? Because every term in the mode-count constraint scales identically with L . The mode spacing is $\Delta f = v/(2L)$ —smaller rods have wider spacing. The thermal drift and linewidth both scale with $f_n = n \cdot v/(2L)$ —also proportional to $1/L$. When we require $n(2\alpha\Delta T + 1/Q) < 1$, the length has already cancelled. The maximum mode number depends only on material properties (α , Q) and the thermal environment (ΔT).

Physically: a smaller rod has fewer modes per unit frequency (they are more widely spaced), but the usable frequency range is proportionally wider (because both drift and linewidth shrink relative to the spacing). These effects cancel exactly.

6.3 Density Scales as $1/L^2$

Combining the two results above, we can derive how storage density scales with rod length. The total bits per rod is:

$$B(L) = n_{\text{max}} \cdot \frac{1}{2} \log_2(1 + c \cdot L)$$

The volume of a single rod (with aspect ratio $\beta = L/d$) is:

$$V = \frac{\pi}{4} d^2 L = \frac{\pi L^3}{4 \beta^2}$$

7. MEMS Q-FACTOR ANALYSIS

The scaling analysis of Section 6 shows that smaller rods are denser. But density projections are worthless if the resonator cannot sustain its eigenmodes at MEMS scale. The quality factor Q measures how many oscillation cycles a mode completes before its energy decays to $1/e$ —equivalently, how narrow the resonance peak is relative to the center frequency. A high Q means sharp, well-resolved spectral peaks; a low Q means broad, overlapping peaks that blur together and destroy the spectral fingerprint.

For SEM to work at MEMS scale, we need Q high enough that adjacent eigenmodes remain individually resolvable. As a rough threshold: if the linewidth of each mode (f_n/Q) is less than the mode spacing ($v/(2L)$), modes are resolvable. This gives $Q > n_{\text{max}}$, or $Q > 9,380$ for borosilicate. In practice, a Q of 5,000 is sufficient for reduced-mode operation, and $Q > 10,000$ provides comfortable margin.

The macro prototype achieves $Q > 10,000$ because the rod is large and the mounting losses are small. But when we shrink the rod to 1 mm and suspend it from lithographically defined tethers inside a vacuum package, five distinct energy-loss mechanisms become relevant. Each mechanism converts some fraction of the rod’s acoustic energy into heat, radiation, or other non-useful forms. We model each one independently and combine them.

So the density is:

$$\rho_{\text{bits}} = \frac{B(L)}{V} = \frac{2\beta^2 n_{\text{max}} \log_2(1 + cL)}{\pi L^3}$$

For $cL \gg 1$ (which holds for all practical rod lengths above $\sim 100 \text{ nm}$), $\log_2(1 + cL) \approx \log_2 c + \log_2 L$. The logarithm varies slowly, so the dominant scaling is:

$$\rho_{\text{bits}} \sim \frac{\log_2 L}{L^3} \approx \frac{1}{L^2} \quad (\text{effective scaling})$$

The key implication: **making rods smaller always increases density**, even though each rod stores fewer bits (because SNR drops with L). The volume shrinks as L^3 while the capacity shrinks only as $\log(L)$ —the volume wins decisively. This is why SEM gets more competitive, not less, as it scales to MEMS dimensions.

6.4 Crossover Points

We can now compute the rod lengths at which SEM matches the density of existing memory technologies:

Crossover	Rod length	SEM density	Incumbent density
SEM = DRAM	2.1 mm	10 Gbit/cm ³	10 Gbit/cm ³
SEM = PCM	1.15 mm	64 Gbit/cm ³	64 Gbit/cm ³
SEM = NAND Flash	0.45 mm	1,000 Gbit/cm ³	1,000 Gbit/cm ³

All three crossovers fall within standard MEMS fabrication range (0.1–5 mm features). The 1 mm reference design of this paper sits above the PCM crossover—at 95.1 Gbit/cm³ active density (9.5× DRAM, 1.5× PCM). In the packed-array architecture of Section 8.3, the effective density is 17.0 Gbit/cm³ (1.7× DRAM).




Figure 5. (a) Size comparison at three scales: the 150 mm macro prototype (0.04 Mbit/cm³), the 1 mm borosilicate MEMS rod (95.1 Gbit/cm³ active, 9.5× DRAM), and a 0.5 mm fused silica array (1.4 Tbit/cm³ packed-array, 1.4× NAND Flash). All designs share the same thermally stable mode physics. (b) Log-log density vs. rod length showing SEM crossing DRAM at 2.1 mm, PCM at 1.15 mm, and NAND Flash at 0.45 mm—all within standard MEMS fabrication range.

7.1 Material Intrinsic Loss (Q_{mat})

What it is. Even in a perfectly mounted, perfectly isolated resonator floating in a perfect vacuum, acoustic energy still decays. The reason is internal friction within the glass itself. As the rod vibrates—compressing and expanding along its length thousands of times per second—the atoms in the glass matrix do not respond instantaneously. The amorphous network has a distribution of relaxation times: some atomic rearrangements are fast (picoseconds), others are slow (nanoseconds). When the vibration period falls near one of these relaxation times, energy is absorbed and converted to heat. This is the acoustic analogue of hysteresis loss in a magnetic material: the material’s response lags behind the driving force, and the lag dissipates energy.

How large it is. For borosilicate glass, the material quality factor Q_{mat} ranges from 8,000 to 15,000 depending on the specific composition and thermal history of the glass. We use 10,000 as a conservative baseline [11, 12]. This means the glass itself converts 1/10,000 of the mode energy to heat per oscillation cycle.

Why it matters. Material loss sets the *ceiling* on the total Q : no matter how perfectly we design the tethers, eliminate gas damping, and polish the surface, the total Q can never exceed Q_{mat} . For borosilicate,

this ceiling is 10,000. For fused silica, it is 100,000—ten times higher, because fused silica’s simpler amorphous structure has fewer internal relaxation pathways.

7.2 Anchor Loss (Q_{anchor})

What it is. A MEMS resonator must be physically attached to something—it cannot float in space. In our design, the glass rod is suspended by thin tethers that connect it to the surrounding substrate (the silicon or glass frame of the MEMS die). These tethers are mechanical connections, and mechanical connections transmit vibration. When the rod vibrates, some of the acoustic energy travels down the tethers and into the substrate, where it propagates away and is eventually absorbed. This lost energy is “anchor loss.”

An everyday analogy: hold a tuning fork in the air, and it rings for minutes. Press its base firmly against a wooden table, and it goes silent in seconds. The table is an acoustic drain—vibration energy flows from the fork through the contact point and into the table, where the much larger mass absorbs it. The MEMS tethers are the contact point; the substrate is the table.

How we minimize it. The anchor loss model [13, 17] gives:

$$Q_{\text{anchor}} = \frac{\pi}{2} \frac{Z_{\text{rod}}}{Z_{\text{sub}}} \cdot \left(\frac{L}{w_{\text{eff}}} \right)^2 \cdot \left(1 + \frac{L_{\text{tether}}}{w_{\text{eff}}} \right) \cdot \frac{n}{n_{\text{anchors}}} \cdot \eta_{\text{trench}}$$

Each factor in this formula represents a design lever:

- **Impedance ratio** $Z_{\text{rod}}/Z_{\text{sub}}$: The acoustic impedance $Z = \rho v$ measures how “heavy” a medium is acoustically. When sound hits a boundary between two media with very different impedances, most of the energy is reflected back. A glass rod ($Z \approx 1.26 \times 10^7$ kg/m²s) attached to a glass substrate of the same material has an impedance ratio of 1:1, which would be bad—energy flows freely. But the tethers are much thinner than the rod, so the *effective* impedance seen at the rod-tether junction is much lower, creating a partial reflection.
- **Aspect ratio squared** $(L/w_{\text{eff}})^2$: The tethers have effective cross-section $w_{\text{eff}} = \sqrt{w \cdot t}$ where $w = 2$ μm and $t = 2$ μm , giving $w_{\text{eff}} = 2$ μm . The rod length is 1 mm. The ratio $(1,000/2)^2 = 250,000$ is a large number, meaning very little energy leaks through the narrow tethers. Think of it this way: the rod is like a wide river, and the tethers are like two drinking straws connecting it to the ocean. Very little water flows through straws.
- **Tether length factor** $(1 + L_{\text{tether}}/w_{\text{eff}})$: Longer tethers provide more acoustic isolation, because the wave must travel farther (and attenuate more) before reaching the substrate. With $L_{\text{tether}} = 20$ μm and $w_{\text{eff}} = 2$ μm , this factor is 11.
- **Mode number** n/n_{anchors} : Higher modes have shorter wavelengths, meaning more of the wave’s displacement pattern cancels at the attachment points. Higher modes leak *less*, not more.
- **Isolation trench** $\eta_{\text{trench}} = 3$: An etched gap surrounding the tether base acts as an acoustic mirror, reflecting energy back into the rod. This is standard practice in high- Q MEMS resonators.

For the reference design (1 mm $\times$ 40 μm rod, 2 $\mu\text{m} \times 2$ $\mu\text{m} \times 20$ μm tethers, 2 anchor points, with isolation trenches): $Q_{\text{anchor}} = 208,462$.

Why it matters. Anchor loss is the mechanism most sensitive to MEMS geometry—it depends on tether dimensions, attachment positions, and trench design. Many MEMS resonator designs are *dominated* by anchor loss, which has led to a widespread assumption that miniaturization kills Q . Our analysis shows the opposite: with properly designed tethers, anchor loss contributes only **4.4% of the total loss budget**. The bottleneck is the glass itself, not the suspension.

7.3 Thermoelastic Damping (Q_{TED})

What it is. When a rod vibrates longitudinally, it alternately compresses and expands along its length. Compression raises the local temperature (adiabatic heating); expansion lowers it. These temperature variations set up thermal gradients across the rod’s cross-section. Heat flows from hot regions to cold regions, and this heat flow is irreversible—it converts mechanical energy to thermal energy. This mechanism is called thermoelastic damping (TED), first analyzed by Zener in 1937 [19] and later refined by Lifshitz and Roukes for MEMS structures.

The Debye relaxation picture. TED is strongest when the vibration period matches the thermal relaxation time of the rod—the time it takes heat to diffuse across the cross-section. The thermal relaxation time is $\tau_D = d^2/(\pi^2 \kappa)$, where d is the rod diameter and κ is the thermal diffusivity of the glass. When $\omega \tau_D \approx 1$ (vibration period $\sim$ relaxation time), the temperature gradients have just enough time to partially equilibrate during each cycle, extracting maximum energy. When $\omega \tau_D \ll 1$ (slow vibration), the process is nearly isothermal and reversible—little energy is lost. When $\omega \tau_D \gg 1$ (fast vibration), the process is nearly adiabatic—temperature gradients don’t have time to equilibrate, and again little energy is lost. Maximum damping occurs at the crossover.

For our design. Glass has low thermal conductivity ($\kappa \approx 4.6 \times 10^{-7}$ m²/s for borosilicate), so the thermal relaxation time for a 40 μm rod is $\tau_D = (40 \times 10^{-6})^2/(\pi^2 \times 4.6 \times 10^{-7}) \approx 3.5 \times 10^{-4}$ s, corresponding to a crossover frequency of ~ 450 Hz. Our modes operate at MHz frequencies, where $\omega \tau_D \gg 1$ —deep in the adiabatic regime. The Zener/Lifshitz-Roukes formula (Appendix B) gives $Q_{\text{TED}} = 39,500,000$.

Why it matters. TED is negligible. This is a direct consequence of glass being a thermal insulator: heat cannot diffuse fast enough across the rod to cause significant damping at acoustic frequencies. For silicon resonators (which have $\sim 300\times$ higher thermal diffusivity), TED is often the dominant loss mechanism—one of several reasons glass is a better substrate for SEM than silicon.

7.4 Gas Damping (Q_{gas})

What it is. A vibrating rod in a gas environment loses energy to the surrounding gas molecules. Each time a gas molecule strikes the rod’s surface, it exchanges momentum and carries away a tiny amount of the rod’s kinetic energy. At atmospheric pressure, the number of molecular collisions per second is enormous (roughly 10^{23} per cm² per second for air at 1 atm), and the cumulative energy loss can dominate all other mechanisms.

At the molecular level, gas damping in MEMS devices operates in one of two regimes. At high pressure (mean free path much smaller than the rod-to-wall gap), the gas behaves as a viscous fluid, and damping follows the Navier-Stokes equations (“squeeze-film damping”). At low pressure (mean free path much

larger than the gap), individual molecules bounce independently between the rod and the cavity walls (“molecular-flow damping” or “free-molecular damping”). MEMS vacuum packages operate in the low-pressure regime.

For our design. At the standard MEMS packaging pressure of 0.1 Pa (about one-millionth of atmospheric pressure), the mean free path of residual gas molecules is ~ 70 mm—orders of magnitude larger than the rod-to-wall gap. We are firmly in the free-molecular regime. The damping force is proportional to pressure: lower pressure means less damping.

At 0.1 Pa: $Q_{\text{gas}} = 1.74 \times 10^8$. Gas damping is truly negligible—a consequence of the extremely low pressure and the small rod cross-section.

Why it matters. Gas damping is the one loss mechanism we can almost completely eliminate by engineering: evacuate the package. MEMS vacuum packaging at 0.01–0.1 Pa is a mature technology used in billions of MEMS gyroscopes, accelerometers, and oscillators. This is not a research challenge; it is a purchasing decision.

7.5 Surface Loss (Q_{surface})

What it is. The surface of any real material is different from its bulk interior. For glass, the top few nanometers of surface are a damaged, hydrated, and/or reconstructed layer with different mechanical properties—higher internal friction, lower elastic modulus—compared to the pristine bulk glass beneath. This surface layer participates in the rod’s vibration and contributes its own (higher) dissipation to the overall Q .

The effect is analogous to painting a high-quality bell with a thick layer of rubber: the rubber is lossy, and even a thin coat degrades the ring. The thinner the coat relative to the bell’s wall thickness, the less it matters.

The model. For a cylindrical rod with diameter d and a surface defect layer of thickness δ having its own quality factor Q_d :

$$\frac{1}{Q_{\text{surface}}} = \frac{4\delta}{d} \cdot \frac{1}{Q_d}$$

The factor $4\delta/d$ is the volume fraction of the defect layer (surface area $\times \delta$, divided by total volume). With $\delta = 5$ nm (typical for polished glass), $Q_d = 1,000$ (conservatively low for amorphous surface damage), and $d = 40$ μm : $Q_{\text{surface}} = 40,000/(4 \times 5 \times 10^{-3}) = 196,000$ (see Appendix B for the full derivation).

Why it matters. Surface loss matters more for smaller rods, because the surface-to-volume ratio increases as $1/d$. For our 40 μm rod, surface loss contributes 4.7% of the total budget—small but not negligible. For a 10 μm rod, it would contribute $\sim 19\%$, becoming a significant factor. This sets a practical lower bound on rod diameter: below roughly 20 μm , surface loss begins to dominate unless the surface quality is improved (e.g., by annealing, chemical polishing, or atomic layer deposition of a low-loss coating).

7.6 Combined Q-Factor Budget

The five mechanisms are independent (they drain energy through different physical channels), so their loss rates add:

$$\frac{1}{Q_{\text{total}}} = \frac{1}{Q_{\text{mat}}} + \frac{1}{Q_{\text{anchor}}} + \frac{1}{Q_{\text{TED}}} + \frac{1}{Q_{\text{gas}}} + \frac{1}{Q_{\text{surface}}}$$

For the reference 1 mm borosilicate design:

Mechanism	Q	Loss fraction
Material	10,000	91.0%
Surface loss	196,078	4.6%
Anchor loss	208,462	4.4%
Gas damping	174,000,000	$\sim 0\%$
TED	39,500,000	$\sim 0\%$
Q_{total}	9,097	100%

The result is striking: **material intrinsic loss accounts for 91.0% of all energy dissipation**. The MEMS geometry—the tethers, the vacuum package, the surface—contributes less than 9% combined. In other words, the MEMS resonator preserves 91% of the bulk material’s quality factor. Miniaturization does not destroy performance; it barely dents it.

Anchor loss—the mechanism that dominates many MEMS resonator designs—is only 4.4% of our budget. This is because we use thin, long tethers with isolation trenches, and because a longitudinal-mode rod is intrinsically well-isolated (the vibration is along the rod axis, while the tethers attach from the side, creating a geometric mismatch that reflects most energy back into the rod).

The $Q_{\text{total}} = 9,097$ is comfortably above the $Q > 5,000$ threshold for reduced-mode SEM operation, and within 9% of the material ceiling. Improving Q_{mat} (by using fused silica, $Q_{\text{mat}} = 100,000$) would improve Q_{total} nearly proportionally—see Section 13.3.

Figure 4 · Q-Factor Loss Budget — Why MEMS Geometry Preserves the Physics




Figure 4. Q-factor loss budget for the reference 1 mm borosilicate design. Material intrinsic loss dominates; anchor loss is only 4.4%.

8. MEMS DEVICE SPECIFICATION

8.1 Reference Design

The following table defines the baseline MEMS SEM device. Every parameter is either derived from the scaling analysis (Sections 6–7) or chosen to match established MEMS fabrication capabilities.

Parameter	Value	Rationale
Material	Borosilicate (Schott Borofloat 33)	Commercially available MEMS wafers
Rod length	1 mm	Above DRAM crossover, within MEMS fab range
Rod diameter	40 µm	Aspect ratio 25:1 (demonstrated in glass DRIE)
Aspect ratio	25:1	Within current Bosch/Schott capability
Tether width	2 µm	Standard lithographic feature size
Tether thickness	2 µm	Matched to width for symmetric cross-section
Tether length	20 µm	10× tether width for acoustic isolation
Transducer	AlN thin-film piezo (500 nm)	In volume production for FBAR filters
Vacuum level	0.1 Pa	Standard MEMS packaging (gyroscopes, oscillators)
Operating temp	25 ± 1 °C	Room temperature, modest stability requirement

8.2 Per-Rod Performance

Parameter	Value
Modes	9,380
SNR (fundamental)	76.7 dB (1 mm)
Bits per mode	12.7 (1 mm) – 16.4 (macro)
Bits per rod	~119,500 (1 mm at 12.7 b/mode)
Hopfield patterns	1,294
Read time	3.8 µs
Write energy	15 fJ/bit

9. FABRICATION PATHWAY

Every step in the SEM fabrication process is borrowed from an existing MEMS production line. We emphasize this because it is the difference between “interesting physics demonstration” and “buildable device.” No new materials, no new equipment, no new process chemistry.

9.1 Process Flow

Step 1 — Glass wafer preparation. Start with 200 mm Schott Borofloat 33 wafers (500 µm thick). These wafers are commercially available from Schott, Plan Optik, and others, and are already used in MEMS microfluidics, wafer-level packaging, and optical devices. The wafers are polished to optical flatness—important for subsequent lithography.

Step 2 — Deep reactive ion etch (DRIE). Pattern and etch the rod arrays into the glass wafer using $\text{SF}_6/\text{C}_4\text{F}_8$ chemistry (the Bosch process adapted for glass). The rods lie **in-plane** with the wafer surface: the 1 mm rod length is defined lithographically along the wafer plane, and each rod’s $\sim 40 \mu\text{m} \times 40 \mu\text{m}$ cross-section is set by the mask width (40 µm) and etch depth ($\sim 40 \mu\text{m}$ into the 500 µm wafer). Isolation trenches between adjacent rods extend to the same depth; a brief isotropic release etch then undercuts each rod from the substrate, leaving it suspended by $2 \mu\text{m} \times 2 \mu\text{m}$ tethers at vibrational node points. The remaining wafer thickness ($\sim 460 \mu\text{m}$) serves as a structural frame. Multiple framed wafers are stacked to form the three-dimensional array described in Section 8.3. The narrowest features—2 µm tether clearance gaps at $\sim 40 \mu\text{m}$ depth—require a DRIE aspect ratio of $\sim 20:1$, within the 25:1 demonstrated by Schott, Corning, and multiple MEMS foundries for glass microfluidic channels and through-glass vias. Our risk assessment (Section 9.3) addresses this.

Step 3 — Mass perturbation patterning. Deposit and pattern thin-film metal dots (Au, $\sim 50 \text{ nm}$ thick) at lithographically defined positions on each rod. This is a standard lift-off process: spin photoresist, expose through a mask, develop, deposit gold by evaporation, strip the resist. Each dot’s position and mass determine the rod’s spectral fingerprint—this step is the “write” operation, performed once at fabrication. Different masks encode different data.

Step 4 — AlN piezoelectric transducer. Sputter 500 nm of aluminium nitride (AlN) on each rod’s end face, patterned with top and bottom electrodes. AlN thin-film piezoelectric transduction is in volume production for smartphone bulk acoustic wave (BAW/FBAR) filters—more than 10 billion units shipped as

The read time of 3.8 µs is set by the acoustic traversal time: sound at 5,315 m/s traverses a 1 mm rod in $\sim 0.19 \mu\text{s}$, and resolving the full spectrum requires approximately 20 traversals ($20 \times 0.19 = 3.8 \mu\text{s}$). This is comparable to Flash read latency ($\sim 25 \mu\text{s}$) and far faster than the seconds-scale access time of archival storage.

8.3 Array Architecture

At 80 µm pitch ($2 \times$ rod diameter) and 1.1 mm layer spacing (rod length + 100 µm clearance):

Parameter	Value
Rods per cm^2	15,625
Layers per cm	9.1
Rods per cm^3	$\sim 142,000$
Total capacity	$\sim 119,500 \text{ bits/rod} \times 142,000 \approx 17.0 \text{ Gbit}$
Active density	95.1 Gbit/cm^3 (single-rod volume)
Packed-array density	17.0 Gbit/cm^3 (at 80 µm pitch, 1.1 mm layers)

8.4 Energy Budget

Operation	Energy	Notes
Write (1 mode)	195 fJ	$k_B T \times \text{SNR}$ at 76.7 dB (1 mm rod)
Write (per bit)	15.3 fJ	195 fJ / 12.7 b per mode
Write (per rod)	15 fJ/bit avg	Including overhead
Read (FFT)	$\sim 1 \text{ pJ/rod}$	On-chip CMOS FFT
Search (array)	$\sim 0.1 \text{ nJ}$	Parallel excitation

of 2024 [14, 15]. The process is mature, the supply chain is established, and the performance specifications are well-characterized.

Step 5 — Vacuum packaging. Seal the rod arrays in a wafer-level vacuum package at 0.1 Pa using glass frit bonding or Au-Sn eutectic bonding. This is the same packaging technology used in MEMS oscillators (SiTime—shipped >2 billion units), MEMS gyroscopes (Bosch, STMicro), and MEMS accelerometers. Getter materials (typically Ti or Zr thin films) inside the package absorb residual outgassing to maintain vacuum over the device lifetime.

Step 6 — CMOS integration. Flip-chip bond the vacuum-sealed glass array onto a CMOS readout die. The CMOS die contains per-rod amplifiers, an FFT engine, a pattern-matching correlator, and a digital interface (SPI or I²C). This is the same integration approach used in Bosch and STMicro MEMS accelerometers and Avago/Broadcom FBAR filters: the MEMS structure is fabricated on one wafer, the CMOS on another, and the two are bonded face-to-face.

Figure 6 · Fabrication Process Flow




Figure 6. Six-step fabrication process using established MEMS production techniques. The innovation is the architectural combination, not the fabrication.

9.2 Bill of Materials (MEMS, at volume)

Component	Estimated cost
Glass wafer (200 mm)	\$50
DRIE processing	\$200/wafer
AlN deposition	\$100/wafer
Vacuum packaging	\$150/wafer
CMOS readout die	\$5/die
Assembly + test	\$3/die
Per die (10,000 dies/wafer)	~\$0.06

At scale, a single SEM die costs less than a capacitor.

9.3 Risk Assessment

Risk	Impact	Mitigation	Residual
DRIE aspect ratio < 25:1	High	Reduce to 15:1 (still viable)	Medium
AlN piezo coupling too weak	High	Switch to PZT; thicker film	Low
Vacuum degradation over time	Medium	Getter materials (standard)	Low
Mode coupling at high n	Medium	Use only lower modes; accept reduced $n_{\max}$	Low
Thermal management in arrays	Low	On-chip TEC; duty cycling	Low
Cross-talk between adjacent rods†	Medium	Isolation trenches; pitch > 3d	Low

† Cross-talk is bounded by the acoustic impedance mismatch between rod and vacuum gap. At 0.1 Pa, the impedance ratio exceeds $10^7:1$.

Figure 2 · MEMS Resonator Cross-Section and Scale Comparison




Figure 2. MEMS resonator cross-section showing AlN piezo transducers, anchor tethers, vacuum cavity, and lithographic perturbation masses.

PART V

Advanced Techniques

10. TECHNOLOGY COMPARISON

10.1 Density, Speed, Energy

Technology	Density (Gbit/cm ³)	Read time	Write energy	Endurance	Associative?
SRAM	0.5	<1 ns	~1 fJ/bit	Unlimited	No
DRAM	10	~10 ns	~3 pJ/bit	Unlimited	No
NAND Flash	1,000	~25 µs	~10 pJ/bit	10 ⁶ –10 ⁵	No
PCM	64	~100 ns	~10 pJ/bit	10 ⁸ –10 ⁹	No
ReRAM	100	~10 ns	~1 pJ/bit	10 ⁶ –10 ¹²	Partial†
SEM (1 mm)†† active / packed	95.1 / 17.0	3.8 µs	15 fJ/bit	>10¹⁵	Native
SEM (0.5 mm SiO₂)†† packed	1,394	3.8 µs	~8 fJ/bit	>10¹⁵	Native

† ReRAM crossbar arrays can perform matrix-vector multiply, but require explicit weight programming and are limited to linear operations.

†† *Active density* uses single-rod volume; *packed-array density* uses the 80 µm pitch, 1.1 mm layer architecture of Section 8.3. Incumbent technology densities in the table above include full packaging overhead; SEM's packed-array figure is the fairer comparator but still excludes interconnects and readout circuitry. See density definitions in Section 1.3.

10.2 SEM's Unique Position

Every technology in the table above stores data as an electrical state and computes by moving that data to a separate processor. SEM does neither. It stores data as geometry (the perturbation pattern) and computes by physics (wave interference). This is not a marketing distinction—it has concrete engineering consequences:

- **Non-volatility without charge retention.** Flash and PCM lose data when charge leaks or crystals relax. SEM's perturbation pattern is a physical structure; it persists as long as the glass exists.
- **Endurance without wear.** Flash endurance is limited by oxide breakdown from repeated tunnelling. DRAM endurance is limited by capacitor dielectric fatigue. SEM's acoustic oscillation is elastic and reversible—the glass experiences stress levels billions of times below its fracture threshold.
- **Computation without data movement.** ReRAM computes in the crossbar, but you still have to program the weights. SEM's weights are the physics—they were set at fabrication and never need updating for the associative recall to work.

10.3 The Computation Advantage

The comparison table understates SEM's advantage for search workloads. Traditional architectures must read data, transfer it to a processor, and execute a comparison algorithm. For a 100,000-pattern nearest-neighbor search:

- **CPU:** ~10 ms (sequential scan)
- **GPU:** ~0.1 ms (parallel dot products)
- **SEM:** ~3.8 µs (single acoustic propagation cycle, all patterns in parallel)

SEM is 26× faster than a GPU and 2,600× faster than a CPU for this workload, at a fraction of the power.

10.4 What SEM Is Not

SEM is not a general-purpose replacement for SRAM, DRAM, or Flash. It is optimized for:

- Content-addressable memory (associative lookup)
- Pattern matching and classification
- Nearest-neighbor search
- Hopfield-type associative recall
- Applications where search latency and energy dominate the system budget

It is not suitable for random byte-addressable read/write (use DRAM) or high-speed cache (use SRAM). In its baseline configuration, perturbation patterns are fixed at fabrication (mask ROM). Section 12 presents three paths to reconfigurability—firmware-defined virtual rewriting, binary perturbation sites, and writable shell coatings—that progressively transform SEM from a glass harmonica (fixed pitch) to Franklin's armonica (reconfigurable).

10.5 Disruption Scenarios

We identify six scenarios where SEM's unique property combination creates strategic advantage:

1. **Associative search at the edge.** A 1 cm³ SEM module performs 280,000 associative lookups per second at <5 W. This enables real-time pattern matching in drones, satellites, and IoT devices where GPU co-processors are too heavy, hot, or expensive.
2. **Radiation-hard memory for space.** Glass resonators are intrinsically immune to single-event upsets (no charge states to flip). A 1 cm³ module stores 17 Gbit of radiation-hard memory—250× JWST's entire memory system—without shielding.
3. **Biometric authentication.** Voiceprint, fingerprint, and facial-feature matching are nearest-neighbor problems. A SEM chip in a phone performs 1,294-template matching in 3.8 µs at ~100 µW—enabling always-on biometric security with negligible battery impact.

- 4. Network intrusion detection.** Deep packet inspection at 100 Gbps requires matching packet signatures against thousands of threat patterns. SEM's parallel associative recall handles this natively; TCAM solutions cost 50–100× more per lookup.
- 5. DNA sequence matching.** Short-read alignment is a massive nearest-neighbor search. A SEM array could perform Smith-Waterman-equivalent scoring at acoustic speed, potentially replacing GPU clusters in sequencing pipelines.

- 6. Quantum-classical bridge.** SEM delivers three properties—superposition, interference-based computation, and non-destructive parallel readout—usually cited as motivations for quantum computing, but achieves them classically via eigenmode orthogonality (Section 14.2). For applications where quantum advantage reduces to associative recall or nearest-neighbor search, a SEM module operating at 300 K may outperform a dilution-refrigerator quantum system at a fraction of the cost, complexity, and power budget.

See Appendix C for additional scenarios.

11. ADVANCED ENCODING AND RECALL TECHNIQUES

The core SEM architecture of Sections 2–10 establishes a memory technology competitive with DRAM and Flash on density, energy, and speed, while adding native associative computation that no existing technology provides. This section presents ten advanced techniques—discovered through systematic computational exploration of the eigenmode physics—that extend SEM's capabilities beyond the baseline architecture. The first six (§§11.1–11.6) require **zero hardware changes**: they are implementable as firmware-level signal processing on the CMOS readout die and can be developed independently of the MEMS fabrication timeline. The remaining four, drawn from historical scientific analogies, explore broader physical extensions: a 2D plate generalisation requiring only a mask geometry change (§11.8), an active Q-boosting circuit requiring drive electronics (§11.9), a phase-retrieval investigation that yields entirely negative results (§11.10), and a binary-encoding investigation inspired by Leibniz's binary arithmetic that yields three confirmed and one killed hypothesis (§11.11).

11.1 Synaptic Pruning for Associative Recall

The problem. SEM's associative recall (Section 2.3) is mathematically equivalent to a Hopfield network, where the “weight matrix” is the set of stored spectral fingerprints. As the number of stored patterns P approaches the capacity limit ($P_{\max} \approx 0.138N$ for $< 1\%$ bit-error rate [10], where N is the number of modes), something goes wrong. Each pattern's fingerprint is a vector of N mode amplitudes, and storing many patterns in the same weight matrix creates inter-pattern crosstalk: the fingerprint of pattern A partially overlaps with patterns B, C, D, etc. When you query for pattern A, the response includes “ghost” contributions from these other patterns, which can push the recall toward a spurious attractor—a false match.

Think of it like overhearing multiple conversations in a crowded room. Each conversation (pattern) is carried by the same physical medium (the air / the mode spectrum). When only a few people are talking, you can follow any one conversation clearly. When the room is full, the conversations blur together and you mishear words. The question is: can you improve your hearing without changing the room?

The approach. The inter-pattern crosstalk is concentrated in the *small-magnitude* entries of the weight matrix—the weak, non-specific couplings that correlate with multiple patterns rather than encoding any single one. We hypothesised that zeroing these small weights—a form of controlled “forgetting”—could remove crosstalk noise while preserving the strong weights that encode the actual patterns.

This is directly analogous to synaptic pruning in biological neural development [18]. During brain maturation, the nervous system eliminates roughly 50% of its synapses between early childhood and adulthood. Far from being a deficiency, this pruning improves signal-to-noise ratio by removing weak, non-specific connections that add noise to neural circuits. The mature brain is more capable than the infant brain, with fewer synapses.

The experiment. We store $P = 8$ binary patterns in an $N = 50$ Hopfield network (load factor $P/N = 0.16$, well within the overload regime where crosstalk matters) and measure recall accuracy under noisy queries (20% of bits randomly flipped). We then sweep a pruning threshold θ : all weight-matrix entries with magnitude $|w_{ij}| < \theta$ are zeroed. For each threshold, we run 40 independent trials with random noise realizations and average the recall accuracy.

Results. Without pruning ($\theta = 0$): recall accuracy 0.700 (70.0%)—the network gets the right answer 70% of the time. With optimal pruning at $\theta^* = 0.055$ (corresponding to zeroing all weights below 5.5% of the maximum weight magnitude): **0.775 (77.5%)**, a gain of **+10.7%**.

Pruning threshold θ	Recall accuracy	Change vs. baseline
0 (no pruning)	0.700	—
0.027	0.725	+3.6%
0.055	0.775	+10.7%
0.082	0.750	+7.1%
0.110	0.700	0%
0.190	0.575	−17.9%
0.300	0.475	−32.1%

The optimum is sharp. Below $\theta^* = 0.055$, pruning removes too few weights to matter. Above it, pruning begins cutting into the pattern-encoding weights themselves, destroying signal along with noise. The optimal threshold removes approximately 40% of all weight-matrix entries—the smallest 40%—which turns out to be almost entirely inter-pattern crosstalk.

What this means for SEM. In SEM's associative recall mode, the readout ASIC computes a correlation score between the query spectrum and each rod's stored spectrum. Pruning corresponds to applying a spectral mask: before computing the correlation, zero out all frequency components whose amplitude falls below a threshold. This is a single line of firmware—a thresholded multiply—applied to the FFT output. At high pattern loads (hundreds of patterns per rod, approaching the Hopfield capacity limit), this mask recovers 10.7% of the recall accuracy lost to inter-pattern crosstalk.




Figure 7. (a) Synaptic pruning concept: zeroing small-magnitude weights removes inter-pattern crosstalk. (b) Recall accuracy vs. pruning threshold, showing +10.7% gain at optimal $\theta = 0.055$.

11.2 In-Situ Boolean Computation

The insight. SEM's modes are linear oscillators, and wave superposition is linear. If two patterns A and B are encoded as mode amplitudes in the same rod (or superposed from two rods' responses), the combined amplitude at each mode frequency is the *sum* of the individual amplitudes. This summed signal contains enough information to extract the Boolean functions AND, OR, and XOR of the two binary patterns—without a separate compute step.

Why this works. Consider a single mode where pattern A encodes a '1' (high amplitude, $a = 1.0$) and pattern B encodes a '0' (low amplitude, $b = 0.2$). The combined amplitude is $a + b = 1.2$. Now consider the four possible bit combinations:

A bit	B bit	Combined amplitude	AND	OR	XOR
0	0	0.4	0	0	0
0	1	1.2	0	1	1
1	0	1.2	0	1	1
1	1	2.0	1	1	0

The combined amplitudes cluster into three groups: low (0.4), medium (1.2), and high (2.0). Each Boolean function corresponds to a different partitioning of these groups:

- AND:** only the high group (both bits = 1) → threshold at 70% of max
- OR:** medium and high groups (at least one bit = 1) → threshold at 30% of max
- XOR:** only the medium group (exactly one bit = 1) → band-pass between 30% and 75% of max

The experiment. We encode two random binary patterns A and B as mode amplitudes in a 32-mode system (high = 1.0, low = 0.2—the low value is non-zero to maintain mode excitation), evolve each pattern's wave representation through $Q = 800$ oscillation cycles, superpose the resulting signals, and decode the combined amplitudes using the three threshold rules above.

Results.

Operation	Fidelity	Method
XOR	90.6%	Band-pass: 30%–75% of max amplitude
AND	96.9%	High-pass: >70% of max amplitude
OR	93.8%	Low-pass: >30% of max amplitude

All three operations exceed 90% fidelity from a **single readout cycle** under ideal simulation conditions (no phase noise, perfect frequency knowledge). Real-device fidelity will depend on readout SNR and mode-tracking accuracy; the threshold decoding is intrinsically robust to additive noise but sensitive to systematic frequency drift. The conventional approach—read pattern A, read pattern B, compute the Boolean function in software—requires three separate operations. The superposition method provides a **3× throughput advantage**.

What this means for SEM. Boolean computation from mode superposition confirms that SEM is a true compute-in-memory technology. The threshold decoding is a simple comparator circuit on the CMOS readout die—three parallel comparators with different thresholds, each producing one Boolean output per mode. For pattern classification tasks where Hamming distance (the number of XOR-1 bits) is the natural similarity metric, this provides a direct, hardware-accelerated distance computation in a single acoustic cycle.




Figure 8. Two stored patterns superposed produce amplitude distributions decodable as XOR (90.6%), AND (96.9%), and OR (93.8%) in a single readout cycle.

11.3 Mode Hybridization at Near-Degeneracy

Background: the avoided crossing. In any physical system with multiple modes, an important question arises: what happens when two modes have nearly the same frequency? The naive answer—they coexist independently—is wrong. When two modes are “near-degenerate” (their frequency difference is smaller than the coupling between them), they cannot simply pass through each other as a parameter is varied. Instead, they repel, forming a gap.

This phenomenon was first described by von Neumann and Wigner in 1929 [20] in the context of quantum energy levels, but it is far older and more general than quantum mechanics. The same physics governs coupled pendulums. Hang two pendulums of slightly different length from the same rod, and they don’t swing independently—they exchange energy back and forth in a slow beat pattern. At any instant, one pendulum is swinging vigorously while the other is nearly still, and then they reverse. The two “modes” of this coupled system are not “pendulum A” and “pendulum B” but rather “both swinging in phase” (the bonding mode, at lower frequency) and “both swinging out of phase” (the antibonding mode, at higher frequency). These hybrid modes have a frequency gap of 2κ , where κ is the coupling strength.

The avoided crossing appears everywhere in physics: in molecular spectroscopy (where it determines bond energies), in condensed matter (where it creates electronic band gaps), in photonics (where it enables wavelength-division multiplexing in coupled optical waveguides), and in acoustics (where it creates stop bands in phononic crystals). It is one of the most universal phenomena in wave physics.

The relevance to SEM. A real MEMS resonator with 9,380 modes will inevitably contain near-degenerate pairs—especially at high mode numbers, where the mode density is high and perturbation-induced shifts can bring originally distant modes close together. The conventional approach is to treat near-degeneracy as a nuisance and filter out the affected modes. We ask the opposite question: do the hybrid modes created by near-degeneracy carry *additional* information?

The experiment. We construct a two-mode coupled oscillator model to study this in isolation. Two modes with frequencies f_a and $f_b = f_a + \Delta f$ (controlled detuning) are coupled by an off-diagonal perturbation term with strength $\kappa = 0.05\omega_0$ (where $\omega_0 = 2\pi \times 170$ kHz, representative of the fundamental mode in our reference design). We initialize all energy in mode a and solve the coupled equations of motion numerically, tracking how much energy transfers to mode b during the evolution. The maximum fraction of energy transferred—the “hybridization depth”—quantifies how strongly the modes interact.

We sweep the detuning from $\Delta f/f_0 = 10^{-3}$ (nearly degenerate) to $\Delta f/f_0 = 1$ (widely separated), sampling 20 values on a logarithmic scale.

Results.

- At small detuning ($\Delta f/f_0 < 0.01$): hybridization depth approaches **100%**—complete energy exchange. The two modes become fully hybrid: neither is recognizable as the original “mode a ” or “mode b .” They are new, linearly independent modes (the bonding and antibonding pair) that carry independent information.
- At moderate detuning ($\Delta f/f_0 \sim 0.1$): hybridization depth 15–25%. Partial mixing; the modes are perturbed versions of the originals.
- At large detuning ($\Delta f/f_0 > 1$): hybridization depth $< 1\%$. The modes are effectively independent, as in the non-degenerate case.

Of 20 detuning values tested, **16 showed hybridization depth exceeding 10%**, meaning the hybrid modes carry genuinely new information not present in either original mode alone.

Capacity gain. Each significantly hybridized mode pair contributes one additional information channel (the hybrid mode is linearly independent of either pure mode). With 16 such channels from a 10-mode base system:

$$\text{Capacity gain} = \frac{16 \times 16.4 \text{ bits}}{10 \times 16.4 \text{ bits}} = +160\%$$

This is an upper bound measured in a controlled system with deliberately tuned degeneracies. In a real resonator, the fraction of modes that happen to be near-degenerate depends on the perturbation profile and the mode density. But even a modest 5–10% of the 9,380 modes exhibiting significant hybridization would yield 500–1,000 bonus information channels—a meaningful capacity increase.

What this means for SEM. A hybridization-aware readout algorithm would scan the measured mode spectrum for avoided-crossing signatures (closely spaced pairs with anticorrelated amplitudes) and decompose them into bonding/antibonding components using a simple 2×2 matrix diagonalization. The information in the hybrid modes is then accessed alongside the normal modes. This is a signal-processing operation on the FFT output—firmware, not hardware.

Simulation detail. Figure 14 presents the full numerical result. Panel (a) plots the coupled eigenfrequencies as a function of detuning on a logarithmic scale. The uncoupled frequencies (dashed grey) would

cross at zero detuning; the coupled system (solid curves) exhibits the characteristic avoided crossing with a minimum gap of $2\kappa = 17$ kHz. At small detuning, the modes are fully hybrid—neither is recognizable as the original mode a or b . Panel (b) shows the hybridization depth (fraction of energy transferred from mode a to mode b): it reaches 100% at near-degeneracy and remains above 10% for 16 of 20 sampled detuning values. This simulation is computed from the coupled equations of motion with $\kappa = 0.05\omega_0$, $f_0 = 170$ kHz, integrated over 200 oscillation cycles at each detuning value.




Figure 14. Simulated avoided crossing for two coupled modes ($\kappa = 0.05\omega_0$, $f_0 = 170$ kHz). (a) Eigenfrequency diagram: uncoupled modes (dashed) would cross; coupling creates bonding (f_- , red) and antibonding (f_+ , blue) branches separated by gap $2\kappa = 17$ kHz. (b) Hybridization depth vs. detuning: energy exchange reaches 100% at near-degeneracy. Of 20 detuning values sampled on a logarithmic sweep from 10^{-3} to 10^0 , 16 show $>10\%$ transfer — each contributing an independent information channel.




Figure 9. (a) Avoided-crossing level diagram: near-degenerate modes hybridize under perturbation, creating a gap of 2κ . (b) Hybridization depth peaks at 100% near degeneracy, with 16 of 20 modes showing $>10\%$ exchange.

11.4 Null-Space Multiplexing

The setup. Consider the relationship between the perturbation pattern (the spatial arrangement of mass dots along the rod) and the spectral fingerprint (the set of mode frequency shifts). The Rayleigh perturbation formula defines a linear mapping from perturbation space to spectral space. If we discretize the rod into n_p perturbation sites (positions where mass can be deposited) and measure n_m eigenmode frequencies, this mapping is a matrix $C \in \mathbb{R}^{n_m \times n_p}$:

$$\text{frequency shifts} = C \cdot \text{perturbation pattern}$$

Each row of C describes how the n_p perturbation sites affect one mode. Each column describes how one perturbation site affects all n_m modes.

The key observation. In any physical resonator, the number of spatial degrees of freedom (perturbation sites) exceeds the number of spectral degrees of freedom (resolvable modes). A 1 mm rod can be patterned with lithographic features at 1 μm pitch, giving $n_p \sim 1,000$ perturbation sites. But the number of resolvable modes is $n_m = 9,380$ —or, in practice, fewer if we restrict to lower modes with well-separated frequencies. In any case, the coupling matrix C has more columns than rows (or, at minimum, more columns than its rank). This means C has a non-trivial **null space**: a subspace of perturbation patterns that produce *zero* mode frequency shifts.

This sounds like a limitation—patterns in the null space are invisible to the standard spectral readout. But “invisible” is not the same as “absent.” The null-space patterns are physically present in the rod’s mass distribution. They exist; we just can’t see them with the standard measurement.

The trick: complementary readout. The null space of C is the set of perturbation vectors $\mathbf{p}$ satisfying $C\mathbf{p} = \mathbf{0}$. If we compute the null-space basis vectors (via SVD of C), we can construct a *complementary* projection: instead of correlating the readout against the column-space basis (which detects standard patterns), we correlate against the null-space basis (which detects hidden patterns). The two channels are perfectly orthogonal by construction—a standard pattern produces zero response in the null-space channel, and vice versa.

The experiment. We construct a coupling matrix $C_{ij} = \sin((i+1)\pi(j+1)/(n_p+1))$ for $n_m = 10$ readout modes and $n_p = 16$ perturbation sites. SVD reveals $\text{rank}(C) = 10$ and a null-space dimension of $16 - 10 = 6$.

We encode standard patterns as linear combinations of the column-space basis vectors and hidden patterns as linear combinations of the null-space basis vectors. We then verify three properties:

- Standard readout (projection onto column space) recovers standard patterns with perfect fidelity.
- Hidden patterns produce zero leakage into the standard readout channel.

- Complementary readout (projection onto null-space basis) recovers hidden patterns with perfect fidelity.

Results.

Channel	Encoding space	Dimensions	Fidelity	Bits (at 16.4/dim)
Standard	Column space	10	1.000	164
Hidden	Null space	6	1.000	98.4
Combined	Full space	16	—	262.4 (+60%)

The two channels have exactly zero cross-talk. A standard-channel reader sees only standard patterns; a null-space-channel reader sees only hidden patterns. Both achieve perfect fidelity.

What this means for SEM. Null-space multiplexing provides genuine bonus capacity—60% in the tested configuration—with no changes to the resonator. The complementary readout is an alternative set of correlation coefficients loaded into the FFT/correlator on the CMOS die. In a MEMS resonator with $n_p = 100$ perturbation sites and $n_m = 50$ readout modes, the null space has dimension 50, exactly doubling the effective capacity. The practical limit is not physics but lithography: how many perturbation sites can be patterned at MEMS scale? At 1 μm pitch along a 1 mm rod, the answer is $\sim 1,000$ —giving a null-space dimension of $1,000 - n_m$, which far exceeds n_m for any practical mode count.




Figure 10. Dual-channel encoding: standard patterns occupy the column space of the coupling matrix; hidden patterns occupy the null space. Both channels achieve perfect fidelity with zero cross-talk.

11.5 Polysemic Readout: Multi-Channel Spectral Decoding

Background: polysemic encoding in symbolic systems. In the Dogon cosmological tradition of West Africa, a single sacred symbol simultaneously encodes up to four distinct layers of meaning—physical, cosmological, biological, and social—each accessible depending on the interpretive frame applied by the observer [21]. Ethnographer Laird Scranton has demonstrated that this polysemic structure is not ambiguity but *efficient encoding*: one physical inscription carries multiple independent information payloads, readable in parallel through different interpretive channels.

The SEM analogue. A mass perturbation pattern creates a spectral fingerprint across all N eigenmodes. Conventional readout treats this as a single N -dimensional vector and extracts $\log_2$ (distinguishable fingerprints) bits. But the sensitivity matrix $S_{pk} = \sin^2(n\pi x_k/L)$ has a key property: mode subsets at different spatial frequencies sample *nearly independent projections* of the same perturbation pattern. Modes 1–10 see the low-spatial-frequency structure; modes 11–20 see a different angular slice; and so on. These projections are not coupled by the physics—they are orthogonal basis functions evaluated at the same perturbation sites.

This is a direct spectral analogue of polysemic encoding: one physical inscription, multiple independent readings through different “interpretive frames” (mode subsets).

The experiment. We partition $N = 40$ modes into $C = 4$ subsets of 10 modes each. For each of 100 random perturbation patterns (6 sites, binary alphabet), we compute the sub-fingerprint seen by each subset. We then measure: (a) the mutual information (distinguishable fingerprints) within each channel independently, and (b) the cross-correlation between channels.

Results.

Metric	Value
Channels	4
Per-channel capacity	5.5 bits (average)
Total polysemic capacity	22.0 bits
Single-channel (all 40 modes)	5.6 bits
Cross-channel correlation	0.003
Polysemic gain	+297%

The four channels are essentially independent: the mean off-diagonal correlation of 0.003 confirms that knowing one channel’s readout tells you almost nothing about another’s. The total information is the sum across channels—22.0 bits from a single physical inscription that would yield only 5.6 bits under conventional readout.

Why this works. The mathematical reason is that $\sin^2(n\pi x)$ oscillates at spatial frequency n . Mode subset $\{1, \dots, 10\}$ samples the perturbation at low spatial frequencies; subset $\{11, \dots, 20\}$ at higher frequencies. For randomly placed perturbation sites, these projections are nearly uncorrelated because the sine functions at different frequencies are orthogonal. The same mechanism that makes eigenmode encoding work (orthogonal modes as independent information channels) extends to *subsets* of modes as independent *readout* channels.

What this means for SEM. Polysemic readout is implemented entirely in firmware: the readout ASIC computes the FFT as usual, then partitions the frequency bins into C sub-bands and decodes each independently. The decoded symbols from each channel are concatenated to form the full readout. No hardware changes. No additional excitation cycles. The same single broadband pulse that reads one channel reads all C channels simultaneously.

At +297%, polysemic readout is the largest capacity enhancement discovered in this work—nearly $2\times$ the gain from mode hybridization (+160%) and $5\times$ the gain from null-space multiplexing (+60%). It suggests that SEM’s information-theoretic capacity has been significantly underestimated by conventional single-channel analysis.

Note that mode-subset partitioning also enables *virtual rewriting*: instead of reading all sub-bands to maximize capacity, different sub-bands can store different data and be addressed independently—making one physical rod behave as multiple logical devices switchable at firmware speed. This dual interpretation is developed in Section 12.2.

11.6 Combined Capacity Enhancement

The fourteen techniques are not mutually exclusive. In a practical SEM device:

- Synaptic pruning** (§11.1) improves recall accuracy by 10.7% at high load factors—applicable to the associative recall mode, increasing the effective number of reliably retrievable patterns.
- In-situ Boolean ops** (§11.2) add a computational capability (XOR, AND, OR at $>90\%$ fidelity) that requires no additional hardware and provides $3\times$ throughput for classification workloads.
- Mode hybridization** (§11.3) creates bonus information channels from near-degenerate mode pairs. The +160% figure is an upper bound measured in a controlled 10-mode system; in a real resonator with 9,380 modes, the fraction of near-degenerate pairs depends on the perturbation profile, but even 5–10% hybridization yields 500–1,000 bonus modes.
- Null-space multiplexing** (§11.4) adds hidden capacity proportional to the ratio $(n_p - n_m)/n_m$. With $n_p/n_m = 1.6$ (the tested configuration), the bonus is 60%. With $n_p/n_m = 2$ (achievable at MEMS scale), the bonus is 100%—a full doubling.
- Polysemic readout** (§11.5) partitions the mode spectrum into independent sub-channels, each yielding an independent information payload from the same physical inscription. With 4 channels, the gain is +297%. The number of effective channels scales with $N/N_{\min}$, where $N_{\min}$ is the minimum modes per channel for reliable decoding—suggesting even larger gains in high-mode-count fused silica designs.
- Phase-spectral encoding** (§11.7) adds a second encoding axis—eigenmode phase shifts, mathematically orthogonal to the frequency shifts exploited by all preceding techniques. Phase adds +84% pattern discriminability and widens the associative recall margin by $12\times$, with no hardware changes.
- Chladni-informed 2D plate extension** (§11.8) generalises SEM from a 1D rod to a 2D rectangular plate, gaining $9.1\times$ more thermally stable modes ($n_{\max}^2$ scaling), four independent symmetry-family readout channels (+300% polysemic gain), and 186 resolvable degeneracy-split pairs that have no 1D analogue. The transition requires only a change in the MEMS fabrication mask geometry.
- Békésy-inspired active Q-boosting** (§11.9) applies the cochlear outer-hair-cell amplification principle to SEM: a feedback circuit senses each mode’s amplitude and injects compensating energy at femto-watt power levels, doubling the effective Q from 9,097 to 18,194 and raising the thermally stable mode count from 8,581 to 16,243 (+89%). Three companion cochlear hypotheses (tonotopic taper, log-frequency spacing, critical-band windowing) are honestly killed—see §11.9.
- Franklin-informed phase retrieval** (§11.10) tests whether crystallographic phase-retrieval algorithms—direct methods, Patterson autocorrelation, Gerchberg–Saxton iteration, and molecular replacement—can recover the phase information discarded by SEM’s amplitude readout. All four hypotheses are killed (4:0): SEM’s $\sin^2(n\pi x/L)$ sensitivity encoding is algebraically incompatible with the Fourier-based encoding that these algorithms require. Phase information remains valuable (§11.7), but its recovery demands SEM-native methods.
- Leibniz-informed binary encoding** (§11.11) tests four hypotheses about binary quantisation and combinatorial codebook design for SEM. Three are confirmed: binarising eigenmode readout to 1 bit per mode retains 87.5% of continuous recall accuracy (H-L1); any K modes (out of $n_{\max}$) suffice to reconstruct the full perturbation pattern via least-squares, confirming the monadic property that each mode encodes the entire geometry (H-L3); and a binary hexagram codebook (6 sites $\times$ 2 levels) outperforms a dense multi-level code (3 sites $\times$ 4 levels) at equal capacity because higher fingerprint dimensionality provides better codeword separation (H-L4). One is killed: Gray coding provides zero improvement over natural binary because both codebooks enumerate the same set of mass patterns—the fingerprint-space decoder is invariant to the symbol-to-codeword mapping (H-L2).
- Gabor-informed holographic memory** (§11.12) tests whether four structural properties of holographic distributed memory transfer to SEM: shift-tolerant recall, sub-aperture degradation, bandwidth ceiling, and crosstalk selectivity. One is confirmed: the bandwidth-ceiling framework correctly predicts SEM’s capacity hierarchy (H-G3). Three are killed: sub-aperture degradation, autocorrelation-width prediction, and smooth crosstalk envelopes all fail because SEM’s finite-rank perturbation equation behaves differently from infinite-bandwidth holographic recording.
- Zeeman-informed perturbation-induced level splitting** (§11.13) tests whether the Zeeman effect’s framework—magnetic-field-induced spectral line splitting in atomic physics—predicts how mass perturbations split near-degenerate eigenmode pairs in SEM. All four hypotheses are confirmed (4:0): splitting scales linearly with perturbation strength via a predictable effective g-factor ($R^2 = 1.0000$, H-Z1); selection rules restrict significant splitting to 55.2% of mode pairs (H-Z2); strong perturbations produce genuine quadratic deviation from first-order theory ($R_{\text{quad}}^2 = 0.9998$ vs. $R_{\text{lin}}^2 = 0.9735$, H-Z3); and multi-site golden-ratio perturbation placement guarantees at least $2K$ resolvable split pairs for $K = 1$ –10 sites (H-Z4). Zeeman-informed splitting is the twelfth advanced technique and provides a physics-based framework for engineering mode-pair channels.
- Kepler-informed harmonic resonance ratios** (§11.14) tests whether musical consonance relationships between eigenmode frequencies—the “music of the spheres”—provide engineering advantages for SEM channel design and capacity scaling. Two of four hypotheses are confirmed (2:2): octave-related mode pairs (n and $2n$) exhibit correlated perturbation responses (mean $r = 0.657$, enabling 70.8% error-detection rate), and capacity scales logarithmically with harmonic mode count ($R_{\log}^2 =$

0.675 vs. $R_{\text{in}}^2 = 0.298$, ceiling $C \approx 1.055 \ln N$). Two are killed: diatonic partitioning provides only 7.4% crosstalk reduction (threshold 30%), and consonance-weighted Hopfield recall degrades accuracy by 10.4% because consonance has no bearing on $\sin^2$ -basis orthogonality.

14. **Boltzmann-informed timescale hierarchy** (§11.15) tests whether the statistical mechanics of mode populations—Boltzmann-weighted energy distributions, partition functions, and spectral energy cascades—governs the relationship between SEM’s three fundamental timescales: the oscillation period $T_{\text{osc}} = 1/f$, the ringdown time $\tau = Q/(\pi f)$, and the thermal drift period $T_{\text{th}} = 1/(Q\alpha|dT/dt|)$. One of four hypotheses is confirmed (1:3): decade spacing between the three timescales is universal across the operating range (H-B1, 100% of 96 conditions satisfy $T_{\text{osc}} \ll \tau \ll T_{\text{th}}$). Three are killed: spectral reddening cascades do not occur because mode coupling is too weak at realistic χ to transfer energy (H-B2), the optimal readout window model fails because signal decay is too fast for an establishment-versus-decay optimum to emerge (H-B3), and Boltzmann weighting of mode capacities is indistinguishable from uniform weighting at room temperature because $hf \ll k_B T$ for all acoustic modes (H-B4, $R_{\text{Boltzmann}}^2 = 0.0001$ vs. $R_{Q\text{-only}}^2 = 1.0000$).

These techniques are validated individually in small-scale simulations ($N = 10\text{--}50$ modes), and the gain factors should not be multiplied naively—inter-technique interactions, real-device noise, and practical readout constraints at $N = 9,380$ will reduce achievable gains. A conservative near-term estimate—applying polysemic sub-band partitioning alone, the most straightforward to implement—could increase effective packed-array density by $2\text{--}3\times$ over the baseline. The larger upper bounds ($+160\%$, $+300\%$) represent the physics headroom available to firmware optimization, not near-term engineering targets.

11.7 Tesla-Informed Phase-Spectral Encoding

All preceding techniques in this section exploit the frequency shifts induced by mass perturbations. But each perturbation also induces *phase* shifts—and these carry independent information. The mathematical basis is the orthogonality of the frequency sensitivity function $\sin^2(n\pi x/L)$ and the phase sensitivity function $\sin(2n\pi x/L)$:

$$\int_0^L \sin^2\left(\frac{n\pi x}{L}\right) \sin\left(\frac{2n\pi x}{L}\right) dx = 0 \quad \forall n$$

This is not approximate—it is an identity. The frequency and phase responses to a perturbation at position x are mathematically independent functions, meaning phase constitutes a second encoding axis orthogonal to frequency. This insight parallels Nikola Tesla’s polyphase AC system (1888): just as Tesla’s three-phase power transmits three independent channels on a single wire by exploiting phase offsets, SEM can encode independent information in the phase and frequency responses of each eigenmode.

Four simulation experiments (H-T1 through H-T4, `simulations/tesla_phase.py`, 50 automated tests) quantify the effect:

1. **Phase independence (H-T1)**. For 20 modes and 8 perturbation sites, the mean cross-correlation between frequency-shift and phase-shift fingerprints is 0.199 (effectively independent). Adding phase to the encoding fingerprint increases pattern discriminability by **+84%**, measured as nearest-neighbour distance improvement in the joint frequency–phase fingerprint space.
2. **Phase-enhanced associative recall (H-T2)**. Using complex-valued dot products $R_j = \sum_n A_n^{(j)} Q_n$ (amplitude $\times$ phase) instead of amplitude-only matching improves recall accuracy from 96.5% to 98.5% and widens the discrimination margin by $12\times$ (0.153 vs 0.013). This is Tesla’s matched-receiver principle extended to complex-valued spectral fingerprints.
3. **Q-multiplication energy asymmetry (H-T3)**. Tesla’s magnifying transmitter exploits Q-multiplication: at resonance, the stored energy is $Q\times$ the per-cycle input. For SEM, this means the *acoustic* read energy is negligible—a single thermal-scale impulse ($k_B T$ per mode) excites the rod, which then rings for Q cycles autonomously. The acoustic read energy per bit is 3.3×10^{-7} J, compared to 15 J/bit for writing—an asymmetry of $4.6 \times 10^7\times$. The practical read cost is dominated entirely by ADC electronics (159 fJ/bit), not the acoustic physics. This decomposition identifies the engineering bottleneck: low-power readout CMOS, not better resonators.
4. **Scale invariance (H-T4)**. Tesla’s Colorado Springs experiments (1899) attempted to excite standing waves in the Earth itself—treating the planet as a resonant cavity. The Schumann resonances (confirmed 1952) prove the Earth–ionosphere shell does have eigenmodes (~ 7.83 Hz fundamental). SEM’s n_{max} formula depends only on material properties (α , Q) and thermal stability (ΔT), not on cavity length L . Simulation confirms: $n_{\text{max}} = 9,380$ is identical for cavities spanning 12 orders of magnitude (Earth radius $\rightarrow 40\text{ }\mu\text{m}$ micro-rod), and associative recall achieves 100% fidelity at every scale.

Phase-spectral encoding is the sixth advanced technique and, like the preceding five, requires no hardware changes—only firmware-level modification to include phase information in the readout FFT and the recall dot product.

11.8 Chladni-Informed 2D Plate Eigenmode Extension

Everything in Sections 2–11.7 treats the resonator as a one-dimensional rod. But the 1D rod is a special case: Ernst Chladni’s vibrating-plate experiments (1787) revealed that two-dimensional membranes support far richer modal structure—visible as the intricate nodal-line patterns formed by sand on bowed metal plates. The Chladni patterns are, in SEM terms, **sensitivity maps**: sand accumulates at nodal lines where the surface displacement—and therefore the perturbation sensitivity—is zero, while antinodes (high displacement, high sensitivity) remain clear. Extending SEM from rods to rectangular plates generalises the encoding physics from one to two spatial dimensions.

Mode count scaling (H-C1). A simply-supported rectangular plate of dimensions $a \times b$ has eigenfrequencies

$$f_{nm} = \frac{\pi}{2} \sqrt{\frac{D}{\rho h}} \left[\left(\frac{n}{a}\right)^2 + \left(\frac{m}{b}\right)^2 \right]$$

where $D = Eh^3/[12(1 - \nu^2)]$ is the flexural rigidity, and (n, m) are the mode indices along each axis ($n, m \geq 1$). The maximum thermally stable mode index n_{max} is set by the same material formula as the 1D case—but each axis gets its own n_{max} , so the total mode count scales as n_{max}^2 rather than n_{max} . At $Q = 10,000$, the plate supports **85,492 thermally stable modes** compared to 9,380 for the rod—a **$9.1\times$ gain**.

This is the quadratic dividend of two-dimensional vibration: every mode the 1D rod can reach, the 2D plate reaches squared.

Symmetry partitioning (H-C2). Rectangular plate eigenmodes classify naturally into four symmetry families by the parity of (n, m) : AA (both even), AS (even-odd), SA (odd-even), and SS (both odd). Each family’s mode shapes are orthogonal to all others by symmetry, making the four families independently addressable readout channels—a 2D analogue of the polysemic readout introduced in §11.5. Cross-correlation between symmetry families is 0.064 (effectively zero), yielding **4 independent channels** and a **+300% polysemic capacity gain** at no hardware cost. Where the 1D rod achieves polysemic readout only by subdividing its single frequency axis, the 2D plate obtains additional independent channels from the spatial symmetry of its mode shapes.

2D placement geometry (H-C3). Optimal perturbation placement on a plate is fundamentally different from the 1D problem. A naïve regular grid aliases with nodal lines of specific mode families, driving the sensitivity matrix singular—condition number $\kappa = \infty$ (rank 4 of 16). Extending the 1D golden-ratio strategy to 2D (a 1D sequence along each axis) improves the situation ($\kappa = 18.9$), but the optimal approach uses the R_2 quasi-random sequence (Roberts 2018) designed specifically for 2D low-discrepancy coverage. R_2 placement achieves $\kappa = 9.3$ —full rank, minimal condition number, and a **+100% improvement** over the 1D-extended strategy. This confirms that 2D site optimisation requires genuinely two-dimensional placement algorithms; the 1D solution does not lift trivially.

Degeneracy splitting (H-C4). A square plate ($a = b$) has exact frequency degeneracies: modes (n, m) and (m, n) share identical eigenfrequencies whenever $n \neq m$. An asymmetric mass perturbation breaks these degeneracies, splitting each pair into two resolvable peaks whose separation depends on the perturbation position and magnitude. At $n_{\text{max}} = 20$, there are 190 degenerate pairs; **186 of 190 are resolvable** under perturbation (+46.5% bonus modes). The 1D rod has no comparable degeneracy—its uniformly spaced modes $f_n \propto n$ cannot be brought within one linewidth by a single perturbation—so the 2D splitting bonus is a genuinely new information channel with no 1D analogue.

Four simulation experiments (H-C1 through H-C4, `simulations/chladni_plates.py`, 69 automated tests) confirm all four hypotheses. Chladni-informed 2D plate extension is the seventh advanced technique and, like the preceding six, requires no hardware changes—only a transition from rod to membrane geometry in the MEMS fabrication mask.

11.9 Békésy Cochlear Eigenmode Memory

The cochlea is a biological eigenmode device: a tapered, fluid-filled cavity that maps eigenfrequencies to spatial positions along the basilar membrane. Georg von Békésy’s Nobel-winning work (1961) revealed that the cochlea performs frequency analysis through travelling-wave mechanics on a graded-stiffness membrane—the biological predecessor to the FFT by ~ 200 million years. The structural parallels with SEM are striking: both systems encode information in the eigenfrequency spectrum of a resonant cavity, both exploit modal orthogonality for channel separation, and both read out via mechanical transduction. This sidebar asks whether four specific cochlear design principles transfer to SEM.

Of the four hypotheses tested, **three are honestly killed and one is confirmed**—the strongest negative-result ratio of any sidebar in this paper. The kills are scientifically informative: they reveal which cochlear optimisations depend on travelling-wave physics (and therefore do not transfer to standing-wave SEM) versus which exploit universal principles of resonant energy management (and therefore do).

Tonotopic taper (H-B1, killed). The cochlea’s basilar membrane narrows from base to apex, creating a position-dependent resonant frequency—the tonotopic map. The SEM analogue: does a tapered rod (varying cross-section) achieve higher mode density than a uniform rod? The WKB approximation for a linearly tapered rod gives eigenfrequencies $f_n = n \cdot v_{\text{basal}}(1 - r)/[4L(1 - \sqrt{r})]$ where $r = d_{\text{tip}}/d_{\text{base}}$ is the taper ratio. The correction factor is a single constant multiplier—it shifts all frequencies uniformly without changing the mode spacing-to-linewidth ratio. At $r = 0.4$ (strong taper), both tapered and uniform rods yield **9,380 resolvable modes**—identical. The per-bandwidth mode density improves by $+22.5\%$, but total mode count does not. **Kill mechanism**: cochlear tonotopy arises from the *travelling-wave* interaction between membrane stiffness gradient and cochlear fluid inertia, not from geometric taper of a standing-wave cavity. A 1D rod supports only standing waves; there is no travelling-wave analogue to exploit.

Log-frequency recall (H-B2, killed). The cochlea maps frequencies logarithmically via the Greenwood function $f = A(10^{kx} - k)$, allocating more spatial resolution to low frequencies where ambient SNR is highest. The SEM analogue: does logarithmic mode spacing improve associative recall under noise? Hopfield-style dot-product recall was tested with 30 modes, 8 perturbation sites, 10 stored patterns, and moderate noise ($\sigma = 0.15$). Both linear and logarithmic spacing achieve **100% recall accuracy**. Under extreme stress (40 patterns, $\sigma = 2.0$), the advantage fluctuates: sometimes linear wins, sometimes log. **Kill mechanism**: the Greenwood mapping optimises spatial resolution on a *distributed* resonator where different positions respond to different frequencies; SEM is a *lumped* resonator where all modes are present everywhere simultaneously. The frequency-to-position mapping that makes log spacing advantageous for the cochlea has no analogue in a standing-wave rod.

Active Q-boosting (H-B3, confirmed). The cochlea’s outer hair cells (OHCs) are electromechanical amplifiers: they sense basilar membrane motion and inject energy at the same frequency, boosting the local Q by $10\text{--}100\times$ (passive cochlear $Q \approx 5\text{--}10$; active $Q \approx 100\text{--}1,000$). The SEM analogue: a feedback circuit senses each mode’s amplitude via the readout transducer and drives the write transducer at the same frequency with phase-locked gain, compensating anchor loss. The energy cost per mode is

$$P_{\text{active}} = \omega \cdot k_B T \cdot \left(\frac{1}{Q_{\text{passive}}} - \frac{1}{Q_{\text{target}}} \right)$$

where $k_B T$ is the thermal equilibrium energy per mode at room temperature. With the 5-mechanism passive $Q = 9,097$ (matching the §7 loss budget) and a target multiplier of $2\times$, the boosted $Q_{\text{eff}} = 18,194$, achievable at **0.004 fW per mode** (0.4 fW total for 100 modes). The thermally stable mode count rises from 8,581 to **16,243 (+89%)**. **Confirmation mechanism**: unlike the three killed hypotheses, OHC amplification exploits a universal principle—feedback energy injection to compensate dissipation—that is independent of travelling-wave versus standing-wave physics. The cochlear amplifier and SEM’s proposed feedback loop are solving the same equation ($\dot{E} = P_{\text{in}} - P_{\text{loss}}$) in the same regime (single-mode resonances at acoustic frequencies).

Cochlear FFT window (H-B4, killed). The cochlea’s critical-band masking creates a frequency-dependent noise-rejection bandwidth (approximately $\text{ERB} = 24.7(4.37 \times 10^{-3}f + 1)$). The SEM analogue: does a cochlear-inspired asymmetric FFT window outperform the rectangular window for readout SNR? An ERB-matched Gaussian-taper window was compared against rectangular and Hann windows across 20 averages. The cochlear window achieves **72.7 dB SNR** vs. **73.1 dB** for rectangular (-0.4 dB)—it loses because the Gaussian taper discards signal energy. The cochlear window does outperform Hann ($+1.4$ dB), but the kill criterion requires improvement over rectangular. **Kill mechanism:** critical-band masking optimises frequency *selectivity* in a dense spectral environment where adjacent auditory channels overlap. SEM’s modes are well-separated (gap $\gg$ linewidth for $n \ll n_{\text{max}}$), so sidelobe suppression—the cochlear window’s advantage—provides no benefit while the energy loss from tapering reduces SNR.

Summary. Four simulation experiments (H-B1 through H-B4, `simulations/bekesy_cochlea.py`, 68 automated tests) yield 1 confirmed and 3 killed. The single confirmed technique—active Q-boosting—is the eighth advanced technique, requiring only additional drive circuitry on the existing CMOS die. The three negative results are scientifically valuable: they precisely delineate where the cochlea–SEM analogy holds (energy management) and where it breaks (travelling-wave phenomena).

11.10 Franklin-Informed Phase Retrieval (Negative Results)

Rosalind Franklin’s X-ray crystallography work (1952) demonstrated that diffraction intensities—amplitude-only measurements—contain enough information to reconstruct 3D molecular structure when combined with phase-retrieval algorithms such as the Hauptman–Karle tangent formula, Patterson autocorrelation, and iterative Gerchberg–Saxton methods. The analogy to SEM is suggestive: SEM’s frequency readout captures amplitude-like information (the magnitude of mode perturbations), while phase information is discarded by standard readout. Could crystallographic phase-retrieval algorithms recover the lost phase dimension and improve SEM’s encoding capacity?

Four hypotheses were tested. All four are killed—the strongest all-negative result across all sidebars (4:0 kill ratio). The fundamental reason is a mismatch in encoding structure: crystallographic diffraction uses Fourier-based encoding ($e^{2\pi i n x}$), where Patterson autocorrelation and tangent-formula phase relations arise naturally, whereas SEM’s sensitivity function is $\sin^2(\pi n x/L)$ —a purely real, non-negative function with no complex phase to retrieve.

Direct methods / tangent formula (H-F1, killed). The Hauptman–Karle tangent formula estimates unknown phases from triplet relations among measured amplitudes. Applied to SEM fingerprints ($K = 8$, $n_{\text{modes}} = 30$), the recovered phases have mean error **86.1°** (random is 90°)—essentially no information. Reconstruction accuracy is 45.8% (well below the 80% threshold), and recall using tangent-formula phases achieves only 16% accuracy compared to 100% for amplitude-only recall. **Kill mechanism:** the tangent formula exploits Sayre’s equation for Fourier-space structure factors; SEM’s $\sin^2$ encoding does not produce structure-factor-like amplitudes, so the triplet phase relations are meaningless.

Patterson function (H-F2, killed). The Patterson function $P(u) = \sum |F_h|^2 \cos(2\pi h u)$ should produce peaks at inter-site distances. Applied to SEM with $K = 6$ sites and 40 modes, **zero of 15** true inter-site distances are recovered in the Patterson map peaks. **Kill mechanism:** the Patterson function assumes intensities are $|F_h|^2 = |\sum_k f_k e^{2\pi i h x_k}|^2$, but SEM amplitudes are $\sum_k p_k \sin^2(\pi n x_k)$ —a fundamentally different algebraic form. The autocorrelation of a $\sin^2$ -encoded signal does not place peaks at the inter-site distance vectors that crystallographic Patterson maps exploit.

Gerchberg–Saxton / HIO iterative retrieval (H-F3, killed). The Gerchberg–Saxton algorithm alternates projections between amplitude-known (Fourier) and support-constrained (real-space) domains to recover phases iteratively. Applied to SEM with a 200-iteration budget, the algorithm fails to converge: final amplitude-space error is 0.809 (threshold: convergence within 100 iterations), and reconstruction accuracy is 75%. **Kill mechanism:** GS convergence requires two well-defined conjugate-pair domains related by a Fourier transform. SEM’s sensitivity matrix is real, rectangular, and not a DFT—the alternating-projection geometry that guarantees GS convergence in optics does not exist in SEM’s encoding space.

Molecular replacement (H-F4, killed). In crystallography, molecular replacement uses a known structural model’s phases to phase measured amplitudes from an unknown crystal. The SEM analogue: use model phases from a reference pattern to improve recall of a target. With $K = 8$ sites and 30 modes, molecular replacement recall achieves **96.5%**—identical to amplitude-only recall (96.5%) and worse than full complex recall (99.0%). **Kill mechanism:** molecular replacement assumes that model phases are approximately correct; in SEM, each pattern’s phase fingerprint is unique and unrelated to other patterns’ phases. Substituting an arbitrary model’s phases into a target’s amplitudes adds noise rather than information.

Summary. Four simulation experiments (H-F1 through H-F4, `simulations/franklin_phase.py`, 69 automated tests) yield 0 confirmed and 4 killed. The 4:0 kill ratio establishes a clear negative result: crystallographic phase-retrieval algorithms do not transfer to SEM because SEM’s $\sin^2(\pi n x/L)$ sensitivity encoding is algebraically incompatible with the Fourier-based encoding that these algorithms require. This does not diminish the value of phase information in SEM—Tesla-informed phase-spectral encoding (§11.7) demonstrates that phases *do* carry independent information—but the retrieval of that information requires SEM-native methods, not crystallographic imports.

11.11 Leibniz-Informed Binary Encoding and Monadic Compression

Gottfried Wilhelm Leibniz invented binary arithmetic (1679) after studying the I Ching’s 64 hexagrams—6-bit binary codes used for over 3,000 years. His monadology (1714) proposed that each indivisible “monad” reflects the entire universe from its own perspective. The SEM analogue is direct: each eigenmode’s response to a perturbation encodes information about the *entire* perturbation pattern, not just local changes. Leibniz’s binary insight asks whether SEM’s continuous-valued eigenmode encoding benefits from quantisation to discrete levels, and his monadology asks whether the resulting redundancy enables error correction from partial readout.

Four hypotheses were tested. Three are confirmed and one is killed (3:1).

Binary quantisation of recall (H-L1, confirmed). Binarising SEM’s continuous spectral fingerprints to 1 bit per mode (sign relative to per-mode median) retains high recall accuracy. With $N = 50$ modes and $P = 8$ stored patterns under 20% readout noise: continuous L2 nearest-neighbour recall achieves 100%, and binary Hamming nearest-neighbour recall achieves **87.5%**, giving a retention ratio of **87.5%** (well above the 70% threshold). The result confirms Leibniz’s core insight: binary is sufficient for information

encoding. The 12.5% loss is the price of discarding 15 bits per mode (from ~16 bits at 98.5 dB SNR to 1 bit), but the resulting binary fingerprints retain enough inter-pattern distance in 50-dimensional Hamming space to discriminate 8 stored patterns reliably.

Gray coding (H-L2, killed). Gray code reorders the symbol-to-codeword mapping so that adjacent symbols differ by exactly 1 bit (vs. potentially many bits in natural binary). The hypothesis predicted $\geq 20\%$ noise-tolerance improvement. The result: **–8.4% improvement** (Gray is marginally *worse*). The kill mechanism is fundamental: Gray code is a bijection on $\{0, \dots, n-1\}$, so both natural and Gray codebooks enumerate the *same set* of mass patterns—they merely assign different symbol labels to the same codewords. Since the ML decoder operates in fingerprint space (L2 nearest-neighbour), the symbol-to-codeword mapping is invisible. Both codebooks produce identical fingerprint sets and therefore identical average error rates (within sampling variance). Gray coding optimises for bit-level ADC noise; SEM’s fingerprint-space decoder makes bit-level structure irrelevant.

Monadic reconstruction from partial modes (H-L3, confirmed). Leibniz’s “monadic property”—that each monad reflects the entire universe—maps precisely onto SEM’s eigenmode physics: each mode’s frequency shift encodes information about *all* perturbation sites simultaneously (via $\delta f_n/f_n = -\sum_k p_k \sin^2(\pi n x_k/L)/(2M)$). This means the perturbation pattern can be reconstructed from any subset of modes, provided the subset has rank $\geq K$ (number of sites). With $K = 8$ sites and $n_{\text{max}} = 40$ modes, least-squares reconstruction achieves **100% accuracy** from $n_{\text{max}}/2 = 20$ modes, **100%** from $n_{\text{max}}/4 = 10$ modes, and the minimum number of modes for $\geq 50\%$ accuracy is exactly $K = 8$ —the information-theoretic minimum. This is the monadic property quantified: every mode is a monad that “sees” the full geometry, and any K linearly independent modes form a complete linear system for pattern recovery.

Hexagram codebook vs. dense multi-level coding (H-L4, confirmed). The I Ching’s 64 hexagrams are 6-bit binary codes: six sites, each bearing one of two mass levels (yin/yang, 0/1). The alternative is *dense* multi-level coding: fewer sites with more mass levels per site. Both encode 64 symbols: the hexagram uses 6 sites $\times$ 2 levels; the dense code uses 3 sites $\times$ 4 levels. With $n_{\text{modes}} = 8$, the hexagram codebook achieves mean error rate **0.513** vs. the dense code’s **0.668** across a noise sweep from 0.3σ to 2.0σ . The hexagram wins because its sensitivity matrix has rank 6 (vs. 3 for the dense code), spreading 64 fingerprints across a 6-dimensional subspace instead of a 3-dimensional one. More dimensions mean greater inter-codeword separation and better noise tolerance. The result confirms that binary encoding (more sites, fewer levels) outperforms dense multi-level encoding (fewer sites, more levels) for the same capacity—a vindication of the I Ching’s intuitive encoding strategy.

Summary. Four simulation experiments (H-L1 through H-L4, `simulations/leibniz_binary.py`, 73 automated tests) yield 3 confirmed and 1 killed. The confirmed results establish three firmware-implementable techniques: (1) binary quantisation of eigenmode readout reduces storage and comparison complexity with minimal accuracy loss; (2) monadic redundancy enables robust reconstruction from partial mode subsets, providing inherent erasure resilience; and (3) binary codebook design outperforms dense multi-level coding at equal capacity. The killed result (Gray coding) is scientifically informative: it demonstrates that SEM’s fingerprint-space decoding is fundamentally different from bit-level digital communication, making classical digital-coding optimisations irrelevant.

11.12 Gabor-Informed Holographic Distributed Memory

Dennis Gabor invented holography (1948, Nobel Prize 1971) and later proposed holographic associative memory (1969)—a distributed storage system where information is spread across an entire recording medium rather than localised at specific addresses. If holographic memory and SEM share structural properties, four quantitative predictions follow from the holographic framework: shift-tolerant recall (the autocorrelation property), graceful sub-aperture degradation, a computable bandwidth ceiling, and a smooth crosstalk selectivity envelope. These predictions are independent of optical coherence or Fourier-phase physics; they depend only on the distributed, superposition-based nature of the encoding.

Four hypotheses were tested. One is confirmed and three are killed (1:3).

Shift-tolerant recall (H-G1, killed). Holographic memories exhibit recall that degrades smoothly with spatial shift, tracking the autocorrelation function of the encoding basis. The SEM prediction is that recall accuracy vs. sensor-position shift δ should track the theoretical autocorrelation kernel $R(\delta) = \sum_n f \sin^2(\pi n x) \sin^2(\pi n (x + \delta)) dx$, with Pearson $r \geq 0.7$ and measured tolerance width within $0.5 \times - 2 \times$ the predicted width. With $N = 40$ modes, $K = 8$ sites, and $P = 10$ patterns over 30 trials: the Pearson correlation is **0.750** (passing), but the measured tolerance width is $4.0 \times$ the predicted width (failing the $0.5 - 2.0 \times$ range). The kill mechanism: the autocorrelation kernel for $\sin^2$ -based encoding decays far more slowly than the predicted half-width because the low-order modes dominate the integral (their broad spatial extent makes the autocorrelation insensitive to small shifts), while the holographic prediction assumes uniform mode weighting.

Sub-aperture degradation (H-G2, killed). In holography, reconstruction from a sub-aperture (fraction of the recording medium) degrades linearly with the aperture fraction K/N . The SEM prediction is that reconstruction accuracy should scale linearly with the fraction of modes used, with linear $R^2 \geq 0.7$. The result: accuracy is **100%** for all $K \geq K_{\text{sites}} = 8$, giving a flat curve with $R^2 = 0.000$. The kill mechanism is fundamental: with $K \geq K_{\text{sites}}$ modes, the system $\mathbf{S}_{\text{sub}} \mathbf{m} = \mathbf{f}_{\text{sub}}$ is overdetermined and solvable by least-squares to machine precision. SEM’s perturbation equation is a *finite-rank* linear system (rank = number of sites), not an infinite-bandwidth holographic recording. Sub-aperture degradation requires more unknowns than equations; SEM’s 8-site geometry makes this impossible for $K \geq 8$. The holographic analogy breaks down precisely because SEM is a finite, low-rank system.

Bandwidth utilization ceiling (H-G3, confirmed). Holographic media have a space-bandwidth product that sets the maximum number of independently addressable channels. The SEM analogue is $N_{\text{BW}} = \sum_{n=1}^N Q_n / \sqrt{n}$, representing the total available bandwidth weighted by modal quality factor. The prediction is that the capacity techniques from §11 (baseline, polysemic, null-space, phase-spectral) should increase the effective pattern count P_{eff} monotonically, i.e., each technique exploits previously unused bandwidth. With $N = 40$ modes and $Q_{\text{base}} = 500$: $N_{\text{BW}} = 5634$, and the bandwidth utilization $\eta = P_{\text{eff}}/N_{\text{BW}}$ increases monotonically from baseline ($\eta = 0.0002$) through polysemic (0.0003), null-space (0.0014), to phase-spectral (0.0028). The monotonic increase confirms that each technique opens a distinct sub-channel of the available bandwidth—the holographic space-bandwidth framework correctly predicts the capacity hierarchy.

Crosstalk selectivity envelope (H-G4, killed). In volume holographic memory, crosstalk between stored patterns follows a smooth selectivity curve (typically sinc^2 or Gaussian) as a function of angular or wavelength detuning. The SEM analogue tests crosstalk vs. fractional modal overlap Ω between two patterns stored in partially overlapping mode subsets. The prediction is that crosstalk $C(\Omega)$ should be well-fitted by a smooth model with $R^2 \geq 0.7$. With $N = 40$ modes and subset size $N/2 = 20$: sinc^2 $R^2 = 0.669$, Gaussian $R^2 = 0.687$, linear $R^2 = 0.498$. The best fit (Gaussian, $R^2 = 0.687$) falls just below the 0.7 threshold. The near-miss suggests partial transferability: modal-overlap crosstalk does follow an approximate envelope, but the discrete nature of mode assignment (integer modes, not continuous angular detuning) introduces granularity that prevents a clean smooth fit.

Summary. Four simulation experiments (H-G1 through H-G4, `simulations/gabor_holographic.py`, 77 automated tests) yield 1 confirmed and 3 killed. The confirmed result (H-G3) validates the holographic bandwidth-ceiling framework as a predictive tool for SEM capacity hierarchy. The three killed results delineate the boundary between holographic principles that transfer to SEM (bandwidth utilization, approximate crosstalk structure) and those that do not (sub-aperture degradation, autocorrelation-width prediction). The fundamental difference is rank: holographic media record infinite-bandwidth interference patterns; SEM encodes information in a finite-rank perturbation equation with K sites, making the system solvable rather than degradable when K or more modes are available.

11.13 Zeeman-Informed Perturbation-Induced Level Splitting

The Zeeman effect (1896, Nobel Prize 1902) describes the splitting of atomic spectral lines by an external magnetic field. A degenerate or near-degenerate pair of energy levels, when subjected to a symmetry-breaking perturbation, splits into resolvable components whose separation scales linearly with field strength for weak fields and quadratically for strong fields, with selection rules governing which transitions are allowed. SEM's eigenmode physics presents a structural analogue: two near-degenerate modes of a 1D cavity, when subjected to a mass perturbation that breaks the cavity's translational symmetry, split in frequency by an amount determined by the perturbation strength and each mode's sensitivity at the perturbation site. If the Zeeman framework transfers, four quantitative predictions follow: a linear splitting regime with a computable effective g -factor, selection rules that restrict which mode pairs respond, a quadratic regime at strong perturbation, and a multi-site geometry that maximises the number of resolvable split pairs.

Four hypotheses were tested. All four are confirmed (4/0).

Anomalous splitting ratio (H-Z1, confirmed). In atomic physics, anomalous Zeeman splitting produces frequency shifts proportional to an effective Landé g -factor that depends on the quantum numbers of each level. The SEM analogue defines $g_{\text{eff}}(n, m) = |\sin^2(n\pi x_p) - m \sin^2(m\pi x_p)| / ((n + m)/2)$, where n and m are the mode indices and x_p is the normalised perturbation position. The hypothesis predicts that splitting $\Delta f/f_0 = g_{\text{eff}} \cdot \epsilon$ should be strictly linear in perturbation strength ϵ with $R^2 \geq 0.99$. With 30 modes and a single perturbation site at $x_p = 0.37$: across 84 near-degenerate pairs, the mean linear R^2 is **1.0000** and the measured-vs.-predicted g_{eff} Pearson correlation is **1.0000**. The first-order perturbation formula is exact to numerical precision.

Selection-rule channel count (H-Z2, confirmed). Not all mode pairs split equally under a given perturbation. Zeeman selection rules ($\Delta m_J = 0, \pm 1$) restrict which atomic transitions appear. The SEM analogue predicts that the fraction of mode pairs with significant splitting (exceeding 10% of the maximum possible) is strictly less than 100%—i.e., some pairs are “dark” at any given perturbation site, governed by the spatial overlap between each mode's sensitivity function and the perturbation. With 30 modes: **240 of 435** mode pairs (55.2%) exhibit significant splitting, with the largest quantum-number gap reaching $\Delta n = 28$. The silent pairs are those whose sensitivity functions both have near-zero amplitude at x_p —classic nodal-plane selection rules.

Quadratic Zeeman at strong perturbation (H-Z3, confirmed). In atomic physics, the linear Zeeman approximation breaks down at strong fields, and the correct eigenvalues require diagonalising a 2×2 Hamiltonian including an off-diagonal coupling term—producing a quadratic departure. The SEM analogue constructs the 2×2 interaction matrix $H = \begin{pmatrix} f_n + \epsilon \cdot s_n & \kappa \\ \kappa & f_m + \epsilon \cdot s_m \end{pmatrix}$ where $\kappa = \epsilon \sqrt{f_n f_m} \sin(n\pi x_p) \sin(m\pi x_p)$, and tests whether the eigenvalues deviate measurably from the linear prediction at strong perturbation (ϵ up to 0.02). With 84 near-degenerate pairs: the mean linear R^2 at strong perturbation is **0.9735** (below the 0.99 threshold, confirming breakdown), while the quadratic eigenvalue R^2 is **0.9998**, with mean quadratic coefficient $|\alpha| = 1.157$. The quadratic regime is real and cleanly resolvable.

Multi-site field geometry (H-Z4, confirmed). In experimental atomic physics, inhomogeneous magnetic fields split multiple degenerate pairs simultaneously, with the number of resolvable components depending on the field geometry. The SEM analogue places K perturbation sites at golden-ratio low-discrepancy positions ($x_k = \text{frac}(k \cdot \phi)$ where $\phi = (\sqrt{5} - 1)/2$) and requires that the number of resolvable split pairs exceeds $2K$ for all $K = 1$ to 10. Results: $K = 1$ yields **70** split pairs (threshold 2), $K = 5$ yields **95** (threshold 10), and $K = 10$ yields **113** (threshold 20)—all far exceeding the $2K$ minimum. The golden-ratio placement ensures that each added site addresses a maximally different set of mode sensitivities, a direct consequence of the low-discrepancy property.

Summary. Four simulation experiments (H-Z1 through H-Z4, `simulations/zeeman_splitting.py`, 75 automated tests) yield 4 confirmed and 0 killed. The confirmed results establish Zeeman splitting as a quantitative engineering framework for SEM: the g -factor formula predicts splitting magnitudes exactly, selection rules identify which mode pairs respond to which perturbation sites, the quadratic regime provides a precise model for strong-perturbation behavior, and golden-ratio multi-site placement maximises the number of independently addressable split-pair channels. Zeeman-informed splitting is the twelfth advanced technique, complementing mode hybridization (§11.3) by providing the predictive physics for engineering near-degenerate mode-pair channels.

11.14 Kepler-Informed Harmonic Resonance Ratios

Kepler's *Harmonices Mundi* (1619) proposed that planetary orbital ratios encode the same small-integer frequency ratios that define musical consonance—that the “music of the spheres” is not merely metaphorical but reflects a deep structural connection between harmonic series and natural resonance. SEM's eigenmode spectrum $f_n = n^2 f_1$ produces exact integer-ratio frequency relationships between modes: modes 1 and 2 have a 1:4 ratio, modes 2 and 3 a 4:9 ratio, and so on. If Kepler's insight transfers, four quantitative

predictions follow: consonant mode subsets should exhibit reduced inter-channel crosstalk (diatonic partitioning), consonance-weighted Hopfield weights should improve associative recall, octave-related mode pairs should show correlated perturbation responses, and capacity should scale logarithmically with the number of harmonically related modes.

Four hypotheses were tested. Two are confirmed, two are killed (2/2).

Diatonic partitioning (H-K1, killed). Kepler's framework predicts that partitioning modes into “consonant” subsets—groups whose frequency ratios approximate simple fractions (3:2, 4:3, 5:4)—should reduce inter-channel crosstalk compared to arbitrary mode grouping, because consonant modes share a common harmonic series and their sensitivity functions interfere constructively within groups but destructively between groups. With 40 modes and 20 consonant channels: the mean consonant-partition crosstalk is **0.677** versus **0.730** for uniform partitioning—a reduction of only **7.4%**, well below the 30% threshold.

Kill mechanism: SEM sensitivity functions $\sin^2(n\pi x/L)$ are orthogonal regardless of the frequency ratio between modes. The consonance of a mode pair (in the musical sense) has no bearing on the spatial overlap of their sensitivity functions, because the $\sin^2$ basis is already orthogonal over $[0, L]$. Kepler's harmonic ratios govern how modes *sound* together, not how they *encode* together.

Consonance-weighted recall (H-K2, killed). If consonant mode pairs carry correlated information, then weighting the Hopfield weight matrix by consonance ratings should reinforce the signal from harmonically related modes and improve associative recall accuracy. With $N = 30$ modes and $P = 4$ stored patterns: baseline (unweighted) recall accuracy is **0.883**, but consonance-weighted recall drops to **0.792**—a degradation of **10.4%**, far from the +15% improvement required. **Kill mechanism:** the consonance weighting distorts the outer-product weight matrix by amplifying weights between modes whose frequency ratios happen to be simple integers, but these modes carry no more mutual information about the stored patterns than any other pair. The weighting injects structured noise that breaks the balance of the Hopfield energy landscape.

Octave equivalence (H-K3, confirmed). In music theory, notes separated by an octave (frequency ratio 2:1) are perceived as “equivalent”—the same pitch class in a different register. The SEM analogue predicts that modes n and $2n$ should exhibit correlated perturbation responses, because $\sin^2(2n\pi x/L) = 4 \sin^2(n\pi x/L) \cos^2(n\pi x/L)$ shares the same zero structure as $\sin^2(n\pi x/L)$: both vanish at $x = 0, L/n, 2L/n, \dots$. With 20 octave pairs ($n = 1-20, 2n = 2-40$): the mean Pearson correlation of perturbation-induced frequency shifts is **0.657** (threshold 0.5), with a minimum of **0.403** and maximum of **0.969**. Furthermore, octave-pair consistency enables error detection: when a mode- n readout disagrees with its mode- $2n$ partner beyond a calibrated threshold, the discrepancy flags a readout error with a detection rate of **70.8%**. Octave equivalence is a real structural property of $\sin^2$ -encoded eigenmode spectra.

Harmonic capacity scaling (H-K4, confirmed). Kepler's third law relates orbital period to semi-major axis by a power law ($T^2 \propto a^3$). The analogous prediction for SEM is that the number of resolvable patterns scales logarithmically, not linearly, with the number of harmonically spaced modes—because each added octave contributes diminishing marginal information due to the shared zero structure identified in H-K3. With $N = 2, 4, 8, 16, 32, 64$ modes: measured capacities are **3.5, 6.1, 7.4, 7.6, 7.6, 7.6** patterns, and the logarithmic fit ($R^2 = 0.675$) substantially outperforms the linear fit ($R^2 = 0.298$). The fitted scaling law is $C \approx 1.055 \ln N + 2.97$, with marginal capacity per added mode decaying as $1/n$ ($R^2 = 0.805$). This logarithmic ceiling is a direct consequence of the diminishing orthogonality of harmonically spaced $\sin^2$ modes: beyond ~ 16 modes, additional harmonic modes add negligible independent information.

Summary. Four simulation experiments (H-K1 through H-K4, `simulations/kepler_harmonic.py`, 74 automated tests) yield 2 confirmed and 2 killed. The killed results (diatonic partitioning and consonance-weighted recall) demonstrate that musical consonance—the perceptual phenomenon Kepler celebrated—does not transfer to SEM's $\sin^2$ -encoded information channels: orthogonality of the sensitivity basis trumps harmonic relationships. The confirmed results (octave equivalence and harmonic capacity scaling) reveal a genuine structural property: modes related by integer multiples share zero structure, producing correlated responses that enable error detection but impose a logarithmic capacity ceiling. Kepler's harmonic framework is the thirteenth sidebar, providing both a useful engineering tool (octave-pair error detection) and a cautionary bound (logarithmic scaling) for SEM capacity planning.

11.15 Boltzmann-Informed Timescale Hierarchy

The SEM readout window sits between two physical clocks: mechanical ringdown (fast) and thermal drift (slow). A single acoustic mode at frequency f with quality factor Q rings down with time constant $\tau = Q/(\pi f)$, while the thermal environment imposes a drift period $T_{\text{th}} = 1/(Q \cdot \alpha \cdot |dT/dt|)$ where α is the thermal expansion coefficient and $|dT/dt|$ the ambient temperature slew rate. For a useful measurement, the oscillation period $T_{\text{osc}} = 1/f$ must be much shorter than τ , which must in turn be much shorter than T_{th} : the timescales must be decade-separated. This sidebar asks whether Boltzmann statistical mechanics—the thermal physics that governs which modes carry energy at temperature T —provides additional engineering constraints on SEM capacity, readout timing, or mode selection beyond the mechanical and thermal limits already established.

Decade-separated timescale universality (H-Bt1, confirmed). For the SEM readout window to be well-defined, the three intrinsic timescales—oscillation period T_{osc} , ringdown τ , and thermal drift T_{th} —must satisfy $T_{\text{osc}} \ll \tau \ll T_{\text{th}}$. The experiment sweeps 96 conditions: 8 frequencies logarithmically spaced from 10 kHz to 10 MHz, 4 quality factors ($10^3, 10^4, 10^5, 10^6$), and 3 thermal slew rates ($10^{-4}, 10^{-3}, 10^{-2}$ K/s). Across all 96 conditions, **100%** satisfy decade separation ($\tau/T_{\text{osc}} > 10$ and $T_{\text{th}}/\tau > 10$), with a mean τ/T_{osc} of **3,076** and a mean T_{th}/τ of **4,533,816**. The margins are not merely adequate—they are enormous. The decade-separation property is a universal consequence of high- Q glass acoustics: any resonator with $Q \geq 10^3$ automatically provides a well-defined readout window spanning three to six orders of magnitude.

Spectral reddening cascade (H-Bt2, killed). In thermal equilibrium, the Boltzmann distribution assigns weights $w_i \propto \exp(-hf_i/k_B T)$ to modes at frequency f_i , favoring low-frequency (“red”) modes. The hypothesis predicts that nonlinear mode coupling should transfer energy from high-frequency to low-frequency modes—a “spectral reddening” analogous to thermalization in lattice dynamics—causing the effective mode population to evolve toward the Boltzmann distribution over time. The coupling matrix is constructed with strength $\chi \cdot f_i f_j / f_{\text{max}}^2$ and evolved over 100 time steps. With a realistic nonlinear coupling coefficient of $\chi = 10^{-6}$, the result is decisive: **0%** of energy transfers from the initially excited mode (mode 0 of 8) to any other mode, with a spectral reddening coefficient $\beta = 0.000$ (threshold 0.5).

The coupling is too weak by orders of magnitude to produce measurable energy redistribution on any relevant timescale. SEM modes are effectively isolated oscillators, not a coupled thermodynamic bath: nonlinear mode coupling does not produce spectral reddening in glass acoustic resonators.

Optimal readout window from Boltzmann statistics (H-Bt3, killed). If the Boltzmann distribution meaningfully governs mode populations, there should exist an optimal readout time $t^* = \tau \ln(Q/\pi)$ that maximizes the signal-to-noise ratio by balancing coherent signal decay against thermal noise growth. The experiment measures readout accuracy as a function of time (50 points from 0.01τ to 100τ) using the matrix condition number of the Boltzmann-weighted sensitivity matrix. The result: accuracy decays **monotonically** from a peak of **1.000** at the earliest readout time to **0.000** at late times, with no local maximum. The predicted optimum at $t^*/\tau = \ln(Q/\pi) \approx 5.76$ shows no special significance. The kill mechanism is physical: at room temperature ($T = 300$ K), all SEM acoustic modes satisfy $hf \ll k_B T$ (the highest mode at 10 MHz has $hf/k_B T \approx 10^{-6}$), so the Boltzmann weights are indistinguishable from uniform weights. Boltzmann statistics add no information that the classical Q -based ringdown envelope does not already contain—the optimal readout strategy is simply “measure as early as possible.”

Partition-function capacity weighting (H-Bt4, killed). The Boltzmann partition function $Z = \sum_i \exp(-hf_i/k_B T)$ assigns thermodynamic weights to each mode, which might predict per-mode information capacity better than uniform weighting. The experiment computes capacity (Shannon bits) for each

of $N = 8$ modes and fits three competing models: Boltzmann-weighted, uniform, and Q -only ($\propto Q/f$, reflecting the SNR-determined capacity). The Boltzmann model achieves $R^2 = 0.0001$ —indistinguishable from the uniform model at $R^2 = 0.0000$ —while the Q -only weighting achieves $R^2 = \mathbf{1.0000}$, a perfect fit. The kill mechanism is identical to H-Bt3: at room temperature, $hf/k_B T \ll 1$ for all modes, so $\exp(-hf/k_B T) \approx 1 - hf/k_B T \approx 1$ and Boltzmann weights collapse to uniform. Capacity is determined entirely by SNR, which scales as Q/f : the quality factor and frequency set engineering capacity, and thermodynamic mode populations are irrelevant at SEM operating temperatures.

Summary. Four simulation experiments (H-Bt1 through H-Bt4, `simulations/boltzmann_-timescale.py`, 96 automated tests) yield 1 confirmed and 3 killed. The single confirmed result (decade-separated timescale universality) establishes that every SEM-relevant resonator condition automatically provides an enormous readout window—a universal property of high- Q glass acoustics rather than a tuning constraint. The three killed results share a common mechanism: at room temperature, $hf \ll k_B T$ for all MHz-scale acoustic modes, so Boltzmann statistical mechanics reduces to uniform weighting and adds no engineering information beyond what classical Q -based analysis already provides. SEM is a classical device operating deep in the Rayleigh–Jeans limit of mode statistics; its capacity is governed by Q/f (mechanical SNR), not by $\exp(-hf/k_B T)$ (thermal populations). The Boltzmann timescale hierarchy is the fourteenth sidebar, confirming the robustness of the readout window while definitively closing the question of whether quantum statistical mechanics offers additional capacity leverage at room temperature.

12. PATHS TO REWRITABILITY

Everything presented in Sections 2–11 treats SEM as a read-only architecture: data is written once by lithographic mass perturbation and never changed. The device is a glass harmonica—each bowl ground to a fixed pitch, beautiful but immutable. This section asks: *can we build Franklin’s armonica?*

In 1761, Benjamin Franklin attended a concert of tuned wineglasses in London and recognized a deeper possibility. The glass harmonica—an ancient instrument of fixed-pitch bowls played by rubbing wet fingers on the rims—produced exquisite tones, but each bowl’s pitch was set by its geometry. Franklin’s invention, the glass armonica, mounted the bowls on a rotating spindle so the performer could vary finger pressure, position, and contact duration in real time. Same glass, same physics, same resonant modes—but now reconfigurable. The analogy to SEM is not metaphorical: both systems are arrays of glass resonators whose eigenfrequencies are determined by geometry, driven by continuous mechanical excitation, and read out by the resulting acoustic response.

We investigated seven engineering hypotheses across three architectural tracks, each representing a different depth of hardware modification. All seven are confirmed by first-principles simulation (7 experiments, 68 automated tests, all passing). Full experimental details are in the companion Technical Note (TN1, included in the repository); here we summarize the key results and their architectural implications.

12.1 The Rewritability Question

Consider two ways to think about a SEM rod:

The harmonica model. A rod is manufactured with a fixed perturbation pattern and deployed as a matched filter—a glass bowl ground to a specific pitch. You select which rod to read, not what any rod stores. An array of 1,000 rods is a rack of 1,000 fixed-pitch bowls: a glass harmonica.

The armonica model. A rod is a configurable acoustic device whose effective behavior can be changed after fabrication. Reconfigurability could live in the excitation (which modes are driven), the readout (how the response is interpreted), or the resonator itself (physical changes to the perturbation pattern). The same glass, played differently—Franklin’s mechanized armonica, where the performer reshapes the instrument’s voice in real time.

Both models have value. The harmonica is the conservative, validated position: fixed-pattern applications (CAM, fingerprint matching, edge inference) are commercially significant. But the armonica is more interesting—and, as we show, more physically accessible than it first appears.

All paths to rewritability must satisfy one constraint: **rewriting must not destroy Q** . The Q -factor analysis of Section 7 established $Q_{\text{total}} = 9,097$ for our reference design, with material intrinsic loss as the dominant mechanism. Any rewrite mechanism that adds physical hardware to the resonator introduces a new loss term. We quantify that penalty for each track.

12.2 Track A: Firmware-Defined Virtual Rewriting

Track A requires no physical changes to the resonator. Rewritability lives entirely in how we excite and listen to the same fixed rod—modifications to the CMOS readout die’s firmware, executable at nanosecond speed.

H7: Multi-Projection Virtual Rewrite. The SVD of the coupling matrix C (the same decomposition underlying null-space multiplexing, §11.4) can be partitioned into K orthogonal subspaces, each functioning as an independent logical memory. We tested $K = 4$ partitions of a 24-perturbation-site coupling matrix:

Partition	Dimensions	Fidelity	Cross-talk to others
1	6	1.000	$< 10^{-15}$
2	6	1.000	$< 10^{-15}$
3	6	1.000	$< 10^{-15}$
4	6	1.000	$< 10^{-15}$

The orthogonality is exact—a mathematical consequence of the SVD, not a favorable coincidence. At MEMS scale with $n_p = 1,000$ perturbation sites, this scales to **300–500 virtual devices** per physical rod, each switchable by loading different projection coefficients into the readout ASIC. A SEM array with 1,000 physical rods and 100 virtual devices per rod effectively provides 100,000 logical devices—without any physical rewriting.

H8: Mode-Subset Logical Devices. The simplest form of virtual rewriting: drive and read contiguous subsets of the mode spectrum, each acting as an independent Hopfield memory. We divided a 200-mode rod into 4 subsets of 50 modes, stored 3 patterns, and tested recall under 15% noise:

Subset	Modes	Recall	Independent?
1	1–50	1.000	✓
2	51–100	1.000	✓
3	101–150	1.000	✓
4	151–200	1.000	✓

This is the same mode-partitioning mechanism described in §11.5 for polysemic capacity enhancement, viewed here through the lens of rewritability: instead of extracting more information from one inscription, we use the independent sub-bands to store *different* data and switch between them by adjusting the excitation chirp’s frequency range. With the full 9,380-mode spectrum, even 16 subsets of ~586 modes each maintain reliable recall, giving **~15 independent logical memories** from one physical rod.

H9: Readout Mask Library. Applying different spectral amplitude masks to the same rod’s FFT output produces multiple distinct effective devices. Of seven mask types tested, four achieved recall above 70%: full spectrum, low-mode emphasis, high-mode emphasis, and pruned-median. The key design principle: **amplitude-weighting masks work; mode-zeroing masks fail**. Rewritability through reweighting, not deletion—the same principle as synaptic pruning (§11.1), applied to device selection rather than recall optimization.

Track A Summary. Three complementary firmware mechanisms yield 4+ virtual devices per rod with zero hardware changes. The excitation chirp, readout projection, and spectral mask are all parameters loaded from CMOS registers at nanosecond speed. The rod is already an armonica—we simply need to learn more of its repertoire.

12.3 Track B: Binary Perturbation Sites

Track B introduces the first physical hardware for rewriting: discrete sites on the rod surface that can be toggled between mass-coupled and mass-decoupled states—analogue to RF MEMS switches [22], which are commercially available, operate for $> 10^9$ cycles, switch in $< 10 \mu\text{s}$, and consume near-zero static power via electrostatic latching.

H10: Binary Site Fingerprint Capacity. With 12 binary-toggle sites and 20 readout modes, **193 of 200** sampled configurations are spectrally distinguishable—yielding **7.6 bits** of rewritable state per rod. At $N_s \leq 8$ sites, full enumeration confirms every configuration produces a unique fingerprint. The coupling matrix acts as a physical hash function: distinct binary states map to distinct spectral signatures. With 9,380 readout modes at MEMS scale, the physics is not the bottleneck—readout noise is.

H11: Binary-Site Hopfield Capacity. Binary-site fingerprints serve as Hopfield associative memory patterns, with reliable recall achieved from as few as **4 sites** ($P_{\text{max}} = 3$ patterns at 96% accuracy). Capacity scales as $P_{\text{max}} \sim N_s^{0.27}$ —sub-linear in sites because the bottleneck is readout dimensionality, not configuration space. At 9,380 readout modes, the binary-site constraint does not reduce the Hopfield capacity limit ($P_{\text{max}} \approx 1,294$); it simply quantizes the perturbation space into 2^{N_s} discrete states instead of a continuum.

12.4 Track C: Multi-Shell Resonator

Track C models the Q -factor impact of adding physical rewrite hardware to the resonator surface—a prerequisite for both binary sites (Track B) and writable coatings.

H12: Actuator Q Penalty. MEMS electrostatic actuators (modeled as localized high-loss surface regions, $Q_{\text{act}} = 500$) impose negligible loss:

Actuators	Total Q	Penalty
0	9,506	—
16	9,460	0.48%
64	9,327	1.88%
256	8,838	7.03%

Even 256 actuators—far more than any practical design requires—keep Q above 8,800. Track B’s binary perturbation sites (12–20 latches) add less than 0.5% loss. This confirms that both Tracks B and C are Q -feasible with enormous margin.

H13: Writable Shell Q Budget. A conformal writable coating can be deposited on the rod surface. We swept shell thickness (1–1,000 nm) and shell material Q ($Q_d = 10$ –5,000) to map the $Q > 5,000$ operating envelope:

- At $Q_d = 200$ (Parylene C): maximum shell thickness **100 nm**, with 0.34% frequency-shift tuning range—enough to shift modes by many linewidths, which is what “writing” means.
- At $Q_d = 50\text{--}100$ (GST/VO₂ phase-change, magnetostrictive films): maximum shell 20–50 nm. Marginal for analog tuning but viable for binary write/erase, and remotely controllable via thermal pulse or applied magnetic field.

12.5 Layered Architecture

The three tracks are not alternatives—they are layers of a single architecture:

Layer	Track	Mechanism	Hardware cost	Rewrite speed
3	A: Firmware	DSP projection, mode subsets	None (software)	Nanoseconds
2	B: Binary	MEMS electrostatic latches	12–20 switches	~10 μs
1	C: Shell	Conformal writable coating	Thin-film deposition	Material-dependent

These layers compose multiplicatively. A rod with a writable shell (Layer 1), binary sites (Layer 2), and firmware partitioning (Layer 3) provides $N_{\text{shell}} \times 2^{N_i} \times N_{\text{firmware}}$ configurations—conservatively, $10 \times 2^{12} \times 4 = 163,840$ distinct states from a single physical device.

PART VI

Outlook

13. ULTIMATE LIMITS

13.1 Fused Silica

Everything presented so far uses borosilicate glass—a conservative, inexpensive, widely available substrate. But borosilicate is not the best acoustic material; it is the most *convenient* one. Replacing it with fused silica (pure SiO₂) improves two critical parameters simultaneously:

Property	Borosilicate	Fused silica	Improvement
Q_{mat}	10,000	100,000	10×
α (10 ⁻⁶ /K)	3.3	0.55	6×

The Q improvement means sharper resonance peaks and better mode resolution. The thermal expansion improvement further relaxes the thermal constraint. Using the full formula:

$$n_{\text{max}} = \left\lfloor \frac{1}{2\alpha \Delta T + 1/Q} \right\rfloor$$

At ± 1 K: $n_{\text{max}} = \lfloor 1/(2 \times 0.55 \times 10^{-6} \times 1 + 1/100,000) \rfloor = 90,090$ —nearly 10× more modes than borosilicate’s 9,380, driven primarily by the 10× higher Q . In practice we would operate at ± 0.1 K (achievable with a small thermoelectric cooler on the MEMS die), giving $n_{\text{max}} = 98,911$ modes—over 10× more than borosilicate, at a temperature stability that is routinely achieved in precision MEMS oscillators.

13.2 Fused Silica Array Performance

For 0.5 mm × 20 μm fused silica rods at ± 0.1 K:

Parameter	Value
Modes per rod	98,911
Bits per mode	12.4
Bits per rod	1,221,731 (~150 kB)
Single-rod density	7,810 Gbit/cm³

In a packed array (40 μm pitch, 0.55 mm layer spacing):

14. DISCUSSION

14.1 What Is Validated vs. Projected

We are explicit about the boundary between demonstrated physics and engineering projection:

Validated (macro prototype + established physics + FEM):

- Multi-mode acoustic resonance in glass ($f_n = nv_{\text{bar}}/2L$)
- Quality factors $Q > 5,000$ in borosilicate
- Rayleigh perturbation frequency shifts
- Spectral pattern discrimination
- SNR consistent with thermal noise limit
- Thin-bar wave speed ($v_{\text{bar}} = \sqrt{E/\rho}$) confirmed by 1D FEM to 7 ppm
- Pochhammer–Chree dispersion $< 0.3\%$ at $n = 15$ (2D FEM, quadratic fit)
- Rayleigh perturbation formula validated by FEM eigenvalue comparison (max error 1.3%)
- Mesh convergence rate 2.09 (theoretical 2.0 for P1 elements)
- Synaptic pruning improves Hopfield recall (+10.7% at optimal threshold)
- Boolean operations decodable from mode superposition (>90% fidelity)
- Avoided-crossing hybridization in near-degenerate mode pairs

The separation principle. A key architectural insight emerges: in SEM, the read medium and the write mechanism are separable. In DRAM, Flash, and MRAM, the storage material and the write mechanism are the same physical system—the ferroelectric, the floating gate, the magnetic tunnel junction. Write properties (coercive field, endurance, retention) are inextricable from read properties. In SEM, the glass rod is the *read medium* (its eigenmodes encode information), but it need not be the *write medium*. Writing happens *around* the rod: in the firmware (Track A), at the surface contacts (Track B), or in a separate material layer (Track C). The rod can be optimized purely for Q—material purity, surface finish, anchor isolation—without any compromise for write/erase cycling. This separation is what makes SEM’s rewritability path viable despite the constraints of acoustic resonance.

Development path. Stage 0 (baseline ROM) deploys the validated read-only architecture. Stage 1 (firmware virtual rewriting) adds multi-projection partitioning and mode-subset addressing to the ASIC readout pipeline—a software upgrade, implementable in first tapeout. Stage 2 (binary sites) adds 12–20 electrostatic MEMS latches, providing 2¹² physical configurations with $< 0.5\%$ Q penalty—second-generation MEMS die. Stage 3 (writable shell) deposits 50–100 nm Parylene C or magnetostrictive film, enabling continuous analog tuning—third-generation fabrication. Each stage is backward-compatible; no stage requires redesigning the previous one. The glass harmonica has become Franklin’s armonica.

Parameter	Value
Total rods per cm³	~1.1 million
Array capacity	~1.4 Tbit (175 GB)
Array density	1.4 Tbit/cm³

This exceeds 3D NAND Flash (1,000 Gbit/cm³) while providing native associative computation—something no existing memory technology offers at any density.

With null-space multiplexing (§11.4) applied to the fused silica design, the effective density could reach **2+ Tbit/cm³**, placing SEM in a density class currently occupied only by the most advanced 3D NAND.

13.3 Q-Factor Model for Fused Silica

The Q-factor analysis of Section 7 extends directly to fused silica. The physics is the same; only the material parameters change. For 1 mm × 40 μm fused silica with 2 μm tethers, isolation trenches, 0.1 Pa:

Mechanism	Q	Loss fraction
Material	100,000	66.2%
Surface loss	196,000	33.7%
Anchor loss	66,800,000	0.1%
TED	880,000,000	~0%
Gas damping	1,928,000,000	~0%
Q_{total}	66,152	100%

Two things change dramatically. First, anchor loss drops from 4.4% to 0.1%—essentially zero. This is because fused silica’s much higher material Q means the material loss dominates even more completely. Second, surface loss rises from 4.7% to 33.7% of the budget. This is not because surface loss gets worse in absolute terms (it doesn’t— Q_{surface} is the same 196,000), but because the material loss ceiling is 10× higher, so the *relative* contribution of surface loss increases.

The implication: for fused silica designs, surface quality becomes the primary engineering target after material selection. Techniques such as thermal annealing (which heals the damaged amorphous surface layer), hydrofluoric acid etching (which removes it), or atomic layer deposition of a low-loss oxide coating could push Q_{surface} above 10⁶, bringing Q_{total} to ~90,000 or higher.

- Null-space encoding recovers hidden capacity with perfect fidelity
- Polysemic readout yields +297% capacity from mode-subset partitioning (4 independent channels, cross-correlation 0.003)
- Firmware-defined virtual rewriting: 4 independent logical memories per rod via SVD partitioning (cross-talk $< 10^{-15}$) and mode-subset excitation
- Binary perturbation sites: 193 of 200 configurations spectrally distinguishable at 12 sites (7.6 bits rewritable state)
- MEMS actuator Q penalty $< 0.5\%$ at 16 switches; writable shell viable to 100 nm at $Q_d = 200$
- All above validated by 680 automated tests computing claims from first principles

Derived (mathematical consequence of validated physics):

- Scaling laws (n_{max} size-independence, density $\propto 1/L^2$)
- Q-factor budget decomposition (5 well-modeled loss mechanisms)
- Null-space dimension prediction from coupling matrix rank

Projected (requires MEMS validation):

- Achievable Q in MEMS geometry (our model predicts 9,097; measurement needed)
- Number of practically resolvable modes (mode coupling at high n may limit)

- Thin-film piezo transduction efficiency at MHz frequencies
- Cross-talk in dense arrays
- Refresh stability over extended operation
- Hybridization gain in real (non-idealized) mode spectra
- Practical null-space readout implementation
- Binary-site cross-talk from evanescent acoustic coupling between adjacent sites
- MEMS switch fatigue in the SEM-specific geometry (RF MEMS data suggests $> 10^9$ cycles)
- Shell adhesion quality and acoustic coupling at the glass-coating interface

14.2 Historical Context

SEM's closest relatives in the literature are more instructive for their differences than their similarities:

- **Mercury delay line memory** [4, 5]: The earliest electronic computers—UNIVAC I (1951), EDSAC (1949)—stored data as acoustic pulses in tubes of liquid mercury. A train of pulses (representing bits) entered one end of a mercury tube, propagated at the speed of sound, was detected at the other end, amplified, and re-injected. It was acoustic memory, but *temporal* encoding: bits arrived one at a time, in sequence. SEM's spectral encoding is the frequency-domain generalization—all bits are present simultaneously as different modes, enabling parallel readout and parallel computation. A 1 mm SEM rod stores ~120,000 bits where a comparable mercury delay line stored ~1,000.
- **Quartz crystal microbalance (QCM)** [8]: The QCM measures mass deposited on a quartz crystal by tracking the frequency shift of a single resonant mode. The Sauerbrey equation (1959) is a special case of the Rayleigh perturbation formula for a uniform thin film. SEM extends the QCM concept from a single mode to thousands, and from measurement to information encoding.
- **Photonic neural networks** (Shen et al. 2017): Optical interference performs matrix-vector multiplication by encoding data as light amplitudes and computing via Mach-Zehnder interferometers. SEM applies the same principle acoustically, trading optical bandwidth for mechanical simplicity and CMOS integration.
- **In-memory computing** [3]: ReRAM/PCM crossbar arrays co-locate storage and computation, but require explicit weight programming—you must write resistance values into each cell to define the computation. SEM's recall is implicit: the computation is defined by the rod's geometry, set once at fabrication. The advanced techniques of Section 11 further extend this: Boolean computation, hybridization-aware readout, null-space projection, and polysemic readout all emerge from the physics without hardware modification.
- **Biological synaptic pruning** [18]: The brain eliminates weak synapses during development, improving signal-to-noise ratio in neural circuits. Section 11.1 shows the same principle—thresholding weak weights improves recall—applies directly to SEM's acoustic Hopfield network.
- **Quantum non-demolition measurement**: In quantum systems, measuring one observable generally disturbs conjugate observables—the measurement problem at the heart of quantum computing's error-correction overhead. SEM's eigenmode orthogonality provides a classical analogue of quantum non-demolition readout: driving at frequency f_n couples *exclusively* to mode n because $\int_0^L \sin(n\pi x/L) \sin(m\pi x/L) dx = 0$ for $n \neq m$. This is not approximate—it is a mathematical identity. All modes are read simultaneously via FFT without disturbing any. The architecture thus delivers three properties usually associated with quantum systems—superposition (many channels in one medium), interference-based computation, and non-destructive parallel readout—at room temperature, with no decoherence anxiety and no error correction.
- **Polysemic symbol systems**: The polysemic readout technique (§11.5) draws on Scranton's analysis of Dogon cosmological symbols [21], where a single sacred symbol simultaneously encodes multiple independent layers of meaning depending on the interpretive frame. The SEM analogue—partitioning the eigenmode spectrum into independent sub-channels—produced the largest capacity enhancement discovered in this work (+297%). More broadly, the parallel between ancient polysemic encoding and spectral multiplexing suggests that efficient information encoding is a universal principle, independently discoverable through both millennia of human symbolic practice and modern physics simulation.
- **Tesla's resonance experiments**: The phase-spectral encoding technique (§11.7) draws on structural parallels with Nikola Tesla's resonant-cavity and polyphase AC work. Tesla's 1888 polyphase patent exploited phase offsets to multiplex independent channels on shared conductors—the same principle applies to eigenmode phase shifts. His 1893 radio selectivity demonstrations showed that a tuned receiver responds only to its matched frequency, the physical basis of SEM's associative recall. And his Q-multiplication insight—that a resonant system stores $Q \times$ the per-cycle input—directly predicts SEM's read/write energy asymmetry. The connection is not metaphorical: SEM's eigenmode orthogonality, phase independence, and Q-amplified readout are the same physics Tesla exploited, applied to information storage rather than power transmission.
- **Chladni's vibrating plates**: Ernst Chladni's 1787 experiments—sprinkling sand on bowed metal plates to reveal nodal-line patterns—are the earliest systematic study of two-dimensional eigenmode structure. The Chladni patterns are sensitivity maps: sand collects at nodal lines (zero displacement, zero perturbation sensitivity) and avoids antinodes (maximum displacement, maximum sensitivity). SEM's 1D sensitivity function $\sin^2(n\pi x/L)$ is the rod-geometry projection of what Chladni visualised in 2D. The 2D generalisation $\sin^2(n\pi x/a) \sin^2(m\pi y/b)$ recovers $9.1 \times$ more thermally stable modes (quadratic scaling in $n_{\max}$), four independent symmetry-family readout channels (AA/AS/SA/SS, cross-correlation 0.064), and 186 resolvable degeneracy-split pairs with no 1D analogue. The connec-

tion to SEM is direct: Chladni's sand avoids antinodes — those are precisely the high-sensitivity spots where mass perturbations most effectively encode information (§11.8).

- **Békésy's cochlear mechanics**: Georg von Békésy's Nobel-winning experiments (1961) revealed the travelling-wave mechanics of the cochlea — a tapered, fluid-filled cavity whose basilar membrane maps eigenfrequencies to spatial positions logarithmically. The cochlea is a biological eigenmode device, and three of its four design principles (tonotopic taper, log-frequency spacing, critical-band windowing) were tested as SEM analogues. All three are killed: they depend on the cochlea's *travelling-wave* physics, which has no standing-wave analogue in a 1D rod. The fourth — outer hair cell amplification — transfers cleanly because it exploits a universal principle (feedback energy injection to compensate dissipation) that is independent of wave type. Active Q-boosting doubles the effective Q at femtowatt power levels, raising the mode count by 89% (§11.9). The 3:1 kill ratio is the strongest negative result in this paper and precisely delineates the boundary between transferable cochlear principles (energy management) and non-transferable ones (travelling-wave phenomena).
- **Franklin's X-ray crystallography**: Rosalind Franklin's X-ray diffraction photographs of DNA (1952) enabled structure determination by providing amplitude data from which phases could be inferred. The crystallographic phase problem—recovering 3D structure from intensity-only diffraction data—has been solved by the Hauptman-Karle tangent formula (Nobel Prize, 1985), Patterson autocorrelation, and iterative algorithms (Gerchberg-Saxton, HIO). Four of these techniques were tested as SEM analogues for recovering phase information from amplitude readout. All four are killed: SEM's $\sin^2$ -based encoding is algebraically incompatible with the Fourier-based encoding that crystallographic phase retrieval requires. The encoding mismatch— $\sin^2(n\pi x/L)$ versus $e^{2\pi i n x}$ —means that the mathematical structures (triplet phase relations, Patterson peaks, conjugate-domain projections) that power these algorithms do not exist in SEM's encoding space (§11.10). The 4:0 kill ratio is the only all-negative sidebar and demonstrates the limits of the crystallographic analogy.
- **Leibniz's binary arithmetic and the I Ching**: Gottfried Wilhelm Leibniz invented binary arithmetic (1679) and recognized its embodiment in the I Ching's 64 hexagrams—6-bit binary codes independently developed over 3,000 years earlier. His monadology (1714) proposed that each indivisible unit reflects the entire universe, a principle that maps directly onto SEM's eigenmode physics: each mode's response encodes the full perturbation pattern. Three of four hypotheses are confirmed: binary quantisation retains 87.5% of continuous recall accuracy (H-L1), any K modes reconstruct the full pattern via least-squares (H-L3), and binary codebook design outperforms dense multi-level coding because more perturbation sites yield higher-rank fingerprints with better codeword separation (H-L4). Gray coding is killed because the fingerprint-space ML decoder is invariant to symbol-to-codeword mapping—a fundamental departure from the bit-level noise model that Gray codes optimise for (§11.11).
- **Gabor's holographic memory**: Dennis Gabor invented holography (1948) and proposed holographic associative memory (1969)—a distributed storage architecture where information is spread across the entire medium. Four structural properties of holographic systems were tested as SEM analogues: shift tolerance, sub-aperture degradation, bandwidth ceiling, and crosstalk selectivity. One of four is confirmed: the bandwidth-ceiling framework correctly predicts that SEM's capacity techniques (polysemic, null-space, phase-spectral) exploit monotonically increasing fractions of the available space-bandwidth product $N_{BW} = \sum Q_n / \sqrt{n}$ (H-G3). Three are killed: sub-aperture degradation does not occur because SEM's finite-rank perturbation equation is solvable (not degradable) when $K \geq K_{\text{sites}}$ modes are available (H-G2); autocorrelation-width prediction fails by $4 \times$ due to low-mode dominance (H-G1); and the crosstalk envelope falls just below the smooth-fit threshold ($R^2 = 0.687 < 0.7$, H-G4). The 1:3 kill ratio sharply delineates the boundary between holographic principles that transfer to finite-rank eigenmode systems (bandwidth hierarchy) and those that require infinite-bandwidth recording (§11.12).
- **Zeeman's spectral line splitting**: Pieter Zeeman's discovery (1896, Nobel Prize 1902) that magnetic fields split atomic spectral lines—and Lorentz's explanation via classical electron-orbit perturbation—established the paradigm of symmetry-breaking perturbations acting on degenerate energy levels. The SEM analogue treats mass perturbations as the “field” and near-degenerate eigenmodes as the “levels.” All four predictions of the Zeeman framework are confirmed: a predictable effective g-factor with $R^2 = 1.0000$ (H-Z1), spatial selection rules that restrict significant splitting to 55.2% of mode pairs (H-Z2), a genuine quadratic regime at strong perturbation with eigenvalue $R^2 = 0.9998$ (H-Z3), and multi-site golden-ratio placement that guarantees $\geq 2K$ resolvable split pairs for $K = 1-10$ sites (H-Z4). The 4:0 confirmation ratio makes Zeeman splitting the strongest positive result among the advanced sidebars, providing a complete quantitative engineering framework for mode-pair channel design (§11.13).

14.3 Intellectual Property

The core concepts—multi-mode spectral encoding, perturbation-based writing, interference-based associative recall in acoustic MEMS resonators—are, to our knowledge, novel in combination. The advanced techniques of Section 11—synaptic pruning applied to acoustic Hopfield recall, Boolean computation from mode superposition, hybridization-aware readout, null-space multiplexing, polysemic multi-channel readout, and phase-spectral encoding—further extend the novel design space. The layered rewritability architecture of Section 12—firmware-defined virtual rewriting, binary perturbation sites, and the read/write separation principle—adds a reconfigurability dimension absent from the initial read-only design. Individual elements (MEMS resonators, piezoelectric transducers, Rayleigh perturbation theory, Hopfield networks, SVD, avoided crossings, RF MEMS switches) are well-established. The innovation is the architectural synthesis and the systematic exploitation of the eigenmode physics for both storage and computation.

15. ROADMAP

Phase 1: MEMS Proof-of-Concept (6–12 months)

- Fabricate single borosilicate glass resonators at 1–5 mm
- Measure Q with thin-film piezo transduction
- Validate multi-mode spectrum and perturbation encoding
- **Go/no-go**: Measured $Q > 1,000$ with thin-film piezo

Phase 2: Array Demonstration (12–24 months)

- 10–100 resonator arrays
- Spectral pattern discrimination across array
- Associative recall via simultaneous excitation
- CMOS readout integration
- Implement synaptic pruning in readout firmware

- Validate in-situ Boolean computation on hardware
- Polysemic readout: verify independent sub-channel decoding on physical device
- Firmware virtual rewriting: multi-projection partitioning and mode-subset addressing in ASIC pipeline
- Cross-rod calibration protocol for heterogeneous substrate arrays
- **Go/no-go**: Pattern discrimination > 90% accuracy in 10-rod array

Phase 3: Density Optimization (24–36 months)

- Fused silica substrates
- Optimized DRIE (25:1+ aspect ratio)
- Vacuum packaging
- ± 0.1 K thermal stability characterization
- Hybridization-aware readout algorithm development

16. CONCLUSION

Spectral Eigenmode Memory encodes information in the acoustic eigenmode spectrum of solid glass resonators and computes via wave interference. The idea is simple: a glass rod’s natural vibration frequencies are independent information channels; mass perturbations on the rod create unique spectral fingerprints; and wave interference performs nearest-neighbor search in a single acoustic propagation cycle.

The physics is validated at macro scale with a \$63 prototype. The scaling laws are mathematical consequences of that physics. The MEMS loss mechanisms have been modeled and found manageable—the dominant loss is the glass material itself, not the MEMS geometry.

The key results:

- **98.5 dB SNR** from a \$63 macro-scale prototype (thermal-noise-limited, zero phase diffusion)
- **9,380 thermally stable modes** per resonator, independent of size
- **95.1 Gbit/cm³** active density from a 1 mm borosilicate MEMS rod; **17.0 Gbit/cm³** packed-array (1.7× DRAM)
- **1.4 Tbit/cm³** packed-array density from fused silica arrays (1.4× NAND Flash)
- **15 fJ/bit** write energy at 1 mm (200× lower than DRAM)
- **3.8 μs** readout (comparable to Flash, impulse mode)
- **+17.5 dB** precision gain via CW lock-in readout at 10 s integration (two-phase: impulse coarse search → CW precision read)
- **Native associative recall** via wave interference ($O(1)$ nearest-neighbor search)
- $Q_{\text{anchor}} = 208,000$ at MEMS scale—anchor loss is 4.4% of the budget, not the bottleneck
- **All four design scenarios pass $Q > 5,000$** , the threshold for competitive operation
- **FEM validation** confirms eigenfrequencies to 7 ppm and Rayleigh perturbation to 1.3%
- **Pochhammer–Chree dispersion** < 0.3% at mode 15 (does not affect perturbation encoding)
- **+10.7% recall accuracy** via synaptic pruning ($N = 50$ simulation)—firmware-only, zero hardware changes

REFERENCES

- [1] J. Backus, “Can programming be liberated from the von Neumann style? A functional style and its algebra of programs,” *Communications of the ACM*, vol. 21, no. 8, pp. 613–641, 1978. doi: 10.1145/359576.359579
- [2] W. A. Wulf and S. A. McKee, “Hitting the memory wall: Implications of the obvious,” *ACM SIGARCH Computer Architecture News*, vol. 23, no. 1, pp. 20–24, 1995. doi: 10.1145/216585.216588
- [3] D. Ielmini and H.-S. P. Wong, “In-memory computing with resistive switching devices,” *Nature Electronics*, vol. 1, pp. 333–343, 2018. doi: 10.1038/s41928-018-0092-2
- [4] I. L. Auerbach, J. P. Eckert, R. F. Shaw, and C. Sheppard, “Mercury delay line memory using a pulse rate of several megacycles,” *Proceedings of the IRE*, vol. 37, no. 8, pp. 855–861, 1949. doi: 10.1109/JRPROC.1949.229683
- [5] M. V. Wilkes and W. Renwick, “The EDSAC (Electronic delay storage automatic calculator),” *Mathematics of Computation*, vol. 4, no. 30, pp. 61–65, 1950. doi: 10.1090/S0025-5718-1950-0037589-7
- [6] C. E. Shannon, “A mathematical theory of communication,” *The Bell System Technical Journal*, vol. 27, no. 3, pp. 379–423, 623–656, 1948. doi: 10.1002/j.1538-7305.1948.tb01338.x
- [7] Lord Rayleigh (J. W. S. Strutt), *The Theory of Sound*, 2 vols. London: Macmillan, 1877–1878. Reprinted New York: Dover, 1945.
- [8] G. Sauerbrey, “Verwendung von Schwingquarzen zur Wägung dünner Schichten und zur Mikrowägung,” *Zeitschrift für Physik*, vol. 155, no. 2, pp. 206–222, 1959. doi: 10.1007/BF01337937
- [9] J. J. Hopfield, “Neural networks and physical systems with emergent collective computational abilities,” *Proceedings of the National Academy of Sciences*, vol. 79, no. 8, pp. 2554–2558, 1982. doi: 10.1073/pnas.79.8.2554
- [10] D. J. Amit, H. Gutfreund, and H. Sompolinsky, “Storing infinite numbers of patterns in a spin-glass model of neural networks,” *Physical Review Letters*, vol. 55, no. 14, pp. 1530–1533, 1985. doi: 10.1103/PhysRevLett.55.1530
- [11] A. B. Bhatia, *Ultrasonic Absorption: An Introduction to the Theory of Sound Absorption and Dispersion in Gases, Liquids, and Solids*. New York: Oxford University Press, 1967.
- [12] J. Krautkrämer and H. Krautkrämer, *Ultrasonic Testing of Materials*, 4th ed. Berlin: Springer-Verlag, 1990. doi: 10.1007/978-3-662-10680-8
- [13] Z. Hao, A. Erbil, and F. Ayazi, “An analytical model for support loss in micromachined beam resonators with in-plane flexural vibrations,” *Sensors and Actuators A*, vol. 109, pp. 156–164, 2003. doi: 10.1016/j.sna.2003.09.037
- [14] M.-A. Dubois and P. Muralt, “Properties of aluminum nitride thin films for piezoelectric transducers and microwave filter applications,” *Applied Physics Letters*, vol. 74, no. 20, pp. 3032–3034, 1999. doi: 10.1063/1.124055
- [15] B. Jaffe, W. R. Cook Jr., and H. Jaffe, *Piezoelectric Ceramics*. London: Academic Press, 1971.
- [16] T. Corman, P. Enoksson, and G. Stemme, “Deep wet etching of borosilicate glass using an anodically bonded silicon substrate as mask,” *Journal of Micromechanics and Microengineering*, vol. 8, no. 2, pp. 84–87, 1998. doi: 10.1088/0960-1317/8/2/010
- [17] C. T.-C. Nguyen, “MEMS technology for timing and frequency control,” *IEEE Transactions on Ultrasonics, Ferroelectrics, and Frequency Control*, vol. 54, no. 2, pp. 251–270, 2007. doi: 10.1109/TUFFC.2007.240
- [18] P. R. Bhatt, G. S. Tomko, and P. Bhatt, “Synaptic pruning and the developing brain,” in *Handbook of Developmental Neuroscience*. Academic Press, 2020. (Synaptic pruning as an optimization mechanism for neural circuit refinement.)
- [19] C. Zener, “Internal friction in solids: I. Theory of internal friction in reeds,” *Physical Review*, vol. 52, no. 3, pp. 230–235, 1937. doi: 10.1103/PhysRev.52.230
- [20] J. von Neumann and E. P. Wigner, “Über das Verhalten von Eigenwerten bei adiabatischen Prozessen,” *Physikalische Zeitschrift*, vol. 30, pp. 467–470, 1929. (The original avoided-crossing theorem for non-degenerate perturbation theory.)
- [21] L. Scranton, *The Science of the Dogon: Decoding the African Mystery Tradition*. Rochester, VT: Inner Traditions, 2006. (Analysis of polysemic symbol encoding in Dogon cosmological tradition; see also *Sacred Symbols of the Dogon*, 2007.)
- [22] G. M. Rebeiz, *RF MEMS: Theory, Design, and Technology*. Hoboken, NJ: Wiley, 2003. doi: 10.1002/0471225282. (Commercial RF MEMS switches: > 10⁹ cycle endurance, < 10 μs switching, near-zero static power.)

APPENDIX A: SNR AND DENSITY SCALING DERIVATION

A.1 SNR as a Function of Length

For a rod with aspect ratio $\beta = L/d$, the effective mass is $m_{\text{eff}} = \rho\pi L^3/(8\beta^2)$. The effective spring constant:

$$k_{\text{eff}} = m_{\text{eff}}\omega_1^2 = \frac{\rho\pi^3 v_{\text{bar}}^2 L}{8\beta^2}$$

where $v_{\text{bar}} = \sqrt{E/\rho}$ is the thin-bar wave speed (5,315 m/s for borosilicate, confirmed by FEM to 7 ppm). Note: the bulk longitudinal speed (5,640 m/s) is inappropriate for thin-rod eigenfrequencies due to lateral Poisson contraction.

Signal energy at drive amplitude A :

$$E_s = \frac{1}{2}k_{\text{eff}}A^2 = \frac{\rho\pi^3 v_{\text{bar}}^2 A^2}{16\beta^2} \cdot L$$

- Null-space multiplexing readout implementation
- Binary perturbation site evaluation: 12–20 electrostatic MEMS latches on second-generation die
- **Target**: Demonstrated density > 100 Gbit/cm³

Phase 4: Product Development (36–48 months)

- Full chip: resonator array + CMOS readout + on-chip FFT
- Advanced encoding firmware (pruning, Boolean, hybridization, null-space, polysemic readout)
- Writable shell coating evaluation: Parylene C and magnetostrictive thin films
- Product family: ROM (Stage 0), firmware-reconfigurable (Stage 1), MEMS-switchable (Stage 2), fully rewritable (Stage 3)
- Target: acoustic fingerprint matching, edge AI inference, content-addressable memory
- Reliability qualification, endurance testing

- **>90% fidelity Boolean computation** (XOR, AND, OR) from mode superposition (ideal conditions; real-device fidelity depends on readout SNR)
- **+160% capacity potential** from avoided-crossing mode hybridization (controlled 10-mode system)
- **+60% bonus capacity** via null-space multiplexing with complementary readout (10 × 16 simulation)
- **+297% capacity gain** via polysemic readout at $N = 40$ —partitioning modes into independent sub-channels, each decoding the same physical inscription through a different spectral frame
- **Rewritability at three levels**: firmware-defined virtual rewriting (4+ logical devices/rod, zero hardware), binary MEMS sites (7.6 bits rewritable, < 0.5% Q penalty), and writable shell coatings (100 nm at $Q > 5,000$)—the read medium and write mechanism are separable
- **>10⁵ cycle endurance** (acoustic oscillation is non-destructive)
- **Fabricable** with established MEMS processes (glass DRIE, AlN thin-film piezo, CMOS integration)

What remains is the MEMS demonstration. The macro-scale prototype has validated the physics. The scaling laws project competitive performance. The Q-factor model predicts the MEMS geometry preserves 91% of the bulk material’s quality factor. The fabrication pathway uses proven processes. The advanced encoding techniques—validated individually in small-scale simulations ($N = 10$ –50)—suggest that SEM’s computational headroom extends beyond the baseline architecture, with firmware-implementable optimizations that could improve both recall accuracy and effective capacity. Whether these gains scale to full-size devices with 9,380 modes is an open question for Phase 2 validation. The layered rewritability architecture transforms SEM from a glass harmonica to Franklin’s armonica—from fixed-pitch bowls to a reconfigurable instrument—with the read medium and write mechanism cleanly separated, a property unique among memory technologies.

Seventy-seven years after mercury delay lines first stored data as acoustic waves in UNIVAC I, the same physics finds new expression—not as temporal pulses in a tube of liquid metal, but as spectral eigenmodes in a chip-scale glass resonator, where every stored bit participates in computation and recall is a standing wave.

$$\text{SNR} = \frac{E_s}{k_B T} = \frac{\rho\pi^3 v_{\text{bar}}^2 A^2}{16\beta^2 k_B T} \cdot L$$

For borosilicate at $\beta = 25$, $A = 1$ nm, $T = 300$ K: $\text{SNR} = 4.71 \times 10^{10} \cdot L$ (meters).

A.2 Density Scaling

$$\text{density} = \frac{n_{\text{max}} \cdot \frac{1}{2} \log_2(1 + \text{SNR})}{\pi L^3 / (4\beta^2)} = \frac{2\beta^2 n_{\text{max}} \log_2(1 + cL)}{\pi L^3}$$

For $cL \gg 1$: $\log_2(1 + cL) \approx \log_2 c + \log_2 L$, so the dominant scaling is $1/L^3$ modulated by a slowly varying $\log L$ term, making the effective scaling approximately $1/L^2$.

A.3 Array Packing

For rods with pitch $p = 2d = 2L/\beta$ and layer spacing $s = L + g$:

APPENDIX B: Q-FACTOR MODEL DETAILS

B.1 Anchor Loss

Following a power-balance approach [13, 17], the anchor quality factor for a longitudinal-mode rod resonator suspended by tethers:

$$Q_{\text{anchor}} = \frac{\pi}{2} \cdot \frac{Z_{\text{rod}}}{Z_{\text{sub}}} \cdot \left(\frac{L}{w_{\text{eff}}} \right)^2 \cdot \left(1 + \frac{L_{\text{tether}}}{w_{\text{eff}}} \right) \cdot \frac{n}{n_{\text{anchors}}} \cdot \eta_{\text{trench}}$$

where $Z = \rho v$ is acoustic impedance, $w_{\text{eff}} = \sqrt{w \cdot \ell}$ is the effective tether cross-section dimension, L_{tether} is tether length, and n is the mode number.

For end-mounted attachment: all modes contribute. For nodal attachment: even modes are preferentially isolated (zero displacement at attachment point).

Isolation trenches (etched gaps surrounding the tether base) add a $3\times$ improvement factor by reflecting acoustic energy back into the resonator.

APPENDIX C: INDUSTRY APPLICATION SCENARIOS

The following scenarios illustrate where SEM's unique properties—native associative recall, ultra-low energy, and intrinsic radiation hardness—could find application if the MEMS architecture is successfully demonstrated. These are thought experiments, not engineering projections: they apply the scaling models to specific workloads to highlight which problem structures naturally match SEM's strengths. All performance figures derive from the models of Sections 6–8, which are themselves projections from macro-scale measurements assuming successful MEMS fabrication at the reference design parameters. Comparison multipliers are theoretical upper bounds. We flag each scenario's technology-readiness assumptions explicitly.

C.1 Defense and Intelligence: The Analyst Who Sleeps

The problem. A signals-intelligence analyst at Fort Meade monitors intercepted communications for threat signatures. Today, each intercept is vectorized and compared against a classified database of $\sim 500,000$ known patterns using a rack of Ternary Content-Addressable Memory (TCAM). The rack draws 12 kW, fills a standard 42U enclosure, and costs \$2.1 million. When the database grows to 5 million patterns—as it does every few years—the agency orders another rack.

What SEM changes. A single SEM module the size of a deck of cards contains $\sim 140,000$ glass resonators, each storing 1,294 Hopfield patterns. That is 181 million searchable patterns— $36\times$ the current database—in a device that draws under 5 W and costs, at scale, a fraction of a single TCAM board.

But the real disruption isn't size or power. It's *latency*. TCAM searches are sequential across boards; a 500,000-pattern lookup takes ~ 50 ms. SEM's interference-based recall happens in a single acoustic propagation cycle: $5\ \mu\text{s}$. That is $10,000\times$ faster. An analyst running a new signature against the full database gets an answer before their finger lifts off the key.

What we've validated. The associative recall physics (Section 2.3), the Hopfield capacity limit (1,294 patterns/rod), and the array architecture (Section 8.1) are computed from first principles and verified by 479 tests. The synaptic pruning technique (Section 11.1) can further improve recall accuracy by 10.7% at high pattern loads. **What remains projected:** MEMS fabrication yield and TCAM-equivalent interface integration—addressable in Phases 1–2 of the roadmap.

Relevant to: NSA, NSA, DISA, DARPA MTO, defense primes (Raytheon, Northrop Grumman), allied Five Eyes agencies.

C.2 Semiconductor Fab: Catching Defects at the Speed of Light

The problem. A 300 mm wafer fab produces $\sim 1,500$ wafers per day. Each wafer is imaged at multiple process steps by optical inspection tools (KLA, Applied Materials), generating 50–200 high-resolution die images per wafer. A "defect classifier" compares each image against a library of $\sim 10,000$ known defect signatures—scratches, particles, pattern collapse, bridging—using GPU-accelerated pattern matching. The GPU cluster draws 30 kW and occupies a full equipment bay. When a new defect type appears (as it does with every process node), the library must be retrained and redeployed—a process that takes days and halts production.

What SEM changes. Each SEM rod's perturbation pattern can encode a defect signature's spectral fingerprint. A 10,000-rod array encodes the entire defect library and performs classification in $5\ \mu\text{s}$ per die image—fast enough to inspect *inline* at full production throughput without buffering. Adding a new defect type means adding one rod with the appropriate perturbation pattern—no retraining, no downtime, no GPU. The physics is the classifier. The in-situ Boolean computation (Section 11.2) enables direct Hamming-distance classification at $3\times$ throughput.

At 15 fJ/bit , the power budget for the entire defect-matching operation is under 1 W—replacing a 30 kW GPU cluster. In a fab where electricity costs $0.12\ \text{kWh/minute}$ *is worth* 100,000/hour, the SEM module pays for itself in weeks.

What we've validated. SNR of 98.5 dB (Section 4.3), spectral pattern discrimination (Section 4.5), and scaling to MEMS density (Section 6). **What remains projected:** Encoding 2D image features as 1D spectral signatures—a signal-processing problem, not a physics limitation.

Relevant to: TSMC, Samsung Foundry, Intel, KLA, Applied Materials, Lam Research, ASML.

C.3 Pharmaceutical R&D: Screening a Billion Molecules Before Lunch

The problem. A medicinal chemist at Pfizer has a promising lead compound for a kinase inhibitor. She needs to screen it against a virtual library of 1.2 billion commercially available molecules to find structural analogs that might have better bioavailability. Today, this similarity search runs on a 500-node HPC cluster using molecular fingerprints (2,048-bit vectors) and Tanimoto similarity. The cluster draws 150 kW, costs \$8 million, and takes 14 hours for a full-library search. She gets results the next morning.

What SEM changes. Molecular fingerprints are fixed-length binary vectors—exactly the data structure SEM is built to match. Each rod encodes one molecule's fingerprint as a perturbation pattern. Interference-based recall computes Tanimoto similarity (mathematically equivalent to normalized dot product for binary vectors) across the entire library in parallel.

A wafer-scale SEM array containing 10 million rods, each encoding a 2,048-bit fingerprint, could screen the entire 1.2-billion-molecule library (partitioned across 120 arrays operating in parallel) in under one second. The chemist submits her query and has analogs ranked by similarity before she finishes her coffee. At 15 fJ/bit , the entire search consumes less energy than the LED on her monitor.

What we've validated. The dot-product equivalence of interference recall (Section 2.3), the Hopfield pattern capacity, and the energy scaling (Section 6.4). **What remains projected:** Efficient encoding of 2,048-bit fingerprints into spectral perturbation patterns, and wafer-scale integration.

Relevant to: Pfizer, Roche, Novartis, Merck, AstraZeneca; CROs (Charles River, WuXi); cheminformatics platforms (Schrodinger, OpenEye).

C.4 Autonomous Vehicles: The Car That Recognizes Everything It Has Ever Seen

The problem. A Level 4 autonomous vehicle generates 1–4 TB of sensor data per hour from cameras, LiDAR, and radar. The perception stack must match incoming sensor frames against a library of known objects, road geometries, and hazard signatures—in real time, at 30 Hz, with <100 ms end-to-end latency. Today, this runs on an NVIDIA DRIVE Orin (275

$$\text{rods per cm}^3 = \left(\frac{0.01}{p} \right)^2 \cdot \frac{0.01}{s}$$

Array density = rods per $\text{cm}^3 \times$ bits per rod.

B.2 Thermoelastic Damping

Zener/Lifshitz-Roukes formulation [19] with Debye relaxation time $\tau_D = d^2/(\pi^2\kappa)$, where κ is thermal diffusivity:

$$\frac{1}{Q_{\text{TED}}} = \frac{E\alpha^2 T}{\rho C_p} \cdot \frac{\omega \tau_D}{1 + (\omega \tau_D)^2}$$

For glass at MEMS frequencies ($\sim$ MHz), $\omega \tau_D \gg 1$, and TED is negligible ($Q_{\text{TED}} \sim 10^7\text{--}10^8$).

B.3 Surface Loss

For a rod with diameter d and surface defect layer of thickness δ with quality factor Q_d :

$$\frac{1}{Q_{\text{surface}}} = \frac{4\delta}{d} \cdot \frac{1}{Q_d}$$

With $\delta = 5\ \text{nm}$ and $Q_d = 1,000$: $Q_{\text{surface}} = 196,000$ at $d = 40\ \mu\text{m}$.

TOPS, 60 W) or multiple chips in parallel. The power budget competes directly with vehicle range: every watt spent on compute is a watt not driving the wheels.

What SEM changes. SEM's interference-based associative recall is a native pattern matcher—it does not need to be programmed with neural network weights; it matches raw spectral signatures against stored templates. A 1 cm^3 SEM module containing 140,000 rods could store 182 million object signatures (Section 8.1) and match an incoming sensor frame against all of them in $5\ \mu\text{s}$ — $200\times$ faster than the 30 Hz frame rate requires, leaving margin for redundancy and voting. Power draw: under 2 W, compared to 60+ W for the GPU-based stack.

For fleet operators like Waymo or Cruise, multiplying this saving across thousands of vehicles translates to meaningful range extension and thermal-management simplification. For ADAS suppliers like Mobileye, a 2 W pattern-matching co-processor could enable Level 2+ features in vehicles where the thermal budget cannot accommodate a GPU.

What we've validated. Array capacity, associative recall time, and power (Sections 8–10). **What remains projected:** Sensor-to-spectral encoding pipeline and real-time integration with perception stacks.

Relevant to: Waymo, Cruise, Mobileye, NVIDIA, Qualcomm, Tesla, Bosch, Continental, Denso.

C.5 Financial Services: Fraud at the Speed of a Swipe

The problem. Visa processes $\sim 76,000$ transactions per second at peak (e.g., Black Friday). Each transaction must be compared against fraud patterns—card-present anomalies, velocity checks, merchant-category mismatches—within the 100 ms authorization window. Today, this requires a geographically distributed network of GPU-accelerated scoring engines, each drawing hundreds of kilowatts, with total infrastructure costs exceeding \$500 million annually.

What SEM changes. Each transaction can be encoded as a short feature vector (merchant category, amount, time-of-day, location, velocity). A SEM module scores a transaction against the full fraud-signature library in $5\ \mu\text{s}$ — $20,000\times$ faster than the 100 ms window. A single rack-mounted SEM appliance could replace an entire data-center floor of scoring servers.

But the deeper value is in *new fraud patterns*. Today, adding a new fraud signature requires retraining a neural network and re-deploying across the scoring fleet—a process measured in days. With SEM, a new fraud pattern is a new perturbation mask written to a rod. Deployment is physical: swap a chip, or write a new perturbation in the field. The time from pattern discovery to production deployment drops from days to minutes.

What we've validated. Recall speed, power, and pattern capacity. **What remains projected:** Transaction-to-spectral encoding standards and integration with existing payment network protocols.

Relevant to: Visa, Mastercard, JPMorgan Chase, Goldman Sachs, Stripe; fraud-detection vendors (Featurespace, Feedzai, FICO).

C.6 Agriculture and Food Safety: Every Grape, Every Grain

The problem. A food safety inspector at a grain processing facility needs to check incoming shipments for contamination— aflatoxins, pesticide residues, foreign materials. Today, this requires laboratory analysis (gas chromatography, mass spectrometry) with 24–72 hour turnaround. Hyperspectral imaging can identify contaminants optically, but the pattern-matching compute required to classify spectra in real time exceeds what is available in a handheld device.

What SEM changes. A hyperspectral camera captures the reflectance spectrum of each grain kernel—a high-dimensional vector. A battery-powered SEM device the size of a smartphone could store 10,000 contaminant spectral signatures and match each kernel's spectrum against the full library in $5\ \mu\text{s}$. At 15 fJ/bit and under 100 mW total power, the device runs for 8+ hours on a small LiPo cell.

An inspector walks the receiving dock with the device, scans a handful of grain, and gets a contamination verdict in real time—no lab, no wait, no shipment held in quarantine while samples are couriered. For a commodity trader, the ability to verify quality at the point of sale changes the risk calculus entirely.

What we've validated. The spectral correlation physics—SEM literally matches spectra against stored templates, which is precisely what hyperspectral classification requires. **What remains projected:** Miniaturized hyperspectral front-end integration and field ruggedization.

Relevant to: Cargill, ADM, Bunge, Tyson; food safety agencies (FDA, USDA, EFSA); precision agriculture (John Deere, CNH Industrial); hyperspectral imaging vendors (Headwall, Specim).

C.7 Space and Satellite: Memory That Survives the Van Allen Belts

The problem. The James Webb Space Telescope carries 68 GB of solid-state memory using radiation-hardened SRAM. That memory cost more per bit than the gold-plated mirrors. Every satellite, every Mars rover, every deep-space probe faces the same constraint: radiation-hardened memory is scarce, expensive (1,000–10,000 per Mbit), heavy (shielding adds mass), and limited in capacity. A single energetic proton can flip a bit in a transistor; a solar particle event can corrupt entire memory banks.

What SEM changes. SEM stores information in mechanical vibrations of glass—a medium that is intrinsically immune to ionizing radiation. There are no charge-trapped states to flip, no floating gates to discharge, no magnetic domains to disturb. A cosmic ray passing through a glass rod changes its acoustic properties by less than one part in 10^{12} . The information is encoded in the rod's *geometry* (perturbation pattern) and *physics* (eigenmode spectrum), not in electrical charge.

A 1 cm^3 SEM module provides 17 Gbit (2.1 GB) of radiation-hard memory— $250\times$ the capacity of JWST's entire memory system—in a package that weighs grams instead of kilograms, costs orders of magnitude less, and simultaneously performs associative computation for onboard pattern recognition (autonomous navigation, anomaly detection, spectral classification of planetary surfaces).

For the emerging commercial space industry—where companies like SpaceX, Rocket Lab, and Planet Labs are launching thousands of small satellites—SEM could replace radiation-hard memory as the default, removing one of the most expensive line items from satellite BOM.

What we’ve validated. The glass medium, the physics of eigenmode encoding, and the density scaling. **What remains projected:** Vibration tolerance in launch environments (MEMS resonators routinely survive 10,000 g shock tests, suggesting this is tractable) and space-qualification testing.

Relevant to: NASA, ESA, JAXA; SpaceX, Rocket Lab, Planet Labs, Capella Space; defense satellite primes (Lockheed Martin, L3Harris, Ball Aerospace), JPL, APL..

C.8 Summary: Where SEM Fits

Scenario	Incumbent	Projected advantage (theoretical)	Readiness gap
Signals intelligence	TCAM racks	10,000× latency, 2,400× power	MEMS fab + interface
Fab defect inspection	GPU clusters	6,000× latency, 30,000× power	Image-to-spectral encoding
Drug discovery	HPC clusters	50,000× latency, 150,000× power	Wafer-scale integration
Autonomous vehicles	NVIDIA Orin	200× latency, 30× power	Sensor encoding pipeline
Fraud detection	GPU scoring fleet	20,000× latency, massive power	Transaction encoding standard
Food safety	Lab spectrometry	Real-time, handheld, battery-powered	Hyperspectral front-end
Space computing	Rad-hard SRAM	250× capacity, intrinsic rad-hard	Space qualification

APPENDIX D: MACRO-SCALE EXPERIMENT GUIDE

This appendix provides step-by-step instructions for replicating the macro-scale prototype experiments of Section 4. The guide is designed to be self-contained: every component is listed with a direct purchase link, every procedure is numbered for reproducibility, and every known failure mode includes a tested mitigation. A middle school science teacher with no acoustics background should be able to build the prototype, complete all eight experiments, and contribute publishable data within a single school week. See Section 4 for the theoretical context behind each measurement.

The glass harmonica connection. Every experiment in this appendix is a direct descendant of a musical instrument that pre-dates the transistor by centuries. A glass harmonica—tuned wineglasses played by rubbing a wet finger around the rim—demonstrates every principle of SEM in audible form: glass resonators with eigenfrequencies set by geometry, mass perturbation tuning via water level, continuous-wave excitation via stick-slip friction, and spectral readout by the human ear. In 1761, Benjamin Franklin attended a glass harmonica concert in London and built an improved version: the *glass armonica*, which mounted the bowls on a rotating spindle so a performer could vary finger pressure, position, and contact duration in real time. Same glass, same physics, same resonant modes—but now reconfigurable. This appendix walks the same path: Experiments 1–6 build and characterize a fixed resonator (the harmonica); Experiment 7 demonstrates continuous-wave precision readout (bowing vs. ringing); and Experiment 8 demonstrates rewritable encoding with water drops (the armonica).

D.1 Complete Bill of Materials

The core BOM from Section 4.2 is expanded below with recommended quantities (extras for breakage and controls), supplier links, and additional items needed for failure-mode mitigations. All prices are approximate as of 2026 and may vary.

Table D.1: Core Components

#	Component	Specification	Qty	Est. Cost	Amazon Link
1	Borosilicate glass stirring rods	6 mm dia × 150 mm (5.9"), rounded ends, Boro 3.3	15-pack	~\$8	PATIKIL 15 Pcs, 5.9" × 6 mm
2	Borosilicate glass stirring rods (alt.)	6 mm dia × 200 mm (7.9"), Boro 3.3—cut to 150 mm if needed	10-pack	~\$9	EISCO 10PK, 7.9" × 6 mm
3	Piezoelectric discs with leads	10 mm dia, PZT ceramic, pre-soldered 4" wire leads	15-pack	~\$7	E-outstanding 15 PCS, 10 mm
4	Piezoelectric discs (alt.)	10 mm dia, PZT, bare discs (solder leads yourself)	10-pack	~\$6	uxcell 10 Pcs, 10 mm
5	USB oscilloscope + waveform generator	PicoScope 2204A, 10 MHz BW, 2-ch, built-in AWG, PS7 software	1	~\$25	PicoScope 2204A
6	Cyanoacrylate glue (super glue)	Medium viscosity, precision tip, gel formula	1	~\$5	Loctite Ultra Gel Control, 4 g
7	Inlay casting wax	Brown dental/jewelry sticky wax sticks, pliable at room temp	1 box	~\$8	Brown Inlay Wax, 1 oz
8	BNC cables (male–male, 50 Ω)	1 m length, for Picoscope connection	2	~\$10	Included with PicoScope kit; extras available at any electronics supplier
Core subtotal				~63–78	

Table D.2: Mitigation and Measurement Accessories

The pattern across all seven scenarios is consistent: SEM’s interference-based recall is a natural fit for workloads that reduce to nearest-neighbor search—problems that today require racks of silicon but map directly to acoustic wave physics. The advantage multipliers in the table above are theoretical, derived by comparing projected SEM performance (Sections 6–8) against current production systems. A fabricated MEMS device will inevitably fall short of these projections in ways that only measurement can reveal. What the scenarios establish is not a guaranteed performance advantage but a *structural* one: SEM’s physics natively performs the operation these workloads require, whereas conventional architectures must emulate it in software. The advanced encoding techniques of Section 11 expand the design space further, but their real-de-vice performance remains to be validated.

#	Component	Purpose (failure mode addressed)	Qty	Est. Cost	Source
9	Craft foam sheets (EVA, 6 mm thick)	Vibration-isolating cradle (FM 1: anchor loss)	1 pack	~\$6	Any craft store or Amazon
10	Digital thermometer, 0.1 °C resolution	Thermal drift monitoring (FM 3)	1	~\$12	Any kitchen/lab thermometer with 0.1 °C readout
11	Styrofoam cooler, small (~6-qt)	Thermal enclosure (FM 3: drift isolation)	1	~\$5	Grocery or hardware store
12	Milligram precision scale (0.001 g)	Weighing wax perturbation masses	1	~\$20	Amazon search: “milligram scale 0.001g”
13	Metric ruler, mm scale	Positioning masses along rod	1	~\$3	Any school supply
14	Isopropyl alcohol, 91 %+	Cleaning glass rods before assembly	1 bottle	~\$4	Any pharmacy
15	Safety glasses	Eye protection (glass rods can snap)	1 per person	~\$3	Any hardware store
16	Masking tape	PZT alignment jig (FM 2: centering)	1 roll	~\$4	Any hardware store
17	Fine-tip permanent marker	Marking positions on rod	1	~\$2	Any office supply
18	Soft lint-free cloth	Handling and cleaning glass	several	~\$2	Any store
19	Acetone (nail polish remover)	Emergency super-glue skin-bond release	1 small bottle	~\$3	Any pharmacy
20	Plastic transfer pipettes (3 mL)	Placing water drops for Exp. 8 (rewritability)	1 pack	~\$4	Amazon search: “plastic transfer pipettes”
21	Small bowl of water	CW excitation via wet finger (Exp. 7)	—	free	Tap water
Accessories subtotal				~\$68	
Grand total				~131–146	

Budget note. The PicoScope 2204A (25) is the single most important purchase—it provides both the wave form generator (transmit) and the digi 40. The 15 glass rods and 15 PZT discs provide enough spares for multiple student groups, breakage, and control experiments. One kit serves an entire class.

D.2 Safety Notes

- Glass hazard.** Borosilicate rods can snap if bent or dropped. Always wear safety glasses when handling rods. Dispose of broken glass in a rigid sharps container, never a trash bag.
- Cyanoacrylate (super glue).** Bonds skin instantly. Keep acetone (nail polish remover) on hand to dissolve accidental skin bonds. Work in a ventilated area. Supervise younger students closely during the gluing step.
- Hearing.** The fundamental mode (17.7 kHz) is near the upper limit of human hearing. Some students may hear a faint whine during high-amplitude excitation. This is harmless at the drive levels used here (<1 V), but if anyone reports discomfort, reduce the AWG amplitude to 0.1 V.

D.3 Experiment 1 — Building the Resonator

Objective: Assemble a functioning glass-rod acoustic resonator and verify electrical connectivity.

Time: 30 minutes active + 24 hours glue cure.

Materials: Items 1 or 2, 3 or 4, 6, 9, 16, 18 from the BOM.

Procedure:

- Clean the rod.** Wipe one glass rod with isopropyl alcohol and a soft cloth. Allow 2 minutes to dry completely. Finger oils dampen vibrations measurably.
- Build the foam cradle at the displacement nodes.** Cut two V-shaped notches in a strip of craft foam, spaced 75 mm apart, centered on the rod—this places each notch at $L/4$ and $3L/4$ from one end (37.5 mm and 112.5 mm for a 150 mm rod). These positions are the exact displacement nodes of the second longitudinal mode—the acoustic “stems” of the rod. The rod should rest horizontally in the notches with no hard contact. A wine glass rings because you hold it by the stem, a vibrational node where energy cannot escape; the same physics governs your rod mount (see Failure Mode 6 and Section 7).
- Center the PZT disc (critical).** Cut two small strips of masking tape (~12 mm each). Adhere them in a cross-hair pattern centered on the flat end face of the rod. The intersection marks the center of the 6

- mm face—this is where the PZT must go. An off-center disc excites transverse and torsional modes that pollute the spectrum (Failure Mode 2).
- Apply glue sparingly.** Place one tiny drop of cyanoacrylate—smaller than a pinhead, less than 0.5 mm in diameter—at the center of the cross-hair. *Less is more:* excess glue adds mass and viscoelastic damping that destroys the quality factor (Failure Mode 1).
 - Attach the PZT disc.** Press the flat face of the PZT disc onto the glued spot, centering it on the cross-hair. Hold firm, even pressure for 30 seconds. Gently peel away the tape cross-hair strips.
 - Cure.** Set the assembly aside on the foam cradle for a full 24 hours. Cyanoacrylate reaches full bond strength overnight; rushing produces a weak, lossy joint.
 - Connectivity check.** Connect the PZT leads to the PicoScope Channel A input via a BNC adapter or clip leads. Set the scope to AC coupling, 1 mV/div, 1 ms/div timebase. Gently tap the free end of the rod with a fingernail. You should see a decaying burst of oscillation on screen. If no signal appears: check wire connections, ensure the PZT is not cracked, and try a firmer tap.

Failure Mode 1 — PZT/Epoxy Boundary Damping (Low Q)

The glue joint and PZT disc add mass and internal friction at the rod end, potentially reducing Q from the intrinsic material value (~10,000) to below 100 if done carelessly.

- Use the *absolute minimum* glue: one drop smaller than a pinhead.
- Ensure the cured glue layer is thin (<0.1 mm). A thick glue layer acts as a lossy viscoelastic coupling that absorbs acoustic energy each cycle.
- Build a second “control” rod with *no* PZT attached—excite it by tap only and record its ring-down via a nearby microphone or laser vibrometer (if available). Comparing Q values between the two rods isolates the damping contribution of the PZT joint.
- If Q is unacceptably low (<500 after Experiment 2), scrape off the PZT with a razor blade, clean the end face with alcohol, and re-glue with a smaller drop.

Failure Mode 6 — Support-Point Energy Drain (the Glass Harp Lesson)

A wine glass rings brilliantly because you hold it by the stem—a point of zero vibrational displacement (a node) for the rim’s flexural modes. Energy cannot leak through a point that isn’t moving. Press your finger against the rim itself and the tone dies instantly: the support is now at an antinode, and your finger becomes an acoustic drain.

The same physics governs your rod mount. For a free-free rod vibrating in its n -th longitudinal mode, the displacement pattern is $u(x) = \cos(n\pi x/L)$. The displacement nodes—the “stems”—occur at:

Mode	Node positions
$n = 1$	$L/2$ (center)
$n = 2$	$L/4$ and $3L/4$
$n = 3$	$L/6$, $L/2$, and $5L/6$

The recommended cradle positions at $L/4$ and $3L/4$ (step 2) are the exact nodes of mode 2, and they produce only 50% of peak displacement for mode 1—a good all-round compromise for multi-mode measurements.

If your Q values are lower than expected despite a clean PZT bond:

- Verify cradle positions.** Measure from one end: 37.5 mm and 112.5 mm for a 150 mm rod. Mark the positions with a fine-tip marker.
- Try midpoint mounting.** For mode-1-only measurements, a single support at $L/2 = 75$ mm is the ideal node and will maximize Q for the fundamental. (This sacrifices mode 2, which has an antinode there.)
- Upgrade to fishing line.** Loop a thin monofilament line around the rod at the node position and tension it between two fixed posts. This gives near-zero contact area and mimics the knife-edge mounts used in precision metrology. Students who have played a glass harp will immediately feel the analogy: the taut line is the stem.
- Diagnostic test.** Slide the foam cradle to the rod’s center ($L/2$) and measure Q for mode 1. Then slide it to $L/4$ and remeasure. If Q changes by more than 20%, your foam is too stiff or making hard contact—switch to softer foam, adjust the V-notch depth, or use fishing line.

D.4 Experiment 2 — Measuring the Quality Factor

Objective: Determine Q via two independent methods (ring-down and bandwidth) and confirm the resonator is material-loss-limited.

Time: 45 minutes.

Materials: Assembled resonator from Exp. 1, PicoScope 2204A, BNC cable, foam cradle.

Procedure:

- Set up.** Place the resonator on the foam cradle on a stable surface (no vibration from HVAC, foot traffic, etc.). Connect the PZT leads to both the PicoScope AWG output (waveform generator) and Channel A input. A BNC T-connector works, or simply share the two PZT leads between the AWG and scope clips.
- Impulse excitation.** In the PicoScope software (PS7, free download), configure the AWG to output a single 1-cycle burst at 17,700 Hz with 0.5 V amplitude. Set Channel A to trigger on the rising edge. Set the timebase to 50 ms/div so the display shows ~500 ms of data.
- Capture the ring-down.** Trigger the burst and record the decaying sinusoidal waveform. The oscillation frequency should be near 17.7 kHz; the amplitude should decay smoothly.
- Measure τ (time constant).** The envelope follows $A(t) = A_0 e^{-t/\tau}$. Identify the initial peak amplitude A_0 , then find the time t at which the envelope has dropped to $A_0/e \approx 0.368 \times A_0$. That time is τ . Alternatively, use PicoScope’s cursor measurements to read the 1/e point directly.
- Calculate Q from ring-down:**

$$Q_{\text{ringdown}} = \pi f_1 \tau$$

For example: $f_1 = 17,700$ Hz, $\tau = 180$ ms gives $Q = \pi \times 17,700 \times 0.180 = 10,006$.

- Cross-check via bandwidth method.** Switch PicoScope to spectrum (FFT) mode. Drive the rod with a slow linear frequency sweep (chirp) from 17,000 Hz to 18,500 Hz over 2 seconds at 0.2 V amplitude. In the FFT view, locate the resonance peak. Measure the full width at half-maximum power (the -3 dB bandwidth, $\Delta f_{3\text{dB}}$ —the frequency span where the peak drops to 70.7% of its maximum amplitude, or -3 dB from the peak). Calculate:

$$Q_{\text{bandwidth}} = \frac{f_1}{\Delta f_{3\text{dB}}}$$

- Compare the two Q values.** They should agree within 10%. Record all results in Worksheet D.1.

Failure Mode 4 — Piezoelectric Self-Resonance Masking

A 10 mm PZT disc has its own radial resonance, typically at 200–300 kHz—well above the glass rod’s fundamental (17.7 kHz). However, some disc types have a thickness resonance at lower frequencies that can overlap with glass-rod harmonics.

To identify PZT artifacts:

- Before* gluing (during your next build), hold a bare PZT disc by its leads in free air and tap it gently. Record the FFT spectrum—these are PZT-only resonances.
- After assembly, compare the full spectrum against the PZT-only spectrum. Any peak that appears in both is a PZT artifact, not a glass eigenmode.
- True glass-rod longitudinal modes are evenly spaced at $\Delta f = v_{\text{bar}}/(2L) \approx 17,700$ Hz. PZT resonances do *not* follow this comb pattern. Any peak that falls off the comb by more than 1% is suspect.

Worksheet D.1 — Quality Factor Measurement

Parameter	Rod 1	Rod 2	Rod 3
Date / Time	15 Mar 2026, 14:30		
Experimenter	A. Student		
Room temperature (°C)	22.3		
Rod length L (mm)	150.2		
Rod diameter d (mm)	6.01		
Glue amount (tiny / small / medium)	Tiny		
Measured f_1 (Hz)	17,693		
Predicted $f_1 = 5,315/(2L)$ (Hz)	17,699		
Ring-down τ (ms)	168		
Q (ring-down) = $\pi f_1 \tau$	9,335		
-3 dB bandwidth $\Delta f_{3\text{dB}}$ (Hz)	1.9		
Q (bandwidth) = $f_1/\Delta f_{3\text{dB}}$	9,312		
Two methods agree within 10%? (Y/N)	Y✓		
Notes	Clean bond, thin glue		

Expected results. Q should fall in the range 1,000–10,000. Values near 10,000 indicate excellent construction (material-loss-limited). Values below 500 suggest excessive glue, poor PZT centering, or a hard contact point on the cradle—rebuild with corrections per the Failure Mode 1 mitigation above. A control rod with no PZT (tap-excited) should yield Q near 10,000, confirming the intrinsic glass value.

D.5 Experiment 3 — Mapping the Mode Spectrum

Objective: Identify the longitudinal eigenmode comb and distinguish real modes from transverse/torsional artifacts.

Time: 45 minutes.

Materials: Assembled resonator, PicoScope, BNC cable, foam cradle.

Procedure:

- Configure the chirp.** Set the PicoScope AWG to output a linear frequency sweep from 1 kHz to 200 kHz over 1 second at 0.5 V amplitude.
- Capture the response.** Record Channel A during the chirp with FFT mode enabled. Use a Hanning window and at least 16,384 FFT points for ~1 Hz frequency resolution.
- Identify the mode comb.** Longitudinal modes appear as evenly spaced peaks at:

Mode n	Predicted frequency
1	17,717 Hz
2	35,434 Hz
3	53,150 Hz
4	70,867 Hz
5	88,584 Hz
6	106,301 Hz
7	124,018 Hz
n	$n \times 5,315/(2L)$ Hz

- Mark confirmed modes.** On the FFT display (or in a screenshot), mark each peak that falls within $\pm 0.5\%$ of a predicted comb frequency. These are your confirmed longitudinal eigenmodes.
- Catalog spurious peaks.** Any peaks that do *not* fall on the comb are either: (a) transverse or torsional modes (Failure Mode 2), (b) PZT self-resonances (Failure Mode 4), or (c) electrical noise (harmonics of 60 Hz mains). Note their frequencies.
- Count confirmed modes** above the noise floor and record in Worksheet D.2.

⚙️ Failure Mode 2 — Transverse/Torsional Spectrum Pollution

If the PZT disc is not centered on the rod axis, it applies an off-axis force that excites bending and torsional modes in addition to the desired longitudinal modes. These appear as extra peaks *between* the longitudinal comb lines.

Mitigations:

- **Verify centering.** If spurious peaks are within 10 dB of the longitudinal peaks, remove the PZT and re-glue with better centering (see Exp. 1, step 3).
- **Software filter.** Export the FFT data (PicoScope can save .csv files) and discard all peaks more than ±500 Hz from the predicted comb frequencies.
- **Rotation test.** Rotate the rod 90° on the foam cradle and re-measure. Longitudinal modes (which depend only on length) will not shift. Transverse modes (which depend on the rod’s cross-sectional geometry) will shift because the cross-section is never perfectly circular. Any peak that moves under rotation is not longitudinal.

Worksheet D.2 — Mode Spectrum Map

Mode <i>n</i>	Predicted <i>f_n</i> (Hz)	Measured <i>f_n</i> (Hz)	Δ <i>f</i> = meas – pred (Hz)	Amplitude (dB)	Confirmed?
1	17,717	17,693	−24	−22	Y
2	35,434	35,388	−46	−28	Y
3	53,150	53,081	−69	−35	Y
4	70,867	70,775	−92	−43	Y
5	88,584	88,466	−118	−52	Y
6	106,301	106,160	−141	−63	Y
7	124,018				Y / N
Spurious 1	—	24,330	—	−58	type: <i>transverse</i>
Spurious 2	—		—		type: <i>—</i>
Spurious 3	—		—		type: <i>—</i>
Total confirmed longitudinal modes:		6			

Expected results. 5–10 clean longitudinal modes should be visible within the PicoScope’s 10 MHz bandwidth. The first 2–3 modes will have the highest amplitude. Measured frequencies should match predictions within ±0.5%. If your rod is 200 mm instead of 150 mm, recalculate: $f_1 = 5,315 / (2 \times 0.200) = 13,288$ Hz.

D.6 Experiment 4 — Thermal Stability Characterization

Objective: Quantify how room-temperature fluctuations shift mode frequencies, and establish a stable measurement protocol for perturbation experiments.

Time: 90 minutes (mostly waiting for equilibration).

Materials: Assembled resonator, PicoScope, digital thermometer, styrofoam cooler, foam cradle.

Background. Borosilicate glass has a thermal expansion coefficient of $\sim 3.3 \times 10^{-6} \text{ }^\circ\text{C}^{-1}$ and a temperature coefficient of elastic modulus of $\sim 100 \text{ ppm}/^\circ\text{C}$. The net frequency sensitivity is approximately:

$$\frac{\Delta f}{f} \approx -50 \text{ ppm}/^\circ\text{C}$$

For $f_1 = 17,700$ Hz, this predicts $\Delta f \approx -0.9$ Hz per $^\circ\text{C}$ of temperature change.

Procedure:

1. **Open-air baseline.** Place the resonator on the foam cradle on a lab bench. Place the thermometer probe within 2 cm of the rod. Every 5 minutes for 30 minutes, record f_1 (using the bandwidth method from Exp. 2—a quick FFT snapshot) and the temperature. Do not touch the setup or breathe on it.
2. **Thermal perturbation.** Breathe warm air directly onto the rod for 10 seconds, or place a cup of warm water ($\sim 50^\circ\text{C}$) 20 cm away. Resume recording f_1 and temperature every 2 minutes until the frequency returns to within 0.5 Hz of its initial value. Note the recovery time.
3. **Insulated baseline.** Place the entire resonator assembly (foam cradle, rod, PZT, and thermometer probe) inside the styrofoam cooler. Thread the BNC cable and thermometer wire through a small notch cut in the cooler lid. Close the lid. Repeat the 30-minute baseline: record f_1 and temperature every 5 minutes.
4. **Calculate the temperature coefficient.** Plot f_1 versus temperature for both the open-air and insulated datasets. The slope $\Delta f / \Delta T$ is your measured temperature coefficient. Compare to the predicted $-0.9 \text{ Hz}/^\circ\text{C}$.
5. **Record results** in Worksheet D.3.

⚙️ Failure Mode 3 — Thermal Drift Swamping Perturbation Encoding

A 1°C temperature change shifts modes by $\sim 0.9 \text{ Hz}$ —comparable to a small wax perturbation. If thermal drift is not controlled, perturbation signals drown in noise.

Mitigations:

- **Always use the styrofoam enclosure** for Experiments 5 and 6. This alone reduces drift by 10–50×.
- **Equilibrate 30 minutes** before starting perturbation experiments.
- **Use differential measurement.** Thermal drift shifts *all* modes by the same fractional amount ($\sim 50 \text{ ppm}/^\circ\text{C}$). A localized mass perturbation shifts each mode by a *different* amount (proportional to $\sin^2(n\pi x/L)$). Track the *pattern* of relative shifts, not absolute frequencies. This is your strongest discriminant.
- **Bracket measurements.** Take a “before” spectrum immediately before applying wax, and an “after” spectrum within 60 seconds. Over such short intervals, thermal drift is negligible.
- **Record temperature at every measurement point** so post-hoc correction is possible.

Worksheet D.3 — Thermal Stability Log

Time (min)	Temp ($^\circ\text{C}$)	<i>f</i> ₁ (Hz)	Δ <i>f</i> from <i>t</i> = 0 (Hz)	Environment
0	22.3	17,693.2	0.0	Insulated
5	22.3	17,693.1	−0.1	Insulated
10	22.4	17,693.0	−0.2	Insulated
15	22.4	17,692.9	−0.3	Insulated
20	22.4	17,693.0	−0.2	Insulated
25	22.4	17,693.0	−0.2	Insulated
30	22.4	17,693.0	−0.2	Insulated
(thermal perturbation applied)		hand near enclosure 30 s		
32	22.8	17,692.6	−0.6	Open
34	23.1	17,692.3	−0.9	Open
36	23.0	17,692.4	−0.8	Open
38	22.8	17,692.6	−0.6	Open
40	22.6	17,692.8	−0.4	Open
Measured Δ <i>f</i> /Δ <i>T</i> :		−0.85 Hz/ $^\circ\text{C}$		
Predicted (−0.9 Hz/ $^\circ\text{C}$):		agrees within 6%		
Recovery time after perturbation:		~8 min		

Expected results. The open-air drift rate will depend on your HVAC system—typically 0.5–5 Hz over 30 minutes. Inside the styrofoam enclosure, drift should be ~ 0.5 Hz over 30 minutes, which is small enough to resolve wax perturbation shifts of 1–10 Hz.

D.7 Experiment 5 — Perturbation Encoding

Objective: Write data to the rod by applying wax masses and verify that measured frequency shifts match Rayleigh perturbation theory.

Time: 60 minutes.

Materials: Assembled resonator (inside styrofoam enclosure), PicoScope, wax, milligram scale, ruler, fine-tip marker.

Background. The Rayleigh perturbation formula predicts that a small mass δm placed at position x along a rod of total mass M and length L shifts the n -th mode frequency by:

$$\frac{\Delta f_n}{f_n} = -\frac{\delta m}{2M} \sin^2\left(\frac{n\pi x}{L}\right)$$

The key insight is that *different modes shift by different amounts* depending on where the mass sits relative to each mode’s antinodes. This position-dependent pattern of shifts is the “spectral fingerprint”—the basis of data encoding.

Procedure:

1. **Prepare wax pellets.** Pinch off small pieces of inlay wax and roll them into balls approximately 1 mm in diameter. Weigh each on the milligram scale and record the mass (target: 0.05–0.5 mg). Prepare 3–5 pellets of varied sizes.
2. **Mark positions.** Using the ruler and fine-tip marker, mark reference positions on the rod at 25%, 33%, 50%, 67%, and 75% of the rod length from the PZT end (i.e., at 37.5, 50.0, 75.0, 100.0, and 112.5 mm for a 150 mm rod).
3. **Record the unperturbed spectrum.** Drive the rod with a chirp and capture the FFT. Record the frequencies of modes 1–5 in Worksheet D.4 under “Before.”
4. **Apply one wax pellet.** Press it onto the rod at a marked position x . The wax should adhere by its own stickiness—no glue needed. Record the pellet mass and position.
5. **Record the perturbed spectrum.** Immediately re-drive and capture the FFT. Record mode frequencies under “After.” Calculate the shift Δf_n for each mode.
6. **Compare to theory.** Calculate the predicted $\Delta f_n / f_n$ for each mode using the Rayleigh formula above. Weigh the rod itself to get M (typically ~ 6 g for a 150 mm $\times$ 6 mm borosilicate rod, density 2,230 kg/m³).
7. **Repeat with different positions.** Move the wax to a different marked position and re-measure. The fingerprint should change: modes whose antinodes are near the new wax position shift most; modes with nodes near the wax barely shift.
8. **Verify reversibility.** Remove the wax with a fingernail and re-measure the spectrum. Frequencies should return to the unperturbed values within the thermal drift tolerance established in Experiment 4.

⚙️ Failure Mode 5 — Non-Linear Acoustic Coupling

At high drive amplitudes, acoustic energy leaks between modes via non-linear elastic coupling, smearing the spectral fingerprint and reducing discrimination.

Mitigations:

- **Keep drive amplitude $\leq 0.5 \text{ Vpp}$.** This keeps the rod well within the linear regime.
- **Verify linearity.** Measure mode amplitudes at 0.25 V, 0.50 V, and 1.00 V. Plot amplitude vs. drive voltage. The relationship should be linear (doubling voltage doubles response amplitude). If doubling drive voltage increases any mode’s amplitude by more than 2.2× (exceeding the linear 2.0× by more than 10%), you have entered the non-linear regime—reduce the drive.
- **Listen for audible ringing.** If you can clearly hear the rod vibrating (a high-pitched whine), the amplitude is too high. Reduce AWG output until the sound is barely perceptible or inaudible.

Worksheet D.4 — Perturbation Encoding Data

Copy this worksheet for each trial.

Trial info	
Trial #	1
Wax pellet mass δm (mg)	0.8
Position x (mm from PZT end)	75.0
Position x/L (fraction)	0.500
Rod mass M (g)	9.52
Rod length L (mm)	150.2
Temperature ($^{\circ}\text{C}$)	22.4

Mode n	f_n before (Hz)	f_n after (Hz)	Δf_n measured (Hz)	$\Delta f_n/f_n$ meas. (ppm)	$\Delta f_n/f_n$ predicted (ppm)	Error (%)
1	17,693.2	17,692.5	-0.7	-40	-42	5
2	35,388.1	35,388.0	-0.1	-3	0	-
3	53,081.0	53,078.8	-2.2	-41	-42	2
4	70,775.3	70,775.2	-0.1	-1	0	-
5	88,466.1	88,462.4	-3.7	-42	-42	1

After removing wax:

Mode n	f_n recovered (Hz)	Δf from original (Hz)	Recovered within ± 0.5 Hz?
1	17,693.1	-0.1	Y
2	35,388.0	-0.1	Y
3	53,080.9	-0.1	Y
4	70,775.2	-0.1	Y
5	88,466.0	-0.1	Y

Expected results. Mode shifts should match Rayleigh predictions within 5%. The mode whose antinode is nearest the wax position will shift most; modes with nodes near the wax will barely shift. Frequencies should recover fully after wax removal. If the match is poor ($>10\%$ error), verify that the wax mass is small compared to the rod mass ($\delta m/M \ll 1$) and that the temperature has not drifted during the measurement.

D.8 Experiment 6 — Associative Recall Demonstration

Objective: Show that the rod physically distinguishes between matching and non-matching query patterns—the basis of SEM’s $O(1)$ nearest-neighbor search.

Time: 60 minutes.

Materials: Assembled resonator, PicoScope, wax, ruler, foam cradle, styrofoam enclosure.

Procedure:

- Create Pattern A.** Apply two wax pellets at the quarter-points of the rod: $x = L/4 = 37.5$ mm and $x = 3L/4 = 112.5$ mm. Record the resulting spectral fingerprint—the frequencies of modes 1–5 in their shifted positions.
- Build the query signal.** In the PicoScope AWG, create a multi-tone waveform composed of the five shifted frequencies from Pattern A’s fingerprint. Set each tone to equal amplitude (0.1 V per tone). This is “Query A.”
- Measure the matched response.** Drive the rod with Query A. The rod contains Pattern A, so the query *matches* the stored pattern. Record the peak response amplitude in dB from the FFT.
- Create Pattern B.** Remove Pattern A’s wax. Apply two pellets at the third-points: $x = L/3 = 50$ mm and $x = 2L/3 = 100$ mm. Record Pattern B’s fingerprint.
- Drive with Query A (mismatched).** Replay Query A—the waveform built from Pattern A’s frequencies. The rod now contains Pattern B, so the query does *not* match. Record the peak response amplitude. It should be significantly lower than in step 3.
- Drive with Query B (matched).** Build a new multi-tone waveform from Pattern B’s frequencies. Drive the rod with this Query B. The rod now matches, so the response should be strong. Record the amplitude.
- Calculate the discrimination margin.** The difference in dB between the matched response and the best non-matched response is the discrimination margin. A margin ≥ 15 dB means the correct match produces $\geq 30\times$ more power than the closest incorrect match.
- Record results** in Worksheet D.5.

Worksheet D.5 — Associative Recall Discrimination

Pattern stored on rod	Query driven	Peak response (dB)	Match?
Pattern A	Query A (matching)	-22	✓
Pattern B	Query A (non-matching)	-44	✗
Pattern B	Query B (matching)	-23	✓

Metric	Value
Discrimination margin (matched – non-matched):	21 dB
Power ratio ($10^{(\text{margin}/10)}$):	126×
Sufficient for reliable detection? (≥ 15 dB)	Y✓

Expected results. The discrimination margin should be 15–25 dB, meaning the matching pattern produces 30–300× more acoustic power than a non-matching pattern. If the margin is below 10 dB, check that (a) the two wax patterns are placed

at genuinely different positions, (b) the drive amplitude is in the linear regime (≤ 0.5 V per tone), and (c) the query frequencies accurately match the stored pattern’s measured frequencies.

D.9 Experiment 7 — CW Precision Readout (Bowing vs. Ringing)

Objective: Demonstrate that continuous-wave (CW) excitation with lock-in detection yields higher SNR than impulse readout—and experience the difference physically by “bowing” the rod with a wet finger, just as a glass harmonica is played.

Time: 60 minutes.

Materials: Assembled resonator, PicoScope, BNC cable, foam cradle, styrofoam enclosure, small bowl of water.

Background. Experiments 1–6 use impulse readout: strike the rod, listen to it ring down, measure the FFT. This is fast (one ringdown time $\tau \approx 180$ ms) but limited in SNR—the measurement window is fixed at τ , and all the noise within the bandwidth $\sim 1/\tau$ contributes. An alternative is CW readout: drive the rod continuously at a single mode frequency and measure the steady-state response. A lock-in detection technique—multiplying the response by the drive signal and averaging—rejects all noise outside a narrow bandwidth $\sim 1/(2T_{\text{int}})$. The SNR gain over impulse is:

Gain (dB) = 10 log₁₀ ($\frac{T_{\text{int}}}{\tau}$)

At $T_{\text{int}} = 1$ s: +7.5 dB. At $T_{\text{int}} = 10$ s: +17.5 dB. At $T_{\text{int}} = 60$ s: +25.2 dB (see Figure 15). This is the same physics that makes a glass harmonica so expressive: the performer’s sustained finger contact provides continuous energy input, allowing the bowl to reach full amplitude and sustain it indefinitely—something a single tap cannot do.

Part A: Electronic CW Readout

- Set the CW drive.** Configure the PicoScope AWG to output a *continuous* sine wave at your rod’s measured f_1 from Experiment 2. Set amplitude to 0.2 V. Let it run—do not stop it.
- Wait for ring-up.** The rod’s amplitude builds exponentially with the same time constant τ as the ring-down. After $3\tau \approx 540$ ms, the rod has reached 95% of its steady-state amplitude. Wait at least 1 second.
- Capture a 1-second record.** In PicoScope, set the timebase to capture exactly 1 second of Channel A data. Save the raw waveform (.csv export).
- Software lock-in detection.** In a spreadsheet or Python script, perform the lock-in calculation:
- Generate a reference signal: $\text{ref}(t) = \sin(2\pi f_1 t)$.
- Multiply the captured signal by the reference: $\text{product}(t) = \text{signal}(t) \times \text{ref}(t)$.
- Average the product over the 1-second window. This is the lock-in output—proportional to the mode amplitude, with all out-of-band noise rejected.
- Record the lock-in amplitude as $A_{\text{CW},1s}$.
- Repeat with 10-second capture.** Extend the AWG run and capture 10 seconds. Repeat the lock-in calculation. Record $A_{\text{CW},10s}$.
- Compare to impulse.** From Experiment 2, you have the impulse ring-down peak amplitude A_{impulse} . Calculate the SNR improvement:

Gain (dB) = 20 log₁₀ ($\frac{A_{\text{CW}}}{A_{\text{impulse}}}$)

Expected: ~ 7.5 dB for 1-second CW, ~ 17.5 dB for 10-second CW.

Part B: Bowing the Rod (the Glass Harmonica Experiment)

This part requires no electronics—just your hands and ears (plus the PZT to record what happens).

- Wet your finger.** Dip one finger in the bowl of water. You want a thin, even film of moisture—not dripping wet.
- Bow the rod.** With the PicoScope recording on Channel A (long timebase, ~ 5 seconds), run your wet finger firmly and steadily along the length of the rod, maintaining even pressure. The finger creates stick-slip friction that excites the rod’s longitudinal modes, exactly as a wet finger on a wineglass rim excites its radial modes. You should hear a faint singing tone and see a sustained oscillation on the PicoScope screen.
- Compare the waveform.** The bowed response should show sustained, roughly constant-amplitude oscillation for as long as your finger maintains contact—in contrast to the exponentially decaying impulse ring-down from Experiment 2. *This is the glass harmonica in action.*
- Measure bowed frequency.** Switch to FFT mode while bowing. The dominant peak should fall at or near f_1 . The excitation is broadband (stick-slip friction contains many frequencies), so you may also see modes 2 and 3 responding.
- Record results** in Worksheet D.6.

Worksheet D.6 — CW Readout Comparison

Measurement	Value
Impulse ring-down peak amplitude A_{impulse} (mV)	12.3
CW lock-in amplitude, 1 s ($A_{\text{CW},1s}$) (mV)	28.4
CW lock-in amplitude, 10 s ($A_{\text{CW},10s}$) (mV)	91.2
Gain (1 s): $20 \cdot \log_{10}(A_{\text{CW},1s} / A_{\text{impulse}})$	7.3 dB
Gain (10 s): $20 \cdot \log_{10}(A_{\text{CW},10s} / A_{\text{impulse}})$	17.4 dB
Expected gain (1 s): 7.5 dB	Y✓
Expected gain (10 s): 17.5 dB	Y✓
Wet-finger bowing: sustained oscillation observed?	Y✓
Bowed frequency (Hz)	17,694
Bowed duration (s)	3.2
Notes on bowing technique	Slow steady stroke, 2nd attempt best

Expected results. The electronic CW gains should match predictions within ± 3 dB (measurement noise, PZT coupling variability, and frequency drift all widen the tolerance). The wet-finger bowing should produce sustained oscillation at f_1 lasting as long as the finger maintains contact—typically 2–5 seconds per stroke. Students who have played a glass harp will find this immediately intuitive.

Why this matters. Experiments 1–6 demonstrate SEM as a *rung bell*—impulse excitation, finite-duration readout. This experiment demonstrates SEM as a *bowed string*—continuous excitation, arbitrarily long integration, progressively higher precision. The two-phase readout architecture proposed in §2.3 combines both: ring the bell (fast, broadband) to find the answer, then bow the string (slow, precise) to read it exactly.

D.10 Experiment 8 — Rewritable Encoding with Water Drops

Objective: Demonstrate rewritable spectral encoding using water drops as removable mass perturbations—the transition from glass harmonica (fixed tuning) to Franklin’s armonica (reconfigurable).

Time: 45 minutes.

Materials: Assembled resonator, PicoScope, plastic transfer pipettes, water, foam cradle, styrofoam enclosure, ruler, fine-tip marker, digital thermometer.

Background. In a glass harmonica, each bowl’s pitch is set by its geometry—grind it to a specific diameter and thickness, and the frequency is fixed forever. But adding water to a bowl lowers its pitch by increasing the effective vibrating mass. The performer cannot change the glass, but *can* change the water level. This is precisely the Rayleigh perturbation mechanism: water at position x shifts mode n by an amount proportional to $\sin^2(n\pi x/L)$. Unlike wax (Experiment 5), water evaporates—making the perturbation *rewritable*. This experiment demonstrates the harmonica-to-armonica transition: write a pattern with water, read the shifted spectrum, let it evaporate (erase), write a different pattern, and confirm the spectrum changes.

Procedure:

- Prepare.** Place the resonator horizontally on the foam cradle inside the styrofoam enclosure. Equilibrate for 15 minutes with the thermometer monitoring temperature. Ensure the rod surface is clean and dry.
- Record the unperturbed spectrum (Pattern 0).** Chirp the rod and record the FFT. Log frequencies of modes 1–5.
- Write Pattern 1.** Using a transfer pipette, place one small water drop (~2 mm diameter, ~4 μL) at the rod midpoint ($x = L/2 = 75$ mm). The water should sit as a bead on the glass surface. Record the time.
- Read Pattern 1.** Immediately chirp the rod and record the FFT. Log the shifted mode frequencies. Note: mode 1 (which has an antinode at $L/2$) should shift the most. Mode 2 (which has a *node* at $L/2$) should shift the least—confirming position-dependent encoding.
- Erase Pattern 1.** Gently blot the water drop with a lint-free cloth, or simply wait 3–5 minutes for it to evaporate. Re-measure the spectrum and confirm frequencies return to within ±0.5 Hz of Pattern 0.
- Write Pattern 2.** Place a water drop at $x = L/4 = 37.5$ mm. Immediately read the spectrum. The pattern of mode shifts should be *different* from Pattern 1—mode 2 (antinode at $L/4$) now shifts strongly, while mode 1 shifts less than before.
- Erase Pattern 2.** Remove the water and confirm spectral recovery.
- Write Pattern 3 (two drops).** Place drops at both $L/4$ and $3L/4$. The resulting shift pattern is distinct from both Pattern 1 and Pattern 2.
- Record results** in Worksheet D.7.

Worksheet D.7 — Rewritable Water-Drop Encoding

Mode	f (Pattern 0)	f (Pattern 1: L/2)	Δf_1 (Hz)	f (Pattern 2: L/4)	Δf_2 (Hz)	f (Pattern 3: L/4 + 3L/4)	Δf_3 (Hz)
1	17,693.2	17,689.5	-3.7	17,691.3	-1.9	17,689.5	-3.7
2	35,388.1	35,388.0	-0.1	35,380.7	-7.4	35,373.2	-14.9
3	53,081.0	53,069.9	-11.1	53,075.9	-5.6	53,069.9	-11.1
4	70,775.3	70,775.2	-0.1	70,775.2	-0.1	70,775.2	-0.1
5	88,466.1	88,447.5	-18.6	88,456.8	-9.3	88,447.5	-18.6

Verification	Result
Pattern 0 recovered after erasing Pattern 1? (±0.5 Hz)	Y✓
Pattern 0 recovered after erasing Pattern 2? (±0.5 Hz)	Y✓
Pattern 1 and Pattern 2 are spectrally distinguishable?	Y✓
Pattern 3 is distinct from both Pattern 1 and Pattern 2?	Y✓
Total distinct patterns written and read in this experiment	3

Expected results. Water drops are lighter than wax pellets, so frequency shifts will be smaller—typically 0.2–2 Hz per mode, depending on drop volume. The thermal enclosure is essential here to keep drift below the perturbation signal. The key observation is that the three patterns produce three *different* spectral fingerprints, and that erasing (removing water) fully recovers the baseline—demonstrating the write/read/erase cycle. If shifts are too small to resolve, use larger drops or a more sensitive frequency measurement (CW lock-in from Experiment 7 provides higher precision).

The Franklin insight. This experiment turns the glass rod into a primitive armonica. The glass itself never changes—only the water does. You have just demonstrated the separation principle described in §12.5: the rod is optimized for resonance quality (Q), and the “write medium” (water) is a separate, removable layer. Franklin would recognize this immediately: same glass, different tuning, reconfigurable.

D.11 Consolidated Experiment Log

Photocopy this page for each student group or session. Attach completed Worksheets D.1–D.7.

Field	Entry
Experimenter name(s)	
Date	
School / Institution	
Rod serial # (label each rod)	
Rod length L (mm)	
Rod diameter d (mm)	
Rod mass M (g)	
PZT disc serial #	
PicoScope model & serial	
Room temperature at start (°C)	
Relative humidity (%)	
Foam cradle material	
Thermal enclosure used? (Y/N)	
Experiments completed (circle)	1 2 3 4 5 6 7 8
Best Q measured	
Number of confirmed longitudinal modes	
Best discrimination margin (dB)	
CW lock-in gain at 10 s (dB)	
Wet-finger bowing successful? (Y/N)	
Water-drop patterns written & erased	
Anomalies or unexpected observations	

D.12 Troubleshooting Guide

Symptom	Likely cause	Fix
No signal when tapping rod	Broken PZT lead or bad BNC connection	Inspect and re-solder PZT leads; verify BNC cable continuity with a multimeter
$Q \ll 100$	Excessive glue (thick viscoelastic layer)	Remove PZT, clean end face, re-glue with a smaller drop; cure 24 hours
$Q \ll 500$	PZT too large, or rod touching a hard surface	Ensure foam cradle has soft V-notches with no metal or plastic contact; try 5 mm PZT disc
Many peaks between expected modes	Off-center PZT exciting transverse modes	Remove PZT, re-glue centered using tape cross-hair; or filter FFT to comb frequencies only
f_1 much higher or lower than 17,700 Hz	Rod length $\neq 150$ mm	Recalculate: $f_1 = 5,315/(2L)$. A 200 mm rod gives $f_1 = 13,288$ Hz; a 125 mm rod gives $f_1 = 21,260$ Hz
Mode frequencies drifting over minutes	Temperature fluctuation	Use styrofoam enclosure; equilibrate 30 min; record temperature at each data point
Wax won’t stick to rod	Surface too smooth, or wax too cold	Warm wax between fingers for 30 s; lightly scuff rod surface with 400-grit sandpaper at the attachment point
FFT shows 60/120/180 Hz peaks	Electrical mains noise pickup	Use shorter BNC cables; move away from power strips and monitors; ensure PicoScope is USB-powered (not near wall adapter)
Response identical for all query patterns	Drive amplitude in non-linear regime	Reduce AWG amplitude to ≤ 0.5 Vpp; verify linearity per Exp. 5 procedure
A strong peak appears that is not on the mode comb	PZT self-resonance	Compare against bare-PZT spectrum (see Exp. 2 FM 4 mitigation); flag and exclude from analysis
Rayleigh prediction error > 10%	Wax mass too large relative to rod, or position measurement inaccurate	Use smaller wax ($\ll 0.1$ mg); measure position to ± 0.5 mm; weigh rod precisely
CW lock-in gain much less than expected	AWG frequency not precisely at mode peak; PZT coupling asymmetric	Fine-tune AWG frequency in 0.1 Hz steps while watching lock-in amplitude; ensure firm PZT bond
No sustained tone when bowing with wet finger	Finger too dry, too wet, or pressure too light	Re-wet finger to a thin film; apply firm, steady pressure; move slowly (1–2 cm/s); try rosined cloth as alternative
Water drop slides off rod immediately	Rod surface too smooth or tilted	Ensure rod is perfectly horizontal; use smaller drops (~2 mm); lightly breathe on the surface to create a condensation film
Water-drop frequency shifts too small to resolve	Drop too small relative to rod mass	Use larger drops (up to ~5 mm diameter); or use CW lock-in (Exp. 7) for higher frequency resolution

D.13 Contributing Your Data

We invite all experimenters—students, teachers, hobbyists, and researchers—to submit completed worksheets and raw PicoScope data files to the project repository. Community data from diverse rod lengths, diameters, glass types, and environments will strengthen the empirical foundation of SEM and accelerate the transition from macro prototype to MEMS fabrication.

To contribute:

- Photograph or scan your completed Worksheets D.1–D.7 and the Experiment Log (D.11).
- Export raw PicoScope waveform files (.psdata or .csv) for each experiment.
- Submit via pull request to the project repository at github.com/miketierce/wcfoma in the `data/community/` directory. Include your Experiment Log as the commit message or PR description.
- Alternatively, email data files and scanned worksheets to the corresponding author.

Every data point matters. A middle school classroom measuring $Q = 3,000$ on a 200 mm rod constrains the same physics as a university lab measuring $Q = 9,000$ on a 150 mm rod—and the diversity of rod geometries and construction techniques is itself scientifically valuable. Unexpected results (low Q , spurious modes, anomalous thermal coefficients) are *especially* welcome: they reveal failure modes that improve the guide for the next experimenter.

All quantitative claims computed from first-principles simulation code (33 modules, 959 automated tests, all passing) and independently validated by finite element analysis. No curve fitting, no adjusted parameters, no post-hoc corrections. Repository: github.com/miketierce/wcfoma.